

Transformer_Captioning

September 10, 2026

```
[ ]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment3/'
FOLDERNAME = 'cs231n/assignments/assignment3/'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the COCO dataset to your Drive
# if it doesn't already exist.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
# !bash get_datasets.sh
!bash get_coco_dataset.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

```
Mounted at /content/drive
/content/drive/My Drive/cs231n/assignments/assignment3/cs231n/datasets
--2026-09-06 12:01:58-- http://cs231n.stanford.edu/coco_captioning.zip
Resolving cs231n.stanford.edu (cs231n.stanford.edu)... 171.64.64.64
Connecting to cs231n.stanford.edu (cs231n.stanford.edu)|171.64.64.64|:80..
connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://cs231n.stanford.edu/coco_captioning.zip [following]
--2026-09-06 12:01:58-- https://cs231n.stanford.edu/coco_captioning.zip
Connecting to cs231n.stanford.edu (cs231n.stanford.edu)|171.64.64.64|:443..
connected.
HTTP request sent, awaiting response... 200 OK
Length: 1035210391 (987M) [application/zip]
Saving to: 'coco_captioning.zip'

coco_captioning.zip 100%[=====>] 987.25M 33.2MB/s in 24s
```

2026-09-06 12:02:22 (40.7 MB/s) - 'coco_captioning.zip' saved
[1035210391/1035210391]

```
Archive:  coco_captioning.zip
  creating: coco_captioning/
  inflating: coco_captioning/coco2014_captions.h5
  inflating: coco_captioning/coco2014_vocab.json
  inflating: coco_captioning/train2014_images.txt
  inflating: coco_captioning/train2014_urls.txt
  inflating: coco_captioning/train2014_vgg16_fc7.h5
  inflating: coco_captioning/train2014_vgg16_fc7_pca.h5
  inflating: coco_captioning/val2014_images.txt
  inflating: coco_captioning/val2014_urls.txt
  inflating: coco_captioning/val2014_vgg16_fc7.h5
  inflating: coco_captioning/val2014_vgg16_fc7_pca.h5
/content/drive/My Drive/cs231n/assignments/assignment3
```

1 Image Captioning with Transformers

You have now implemented a vanilla RNN and for the task of image captioning. In this notebook you will implement key pieces of a transformer decoder to accomplish the same task.

NOTE: This notebook will be primarily written in PyTorch rather than NumPy, unlike the RNN notebook.

```
[ ]: # Setup cell.
import time, os, json
import numpy as np
import matplotlib.pyplot as plt
import sys
import types
import importlib

from cs231n.gradient_check import eval_numerical_gradient, \
    eval_numerical_gradient_array
from cs231n.transformer_layers import *
from cs231n.captioning_solver_transformer import CaptioningSolverTransformer
from cs231n.classifiers.transformer import CaptioningTransformer
from cs231n.coco_utils import load_coco_data, sample_coco_minibatch, \
    decode_captions
from cs231n.image_utils import image_from_url

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # Set default size of plots.
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'
```

```

if "imp" not in sys.modules:
    imp = types.ModuleType("imp")
    imp.reload = importlib.reload
    sys.modules["imp"] = imp

%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

```

2 COCO Dataset

As in the previous notebooks, we will use the COCO dataset for captioning.

```

[ ]: # Load COCO data from disk into a dictionary.
data = load_coco_data(pca_features=True)

# Print out all the keys and values from the data dictionary.
for k, v in data.items():
    if type(v) == np.ndarray:
        print(k, type(v), v.shape, v.dtype)
    else:
        print(k, type(v), len(v))

```

```

base dir /content/drive/MyDrive/cs231n/assignments/assignment3/cs231n/datasets/
coco_captioning
train_captions <class 'numpy.ndarray'> (400135, 17) int32
train_image_idxes <class 'numpy.ndarray'> (400135,) int32
val_captions <class 'numpy.ndarray'> (195954, 17) int32
val_image_idxes <class 'numpy.ndarray'> (195954,) int32
train_features <class 'numpy.ndarray'> (82783, 512) float32
val_features <class 'numpy.ndarray'> (40504, 512) float32
idx_to_word <class 'list'> 1004
word_to_idx <class 'dict'> 1004
train_urls <class 'numpy.ndarray'> (82783,) <U63
val_urls <class 'numpy.ndarray'> (40504,) <U63

```

3 Transformer

As you have seen, RNNs are incredibly powerful but often slow to train. Further, RNNs struggle to encode long-range dependencies (though LSTMs are one way of mitigating the issue). In 2017, Vaswani et al introduced the Transformer in their paper [“Attention Is All You Need”](#) to a) introduce parallelism and b) allow models to learn long-range dependencies. The paper not only led to famous models like BERT and GPT in the natural language processing community, but also an explosion

of interest across fields, including vision. While here we introduce the model in the context of image captioning, the idea of attention itself is much more general.

4 Transformer: Multi-Headed Attention

4.0.1 Dot-Product Attention

Recall that attention can be viewed as an operation on a query $q \in \mathbb{R}^d$, a set of value vectors $\{v_1, \dots, v_n\}, v_i \in \mathbb{R}^d$, and a set of key vectors $\{k_1, \dots, k_n\}, k_i \in \mathbb{R}^d$, specified as

$$c = \sum_{i=1}^n v_i \alpha_i \quad (1)$$

$$\alpha_i = \frac{\exp(k_i^\top q)}{\sum_{j=1}^n \exp(k_j^\top q)} \quad (2)$$

$$(3)$$

where α_i are frequently called the “attention weights”, and the output $c \in \mathbb{R}^d$ is a correspondingly weighted average over the value vectors.

4.0.2 Self-Attention

In Transformers, we perform self-attention, which means that the values, keys and query are derived from the input $X \in \mathbb{R}^{\ell \times d}$, where ℓ is our sequence length. Specifically, we learn parameter matrices $V, K, Q \in \mathbb{R}^{d \times d}$ to map our input X as follows:

$$v_i = Vx_i \quad i \in \{1, \dots, \ell\} \quad (4)$$

$$k_i = Kx_i \quad i \in \{1, \dots, \ell\} \quad (5)$$

$$q_i = Qx_i \quad i \in \{1, \dots, \ell\} \quad (6)$$

4.0.3 Multi-Headed Scaled Dot-Product Attention

In the case of multi-headed attention, we learn a parameter matrix for each head, which gives the model more expressivity to attend to different parts of the input. Let h be number of heads, and Y_i be the attention output of head i . Thus we learn individual matrices Q_i, K_i and V_i . To keep our overall computation the same as the single-headed case, we choose $Q_i \in \mathbb{R}^{d \times d/h}$, $K_i \in \mathbb{R}^{d \times d/h}$ and $V_i \in \mathbb{R}^{d \times d/h}$. Adding in a scaling term $\frac{1}{\sqrt{d/h}}$ to our simple dot-product attention above, we have

$$Y_i = \text{softmax} \left(\frac{(XQ_i)(XK_i)^\top}{\sqrt{d/h}} \right) (XV_i) \quad (7)$$

where $Y_i \in \mathbb{R}^{\ell \times d/h}$, where ℓ is our sequence length.

In our implementation, we apply dropout to the attention weights (though in practice it could be used at any step):

$$Y_i = \text{dropout} \left(\text{softmax} \left(\frac{(XQ_i)(XK_i)^\top}{\sqrt{d/h}} \right) \right) (XV_i) \quad (8)$$

Finally, then the output of the self-attention is a linear transformation of the concatenation of the heads:

$$Y = [Y_1; \dots; Y_h]A \quad (9)$$

were $A \in \mathbb{R}^{d \times d}$ and $[Y_1; \dots; Y_h] \in \mathbb{R}^{\ell \times d}$.

Implement multi-headed scaled dot-product attention in the `MultiHeadAttention` class in the file `cs231n/transformer_layers.py`. The code below will check your implementation. The relative error should be less than $7\text{e-}3$.

```
[ ]: torch.manual_seed(231)

# Choose dimensions such that they are all unique for easier debugging:
# Specifically, the following values correspond to N=1, H=2, T=3, E//H=4, and
# E=8.
batch_size = 1
sequence_length = 3
embed_dim = 8
attn = MultiHeadAttention(embed_dim, num_heads=2)
attn.eval()

# Self-attention.
data = torch.randn(batch_size, sequence_length, embed_dim)
self_attn_output = attn(query=data, key=data, value=data)

# Masked self-attention.
mask = torch.randn(sequence_length, sequence_length) < 0.5
masked_self_attn_output = attn(query=data, key=data, value=data, attn_mask=mask)

# Attention using two inputs.
other_data = torch.randn(batch_size, sequence_length, embed_dim)
attn_output = attn(query=data, key=other_data, value=other_data)

expected_self_attn_output = np.asarray([[
    [-0.1030,  0.2090,  0.7481, -0.4116, -0.2772,  0.2357, -0.2740, -0.2798],
    [-0.1794,  0.1637,  0.6930, -0.3435, -0.2300,  0.2457, -0.2504, -0.2323],
    [-0.1861,  0.1567,  0.6677, -0.3078, -0.2731,  0.2302, -0.2714, -0.2082]]])

expected_masked_self_attn_output = np.asarray([[
    [-0.1209,  0.1806,  0.8056, -0.4654, -0.2262,  0.2552, -0.2544, -0.2812],
    [-0.2374,  0.1183,  0.7423, -0.3777, -0.1460,  0.2716, -0.2213, -0.2077],
    [-0.0851,  0.2423,  0.7941, -0.3778, -0.4924,  0.1512, -0.4241, -0.1740]]])
```

```

expected_attn_output = np.asarray([[
    [-0.1965, 0.0702, 0.5227, -0.1388, -0.3530, 0.2163, -0.2643, -0.1902],
    [-0.3592, -0.0899, 0.5836, 0.0029, -0.3643, 0.2667, -0.1973, -0.1698],
    [-0.3795, 0.0433, 0.5634, -0.0749, -0.4379, 0.2714, -0.2207, -0.1980]]])

print('self_attn_output error: ', rel_error(expected_self_attn_output,
    ↪self_attn_output.detach().numpy()))
print('masked_self_attn_output error: ',
    ↪rel_error(expected_masked_self_attn_output, masked_self_attn_output.detach().
    ↪numpy()))
print('attn_output error: ', rel_error(expected_attn_output, attn_output.
    ↪detach().numpy()))

```

```

self_attn_output error: 0.00017949659048965922
masked_self_attn_output error: 9.809519326315612e-05
attn_output error: 0.006226898176165351

```

5 Positional Encoding

While transformers are able to easily attend to any part of their input, the attention mechanism has no concept of token order. However, for many tasks (especially natural language processing), relative token order is very important. To recover this, the authors add a positional encoding to the embeddings of individual word tokens.

Let us define a matrix $P \in \mathbb{R}^{l \times d}$, where $P_{ij} =$

$$\begin{cases} \sin\left(i \cdot 10000^{-\frac{j}{d}}\right) & \text{if } j \text{ is even} \\ \cos\left(i \cdot 10000^{-\frac{(j-1)}{d}}\right) & \text{otherwise} \end{cases}$$

Rather than directly passing an input $X \in \mathbb{R}^{l \times d}$ to our network, we instead pass $X + P$.

Implement this layer in `PositionalEncoding` in `cs231n/transformer_layers.py`. Once you are done, run the following to perform a simple test of your implementation. You should see errors on the order of $e-3$ or less.

```

[ ]: torch.manual_seed(231)

batch_size = 1
sequence_length = 2
embed_dim = 6
data = torch.randn(batch_size, sequence_length, embed_dim)

pos_encoder = PositionalEncoding(embed_dim)
pos_encoder.eval()
output = pos_encoder(data)

```

```

expected_pe_output = np.asarray([[-1.1106, 1.0014, 1.5280, -0.0778, -0.6964,
↪ 1.1455],
                                [ 0.8125, -0.4303, 0.4981, 0.7320, 1.1380,
↪ 1.5331]])
print('pe_output error: ', rel_error(expected_pe_output, output.detach().
↪ numpy()))

```

pe_output error: 0.00021699471301977143

6 Inline Question 1

Several key design decisions were made in designing the scaled dot product attention we introduced above. Explain why the following choices were beneficial: 1. Using multiple attention heads as opposed to one. 2. Dividing by $\sqrt{d/h}$ before applying the softmax function. Recall that d is the feature dimension and h is the number of heads. 3. Adding a linear transformation to the output of the attention operation.

Only one or two sentences per choice is necessary, but be sure to be specific in addressing what would have happened without each given implementation detail, why such a situation would be suboptimal, and how the proposed implementation improves the situation.

Your Answer:

1. Using multiple attention heads as opposed to one. Splitting the d -dimensional representation into h heads, each of dimension $d_{qk} = d_v = d/h$, lets each head dedicate its own subspace to modeling a *different* aspect of the sequence (e.g., distinct syntactic or semantic roles), since each head produces its own “relevance distribution” over all tokens. With only one head, all d dimensions would have to share a single attention distribution, which suppresses this “dedicated dimensions for dedicated semantics” effect and weakens the model’s representational capacity.

2. Dividing by $\sqrt{d/h}$ before applying the softmax function. If the components of a query vector $q \in \mathbb{R}^{d_{qk}}$ and a key vector $k^\top \in \mathbb{R}^{d_{qk}}$ are i.i.d. standard normal, the dot product $x = qk^\top = \sum_{i=1}^{d_{qk}} q_i k_i^\top$ has $E(x) = 0$ and $\text{Var}(x) = d_{qk} = d/h$, which is large, so x is likely to take an extreme value and $\text{softmax}(x)$ would saturate near 1, letting a single token monopolize almost all of the attention probability (and killing gradients). Dividing by $\sqrt{d/h}$ rescales the variance to 1, so the softmax outputs stay moderate and well-conditioned.

3. Adding a linear transformation to the output of the attention operation. The output projection $W_{\text{Output}} \in \mathbb{R}^{h d_v \times d}$ mixes the concatenated per-head outputs $\text{concatenate}(\text{head}_1, \dots, \text{head}_h) \in \mathbb{R}^{n \times h d_v}$, allowing features learned independently in different heads’ subspaces to interact, and maps the result back to $\mathbb{R}^{n \times d}$ so it can be combined with the residual stream. Without it, each head would remain an isolated information channel unable to exchange information with the others, and the output could not even be residually added to $X \in \mathbb{R}^{n \times d}$, both of which would reduce expressiveness.

7 Transformer Decoder Block

Transformer decoder layer consists of three modules: (1) self attention to process input sequence of vectors, (2) cross attention to process based on available context (i.e. image features in our

case), (3) feedforward module to process each vector of the sequence independently. Complete the implementation of `TransformerDecoderLayer` in `cs231n/transformer_layers.py` and test it below. The relative error should be less than $1e-6$.

The Transformer decoder layer has three main components: (1) a self-attention module that processes the input sequence of vectors, (2) a cross-attention module that incorporates additional context (e.g., image features in our case), and (3) a feedforward module that independently processes each vector in the sequence. Complete the implementation of `TransformerDecoderLayer` in `cs231n/transformer_layers.py` and test it below. The relative error should be less than $1e-6$.

```
[ ]: torch.manual_seed(231)
      np.random.seed(231)

      N, T, TM, D = 1, 4, 5, 12

      decoder_layer = TransformerDecoderLayer(D, 2, 4*D)
      decoder_layer.eval()
      tgt = torch.randn(N, T, D)
      memory = torch.randn(N, TM, D)
      tgt_mask = torch.randn(T, T) < 0.5

      output = decoder_layer(tgt, memory, tgt_mask)

      expected_output = np.asarray([
          [[ 0.27005595, -0.21056601,  0.63376445, -0.35391954,  0.60071295,
            -1.7566615, -0.1314159, -1.7659538, -0.5575517,  0.680143,
             1.8622686,  0.7291234],
          [-0.6080135,  0.4421802, -0.07590749, -0.12093157,  0.94722354,
            -1.0223433,  0.47332686, -1.3958666,  2.1890075, -1.0380646,
             0.9437816, -0.7343922],
          [-0.89690155,  0.60447216,  0.10688882,  0.03492759,  1.8388084,
            -0.9772293,  0.7621334, -0.668234,  1.3260399, -1.8417019,
             0.21718435, -0.506388],
          [-0.70961505, -0.02918372, -0.39894593, -0.8069395,  0.04967685,
             0.29345992,  0.7871747, -1.3172938,  1.4744806,  1.2633642,
             1.1611108, -1.767289]]
      ])
      print('error: ', rel_error(expected_output, output.detach().numpy()))
```

error: 8.386729410631422e-08

8 Transformer for Image Captioning

Now that you have implemented the previous layers, you can combine them to build a Transformer-based image captioning model. Open the file `cs231n/classifiers/transformer.py` and look at the `CaptioningTransformer` class.

Implement the `forward` function of the class. After doing so, run the following to check your forward pass using a small test case; you should see error on the order of $e-5$ or less.


```
[ ]: torch.manual_seed(231)
      np.random.seed(231)

      N, D, W = 4, 20, 30
      word_to_idx = {'<NULL>': 0, 'cat': 2, 'dog': 3}
      V = len(word_to_idx)
      T = 3

      transformer = CaptioningTransformer(
          word_to_idx,
          input_dim=D,
          wordvec_dim=W,
          num_heads=2,
          num_layers=2,
          max_length=30
      )
      transformer.eval()

      features = torch.randn(N, D)
      captions = torch.randint(0, V, (N, T))

      scores = transformer(features, captions)

      expected_scores = np.asarray([
          [[ 0.62165695, -0.6203744, -1.116114 ],
           [ 0.6641097, -0.60309553, -1.3479083 ],
           [ 0.62117404, -0.79205745, -1.1095243 ]],
          [[ 0.62160987, -0.6214724, -1.1171186 ],
           [ 0.6093992, -0.61166143, -1.3365762 ],
           [ 0.62108713, -0.79322916, -1.1104958 ]],
          [[ 0.6406587, -0.57419956, -1.1300566 ],
           [ 0.6075228, -0.60872054, -1.3372457 ],
           [ 0.6471472, -0.8270494, -1.1064028 ]],
          [[ 0.64029807, -0.5755552, -1.1284486 ],
           [ 0.60708034, -0.6099963, -1.3356377 ],
           [ 0.59367496, -0.8346836, -1.0914496 ]]]
      ])

      print('scores error: ', rel_error(expected_scores, scores.detach().numpy()))
```

scores error: 1.4597957790730836e-07

9 Overfit Transformer Captioning Model on Small Data

Run the following to overfit the Transformer-based captioning model on the same small dataset as we used for the RNN previously.

```
[ ]: torch.manual_seed(231)
      np.random.seed(231)

      data = load_coco_data(max_train=50)

      transformer = CaptioningTransformer(
          word_to_idx=data['word_to_idx'],
          input_dim=data['train_features'].shape[1],
          wordvec_dim=256,
          num_heads=2,
          num_layers=2,
          max_length=30
      )

      transformer_solver = CaptioningSolverTransformer(transformer, data,
          ↪idx_to_word=data['idx_to_word'],
              num_epochs=100,
              batch_size=25,
              learning_rate=0.001,
              verbose=True, print_every=10,
          )


      transformer_solver.train()

      # Plot the training losses.
      plt.plot(transformer_solver.loss_history)
      plt.xlabel('Iteration')
      plt.ylabel('Loss')
      plt.title('Training loss history')
      plt.show()
```

```
base dir /content/drive/MyDrive/cs231n/assignments/assignment3/cs231n/datasets/
coco_captioning
```

```
(Iteration 1 / 200) loss: 5.045327
(Iteration 11 / 200) loss: 2.808727
(Iteration 21 / 200) loss: 2.022354
(Iteration 31 / 200) loss: 1.670440
(Iteration 41 / 200) loss: 1.330231
(Iteration 51 / 200) loss: 1.236568
(Iteration 61 / 200) loss: 0.843831
(Iteration 71 / 200) loss: 0.702294
(Iteration 81 / 200) loss: 0.516776
(Iteration 91 / 200) loss: 0.438581
(Iteration 101 / 200) loss: 0.376332
(Iteration 111 / 200) loss: 0.153824
(Iteration 121 / 200) loss: 0.139354
```

```
(Iteration 131 / 200) loss: 0.104465
(Iteration 141 / 200) loss: 0.084865
(Iteration 151 / 200) loss: 0.087971
(Iteration 161 / 200) loss: 0.053218
(Iteration 171 / 200) loss: 0.046330
(Iteration 181 / 200) loss: 0.043136
(Iteration 191 / 200) loss: 0.027111
```



Print final training loss. You should see a final loss of less than 0.05 .

```
[ ]: print('Final loss: ', transformer_solver.loss_history[-1])
```

Final loss: 0.022365998

10 Transformer Sampling at Test Time

The sampling code has been written for you. You can simply run the following to compare with the previous results with the RNN. As before the training results should be much better than the validation set results, given how little data we trained on.

```
[ ]: # If you get an error, the URL just no longer exists, so don't worry!
# You can re-sample as many times as you want.
for split in ['train', 'val']:
    minibatch = sample_coco_minibatch(data, split=split, batch_size=2)
    gt_captions, features, urls = minibatch
    gt_captions = decode_captions(gt_captions, data['idx_to_word'])

    sample_captions = transformer.sample(features, max_length=30)
    sample_captions = decode_captions(sample_captions, data['idx_to_word'])

    for gt_caption, sample_caption, url in zip(gt_captions, sample_captions,
↪urls):
        img = image_from_url(url)
        # Skip missing URLs.
        if img is None: continue
        plt.imshow(img)
        plt.title('%s\n%s\nGT:%s' % (split, sample_caption, gt_caption))
        plt.axis('off')
        plt.show()
```

URL Error: Bad Gateway

http://farm1.staticflickr.com/202/487987371_489a65d670_z.jpg

train
a <UNK> decorated living room with a big tv in it <END>
GT:<START> a <UNK> decorated living room with a big tv in it <END>



URL Error: Bad Gateway

http://farm1.staticflickr.com/25/44101107_9491d72776_z.jpg

val
a man <UNK> with a bite on the busy while having his picture taken <END>
GT:<START> a group of people <UNK> outside by a wall <END>



11 Vision Transformer (ViT)

Dosovitskiy et. al. showed that applying a transformer model on a sequence of image patches (referred to as Vision Transformer) not only achieves impressive performance but also scales more effectively than convolutional neural networks when trained on large datasets. We will build a version of Vision Transformer using our existing implementation of transformer components and train it on the CIFAR-10 dataset.

Vision Transformer converts input image into a sequence of patches of fixed size and embed each patch into a latent vector. In `cs231n/transformer_layers.py`, complete the implementation of `PatchEmbedding` and test it below. You should see relative error less than $1e-4$.

```
[ ]: from cs231n.transformer_layers import PatchEmbedding

torch.manual_seed(231)
np.random.seed(231)

N = 2
```

```

HW = 16
PS = 8
D = 8

patch_embedding = PatchEmbedding(
    img_size=HW,
    patch_size=PS,
    embed_dim=D
)

x = torch.randn(N, 3, HW, HW)
output = patch_embedding(x)

expected_output = np.asarray([
    [[-0.6312704 ,  0.02531429,  0.6112642 , -0.49089882,
      0.01412961, -0.6959372 , -0.32862484, -0.45402682],
     [ 0.18816411, -0.08142513, -0.9829535 , -0.23975623,
     -0.23109074,  0.97950286, -0.40997326,  0.7457837 ],
     [ 0.01810865,  0.15780598, -0.91804236,  0.36185235,
      0.8379501 ,  1.0191797 , -0.29667392,  0.20322265],
     [-0.18697818, -0.45137224, -0.40339014, -1.4381214 ,
     -0.43450755,  0.7651071 , -0.83683825, -0.16360264]]],

    [[-0.39786366,  0.16201034, -0.19008337, -1.0602452 ,
     -0.28693503,  0.09791763,  0.26614824,  0.41781986],
     [ 0.35146567, -0.4469593 , -0.1841726 ,  0.45757473,
     -0.61304873, -0.29104248, -0.16124889, -0.14987172],
     [-0.2996967 ,  0.27353522, -0.09929767,  0.01973832,
     -1.2312065 , -0.6374332 , -0.22963578,  0.55696607],
     [-0.93818814,  0.02465284, -0.21117875,  1.1860403 ,
     -0.06137538, -0.21062079, -0.094347 ,  0.50032747]]])

print('error: ', rel_error(expected_output, output.detach().numpy()))

```

error: 9.182286955268188e-08

The sequence of patch vectors is processed by transformer encoder layers, each consisting of a self-attention and a feed-forward module. Since all vectors attend to one another, attention masking is not strictly necessary. However, we still implement it for the sake of consistency.

Implement `TransformerEncoderLayer` in `cs231n/transformer_layers.py` and test it below. You should see relative error less than $1e-6$.

```

[ ]: torch.manual_seed(231)
     np.random.seed(231)

     from cs231n.transformer_layers import TransformerEncoderLayer

```

```

N, T, TM, D = 1, 4, 5, 12

encoder_layer = TransformerEncoderLayer(D, 2, 4*D)
encoder_layer.eval()
x = torch.randn(N, T, D)
x_mask = torch.randn(T, T) < 0.5

output = encoder_layer(x, x_mask)

expected_output = np.asarray([
    [-0.48887438, -0.16470742, 0.7249629, -1.0907345, 1.5764195,
     0.17484017, 0.8643941, -0.5209561, -1.6651545, -1.1943513,
     1.4809005, 0.30326116],
    [-0.14663944, 1.052812, -0.96227, -0.20350897, -2.1478589,
     1.2637341, 0.04015125, -0.945255, 1.408903, -0.43051198,
     0.37110543, 0.69933903],
    [0.08718029, 0.91050833, -1.208469, 1.1440394, -1.9794976,
     -0.38342738, -0.47407156, -0.17345548, 0.11738186, 1.9534125,
     0.3525986, -0.3461999],
    [-0.81389004, 0.76213264, -0.7094313, -1.5332246, 1.3837544,
     0.4091982, 0.6569406, 0.7305365, -1.197077, -1.3698943,
     0.58292854, 1.0980265]]
])

print('error: ', rel_error(expected_output, output.detach().numpy()))

```

error: 4.162545799227941e-08

Take a look at the VisionTransformer implementation in `cs231n/classifiers/transformer.py`.

For classification, ViT divides the input image into patches and processes the sequence of patch vectors using a transformer. Finally, all the patch vectors are average-pooled and used to predict the image class. We will use the same 1D sinusoidal positional encoding to inject ordering information, though 2D sinusoidal and learned positional encodings are also valid choices.

Complete the ViT forward pass and test it below. You should see relative error less than $1e-6$.

```

[ ]: torch.manual_seed(231)
      np.random.seed(231)
      from cs231n.classifiers.transformer import VisionTransformer

      imgs = torch.randn(3, 3, 32, 32)
      transformer = VisionTransformer()
      transformer.eval()

      scores = transformer(imgs)

      expected_scores = np.asarray(
          [[-0.15802915, 0.15734261, -0.06002094, -0.1860142, -0.09376919,

```



```

-0.08837911, 0.16224845, 0.0168601, 0.20606893, 0.13079852],
[-0.12838936, 0.20798579, -0.0620573, -0.11705838, -0.14088857,
-0.11387511, 0.18839046, 0.03078492, 0.17386372, 0.06414902],
[-0.09435399, 0.20495278, -0.03362986, -0.12232997, -0.15535004,
-0.10380758, 0.21114832, 0.03481449, 0.18798208, 0.12231653]])

print('scores error: ', rel_error(expected_scores, scores.detach().numpy()))

```

scores error: 2.3590082771647052e-07

We will first verify our implementation by overfitting it on one training batch. Tune learning rate and weight decay accordingly.

```

[ ]: from torchvision import transforms
from torchvision.datasets import CIFAR10
from tqdm.auto import tqdm
from torch.utils.data import DataLoader

train_data = CIFAR10(root='data', train=True, transform=transforms.ToTensor(),
    ↪download=True)
test_data = CIFAR10(root='data', train=False, transform=transforms.ToTensor(),
    ↪download=True)

```

100%| | 170M/170M [05:04<00:00, 559kB/s]

```

[ ]: learning_rate = 1e-3 # Experiment with this
weight_decay = 1.e-4 # Experiment with this

batch = next(iter(DataLoader(train_data, batch_size=64, shuffle=False)))
model = VisionTransformer(dropout=0.0)
loss_criterion = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate,
    ↪weight_decay=weight_decay)
model.train()

epochs = 100
for epoch in range(epochs):
    imgs, target = batch
    out = model(imgs)
    loss = loss_criterion(out, target)

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    top1 = (out.argmax(-1) == target).float().mean().item()
    if epoch % 10 == 0:

```

```
print(f"[{epoch}/{epochs}] Loss {loss.item():.6f}, Top-1 Accuracy: {top1:.3f}")
```

```
[0/100] Loss 2.324255, Top-1 Accuracy: 0.062
[10/100] Loss 2.088192, Top-1 Accuracy: 0.234
[20/100] Loss 1.914399, Top-1 Accuracy: 0.297
[30/100] Loss 1.613288, Top-1 Accuracy: 0.453
[40/100] Loss 1.340390, Top-1 Accuracy: 0.547
[50/100] Loss 0.782776, Top-1 Accuracy: 0.828
[60/100] Loss 0.390877, Top-1 Accuracy: 0.938
[70/100] Loss 0.119019, Top-1 Accuracy: 1.000
[80/100] Loss 0.039872, Top-1 Accuracy: 1.000
[90/100] Loss 0.019011, Top-1 Accuracy: 1.000
```

```
[ ]: # You should get perfect 1.00 accuracy
print(f"Overfitting ViT on one batch. Top-1 accuracy: {top1}")
```

Overfitting ViT on one batch. Top-1 accuracy: 1.0

Now we will train it on the entire dataset.

```
[ ]: from cs231n.classification_solver_vit import ClassificationSolverViT

#####
# TODO: Train a Vision Transformer model that achieves over 0.45 test #
# accuracy on CIFAR-10 after 2 epochs by adjusting the model architecture #
# and/or training parameters as needed. #
# #
# Note: If you want to use a GPU runtime, go to `Runtime > Change runtime #
# type` and set `Hardware accelerator` to `GPU`. This will reset Colab, #
# so make sure to rerun the entire notebook from the beginning afterward. #
#####

learning_rate = 1e-4
weight_decay = 0.0
batch_size = 64
model = VisionTransformer()

#####
#                               END OF YOUR CODE                               #
#####

solver = ClassificationSolverViT(
    train_data=train_data,
```

```

    test_data=test_data,
    model=model,
    num_epochs = 2, # Don't change this
    learning_rate = learning_rate,
    weight_decay = weight_decay,
    batch_size = batch_size,
)

solver.train('cuda' if torch.cuda.is_available() else 'cpu')

```

```

0%|          | 0/782 [00:00<?, ?it/s]
0%|          | 0/157 [00:00<?, ?it/s]
0%|          | 0/782 [00:00<?, ?it/s]
0%|          | 0/157 [00:00<?, ?it/s]

```

```
[ ]: print(f"Accuracy on test set: {solver.results['best_test_acc']}")
```

Accuracy on test set: 0.4047

12 Inline Question 2

Despite their recent success in large-scale image recognition tasks, ViTs often lag behind traditional CNNs when trained on smaller datasets. What underlying factor contribute to this performance gap? What techniques can be used to improve the performance of ViTs on small datasets?

Your Answer:

The underlying factor is the **lack of inductive bias**. CNNs build translation equivariance, locality (local receptive fields), and hierarchical feature composition directly into their architecture, so these visual priors need not be learned from data. A ViT has almost no image-specific prior — each patch must learn spatial relationships from scratch through global self-attention, which requires far more training data to generalize; on small datasets it underfits or overfits to spurious patterns, hence lagging behind CNNs. Techniques that help ViTs on small datasets: - **Strong pre-training then fine-tuning** on large corpora (e.g., ImageNet-21k, JFT-300M), which supplies the “missing” data. - **Aggressive data augmentation** (RandAugment, Mixup, CutMix, random erasing) to artificially enlarge the effective dataset. - **Regularization**: weight decay, dropout, stochastic depth, label smoothing. - **Knowledge distillation** from a strong CNN teacher (e.g., the distillation token in DeiT). - **Architectural priors**: hybrid designs with a convolutional patch-embedding stem, relative/2D-aware position encodings, or local (windowed) attention to inject locality back into the model.

13 Inline Question 3

How does the computational cost of the self-attention layers in a ViT change if we independently make the following changes? Please ignore the computation cost of QKV and output projection.

- (i) Double the hidden dimension.

- (ii) Double the height and width of the input image.
- (iii) Double the patch size.
- (iv) Double the number of layers.

Your Answer:

Per layer, ignoring QKV/output projections, the self-attention cost is dominated by $QK^\top$ and $\text{softmax}(\cdot)V$, each $O(n^2d)$, so the total is $O(n^2d)$, where $n = N + 1$ tokens ($N = HW/P^2$ patches plus the class token) and d is the hidden/embedding dimension.

(i) Double the hidden dimension.

$$O(n^2 \cdot 2d) = 2 O(n^2d) \quad \Rightarrow \quad \mathbf{2 \times \text{cost.}}$$

(ii) Double the height and width of the input image. $N' = \frac{(2H)(2W)}{P^2} = 4N$, so $n' \approx 4n$ (for large N , the class token's +1 is negligible):

$$O((4n)^2d) = 16 O(n^2d) \quad \Rightarrow \quad \mathbf{16 \times \text{cost.}}$$

(iii) Double the patch size. $N' = \frac{HW}{(2P)^2} = N/4$, so $n' \approx n/4$:

$$O\left(\left(\frac{n}{4}\right)^2 d\right) = \frac{1}{16} O(n^2d) \quad \Rightarrow \quad \mathbf{1/16 \times \text{cost.}}$$

(iv) Double the number of layers. The per-layer cost is unchanged, but the total over all layers becomes $2L \cdot O(n^2d) \Rightarrow \mathbf{2 \times \text{total cost.}}$

[]:

Self_Supervised_Learning

September 10, 2026

```
[ ]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment3/'
FOLDERNAME = "cs231n/assignments/assignment3/"
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This downloads the COCO dataset to your Drive
# if it doesn't already exist.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
!bash get_datasets.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

```
Mounted at /content/drive
/content/drive/My Drive/cs231n/assignments/assignment3/cs231n/datasets
--2026-09-06 13:44:38-- http://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
Resolving www.cs.toronto.edu (www.cs.toronto.edu)... 128.100.3.30
Connecting to www.cs.toronto.edu (www.cs.toronto.edu)|128.100.3.30|:80...
connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://cave.cs.toronto.edu/kriz/cifar-10-python.tar.gz [following]
--2026-09-06 13:44:38-- https://cave.cs.toronto.edu/kriz/cifar-10-python.tar.gz
Resolving cave.cs.toronto.edu (cave.cs.toronto.edu)... 128.100.3.67
Connecting to cave.cs.toronto.edu (cave.cs.toronto.edu)|128.100.3.67|:443...
connected.
HTTP request sent, awaiting response... 200 OK
Length: 170498071 (163M) [application/x-gzip]
Saving to: 'cifar-10-python.tar.gz'

cifar-10-python.tar 100%[=====>] 162.60M 827KB/s in 4m 8s
```

```
2026-09-06 13:48:46 (673 KB/s) - 'cifar-10-python.tar.gz' saved
[170498071/170498071]
```

```
cifar-10-batches-py/
cifar-10-batches-py/data_batch_4
cifar-10-batches-py/readme.html
cifar-10-batches-py/test_batch
cifar-10-batches-py/data_batch_3
cifar-10-batches-py/batches.meta
cifar-10-batches-py/data_batch_2
cifar-10-batches-py/data_batch_5
cifar-10-batches-py/data_batch_1
--2026-09-06 13:48:49-- http://cs231n.stanford.edu/imagenet_val_25.npz
Resolving cs231n.stanford.edu (cs231n.stanford.edu)... 171.64.64.64
Connecting to cs231n.stanford.edu (cs231n.stanford.edu)|171.64.64.64|:80...
connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://cs231n.stanford.edu/imagenet_val_25.npz [following]
--2026-09-06 13:48:50-- https://cs231n.stanford.edu/imagenet_val_25.npz
Connecting to cs231n.stanford.edu (cs231n.stanford.edu)|171.64.64.64|:443...
connected.
HTTP request sent, awaiting response... 200 OK
Length: 3940548 (3.8M)
Saving to: 'imagenet_val_25.npz'
```

```
imagenet_val_25.npz 100%[=====>] 3.76M --.-KB/s in 0.1s
```

```
2026-09-06 13:48:50 (28.7 MB/s) - 'imagenet_val_25.npz' saved [3940548/3940548]
```

```
/content/drive/My Drive/cs231n/assignments/assignment3
```

0.1 Using GPU

Go to Runtime > Change runtime type and set Hardware accelerator to GPU. This will reset Colab. **Rerun the top cell to mount your Drive again.**

1 Self-Supervised Learning

1.1 What is self-supervised learning?

Modern day machine learning requires lots of labeled data. But often times it's challenging and/or expensive to obtain large amounts of human-labeled data. Is there a way we could ask machines to automatically learn a model which can generate good visual representations without a labeled dataset? Yes, enter self-supervised learning!

Self-supervised learning (SSL) allows models to automatically learn a “good” representation space using the data in a given dataset without the need for their labels. Specifically, if our dataset were

a bunch of images, then self-supervised learning allows a model to learn and generate a “good” representation vector for images.

The reason SSL methods have seen a surge in popularity is because the learnt model continues to perform well on other datasets as well i.e. new datasets on which the model was not trained on!

1.2 What makes a “good” representation?

A “good” representation vector needs to capture the important features of the image as it relates to the rest of the dataset. This means that images in the dataset representing semantically similar entities should have similar representation vectors, and different images in the dataset should have different representation vectors. For example, two images of an apple should have similar representation vectors, while an image of an apple and an image of a banana should have different representation vectors.


1.3 Contrastive Learning: SimCLR

Recently, [SimCLR](#) introduces a new architecture which uses **contrastive learning** to learn good visual representations. Contrastive learning aims to learn similar representations for similar images and different representations for different images. As we will see in this notebook, this simple idea allows us to train a surprisingly good model without using any labels.

Specifically, for each image in the dataset, SimCLR generates two differently augmented views of that image, called a **positive pair**. Then, the model is encouraged to generate similar representation vectors for this pair of images. See below for an illustration of the architecture (Figure 2 from the paper).

```
[ ]: # Run this cell to view the SimCLR architecture.  
from IPython.display import Image  
Image('images/simclr_fig2.png', width=500)
```

```
[ ]:
```



Given an image $\mathbf{x}$, SimCLR uses two different data augmentation schemes $\mathbf{t}$ and $\mathbf{t}'$ to generate the positive pair of images $\tilde{x}_i$ and $\tilde{x}_j$. f is a basic encoder net that extracts representation vectors from the augmented data samples, which yields h_i and h_j , respectively. Finally, a small neural network projection head g maps the representation vectors to the space where the contrastive loss is applied. The goal of the contrastive loss is to maximize agreement between the final vectors $z_i = g(h_i)$ and $z_j = g(h_j)$. We will discuss the contrastive loss in more detail later, and you will get to implement it.

After training is completed, we throw away the projection head g and only use f and the representation h to perform downstream tasks, such as classification. You will get a chance to finetune a layer on top of a trained SimCLR model for a classification task and compare its performance with a baseline model (without self-supervised learning).

1.4 Pretrained Weights

For your convenience, we have given you pretrained weights (trained for ~18 hours on CIFAR-10) for the SimCLR model. Run the following cell to download pretrained model weights to be used later. (This will take ~1 minute)

```
[ ]: %%bash
DIR=pretrained_model/
```



```

if [ ! -d "$DIR" ]; then
    mkdir "$DIR"
fi

URL=http://downloads.cs.stanford.edu/downloads/cs231n/pretrained_simclr_model.
    ↪pth
# Try this if above doesn't work.
# URL=http://cs231n.stanford.edu/2025/storage/a3/pretrained_simclr_model.pth
FILE=pretrained_model/pretrained_simclr_model.pth
if [ ! -f "$FILE" ]; then
    echo "Downloading weights..."
    wget "$URL" -O "$FILE"
fi

```

Downloading weights...

--2026-09-06 13:49:46--

http://downloads.cs.stanford.edu/downloads/cs231n/pretrained_simclr_model.pth

Resolving downloads.cs.stanford.edu (downloads.cs.stanford.edu)... 171.64.64.22

Connecting to downloads.cs.stanford.edu

(downloads.cs.stanford.edu)|171.64.64.22|:80... connected.

HTTP request sent, awaiting response... 200 OK

Length: 98808459 (94M) [application/octet-stream]

Saving to: 'pretrained_model/pretrained_simclr_model.pth'

0K	...	...	0%	1.63M	58s
50K	...	...	0%	3.23M	43s
100K	...	...	0%	63.0M	29s
150K	...	...	0%	3.41M	29s
200K	...	...	0%	40.9M	24s
250K	...	...	0%	64.8M	20s
300K	...	...	0%	69.4M	17s
350K	...	...	0%	3.89M	18s
400K	...	...	0%	42.7M	16s
450K	...	...	0%	61.5M	15s
500K	...	...	0%	38.5M	14s
550K	...	...	0%	75.0M	13s
600K	...	...	0%	83.4M	12s
650K	...	...	0%	90.1M	11s
700K	...	...	0%	4.23M	12s
750K	...	...	0%	78.3M	11s
800K	...	...	0%	79.4M	10s
850K	...	...	0%	71.3M	10s
900K	...	...	0%	88.5M	9s
950K	...	...	1%	67.9M	9s
1000K	...	...	1%	41.5M	9s
1050K	...	...	1%	67.6M	8s
1100K	...	...	1%	86.8M	8s

1150K	...	...	...	...	...	1%	89.5M	8s
1200K	...	...	...	...	...	1%	93.0M	8s
1250K	...	...	...	...	...	1%	77.7M	7s
1300K	...	...	...	...	...	1%	86.9M	7s
1350K	...	...	...	...	...	1%	90.3M	7s
1400K	...	...	...	...	...	1%	8.34M	7s
1450K	...	...	...	...	...	1%	31.9M	7s
1500K	...	...	...	...	...	1%	40.9M	7s
1550K	...	...	...	...	...	1%	82.8M	6s
1600K	...	...	...	...	...	1%	66.4M	6s
1650K	...	...	...	...	...	1%	67.7M	6s
1700K	...	...	...	...	...	1%	86.6M	6s
1750K	...	...	...	...	...	1%	87.4M	6s
1800K	...	...	...	...	...	1%	88.7M	6s
1850K	...	...	...	...	...	1%	77.4M	6s
1900K	...	...	...	...	...	2%	79.0M	6s
1950K	...	...	...	...	...	2%	86.1M	5s
2000K	...	...	...	...	...	2%	52.1M	5s
2050K	...	...	...	...	...	2%	35.3M	5s
2100K	...	...	...	...	...	2%	87.4M	5s
2150K	...	...	...	...	...	2%	76.0M	5s
2200K	...	...	...	...	...	2%	87.4M	5s
2250K	...	...	...	...	...	2%	67.0M	5s
2300K	...	...	...	...	...	2%	22.9M	5s
2350K	...	...	...	...	...	2%	35.7M	5s
2400K	...	...	...	...	...	2%	42.3M	5s
2450K	...	...	...	...	...	2%	58.6M	5s
2500K	...	...	...	...	...	2%	82.9M	5s
2550K	...	...	...	...	...	2%	40.6M	5s
2600K	...	...	...	...	...	2%	76.2M	5s
2650K	...	...	...	...	...	2%	81.5M	4s
2700K	...	...	...	...	...	2%	89.5M	4s
2750K	...	...	...	...	...	2%	75.3M	4s
2800K	...	...	...	...	...	2%	91.0M	4s
2850K	...	...	...	...	...	3%	91.1M	4s
2900K	...	...	...	...	...	3%	90.0M	4s
2950K	...	...	...	...	...	3%	71.7M	4s
3000K	...	...	...	...	...	3%	79.8M	4s
3050K	...	...	...	...	...	3%	39.9M	4s
3100K	...	...	...	...	...	3%	81.2M	4s
3150K	...	...	...	...	...	3%	76.5M	4s
3200K	...	...	...	...	...	3%	83.4M	4s
3250K	...	...	...	...	...	3%	79.9M	4s
3300K	...	...	...	...	...	3%	76.9M	4s
3350K	...	...	...	...	...	3%	49.1M	4s
3400K	...	...	...	...	...	3%	90.3M	4s
3450K	...	...	...	...	...	3%	85.1M	4s
3500K	...	...	...	...	...	3%	50.6M	4s

3550K	...	...	...	...	3%	24.8M	4s
3600K	...	...	...	...	3%	73.5M	4s
3650K	...	...	...	...	3%	79.0M	4s
3700K	...	...	...	...	3%	80.3M	4s
3750K	...	...	...	...	3%	73.6M	4s
3800K	...	...	...	...	3%	94.0M	3s
3850K	...	...	...	...	4%	88.6M	3s
3900K	...	...	...	...	4%	87.8M	3s
3950K	...	...	...	...	4%	74.9M	3s
4000K	...	...	...	...	4%	43.1M	3s
4050K	...	...	...	...	4%	81.6M	3s
4100K	...	...	...	...	4%	46.0M	3s
4150K	...	...	...	...	4%	63.7M	3s
4200K	...	...	...	...	4%	79.2M	3s
4250K	...	...	...	...	4%	87.4M	3s
4300K	...	...	...	...	4%	93.4M	3s
4350K	...	...	...	...	4%	74.5M	3s
4400K	...	...	...	...	4%	68.3M	3s
4450K	...	...	...	...	4%	88.7M	3s
4500K	...	...	...	...	4%	88.8M	3s
4550K	...	...	...	...	4%	79.4M	3s
4600K	...	...	...	...	4%	41.8M	3s
4650K	...	...	...	...	4%	83.0M	3s
4700K	...	...	...	...	4%	86.2M	3s
4750K	...	...	...	...	4%	81.0M	3s
4800K	...	...	...	...	5%	75.3M	3s
4850K	...	...	...	...	5%	77.8M	3s
4900K	...	...	...	...	5%	87.5M	3s
4950K	...	...	...	...	5%	76.8M	3s
5000K	...	...	...	...	5%	91.5M	3s
5050K	...	...	...	...	5%	88.4M	3s
5100K	...	...	...	...	5%	1.20M	4s
5150K	...	...	...	...	5%	71.2M	4s
5200K	...	...	...	...	5%	65.1M	4s
5250K	...	...	...	...	5%	86.2M	3s
5300K	...	...	...	...	5%	86.1M	3s
5350K	...	...	...	...	5%	64.7M	3s
5400K	...	...	...	...	5%	92.1M	3s
5450K	...	...	...	...	5%	85.6M	3s
5500K	...	...	...	...	5%	82.5M	3s
5550K	...	...	...	...	5%	49.6M	3s
5600K	...	...	...	...	5%	33.3M	3s
5650K	...	...	...	...	5%	50.5M	3s
5700K	...	...	...	...	5%	60.8M	3s
5750K	...	...	...	...	6%	79.9M	3s
5800K	...	...	...	...	6%	84.6M	3s
5850K	...	...	...	...	6%	89.4M	3s
5900K	...	...	...	...	6%	85.9M	3s

5950K	...	...	...	...	6%	79.0M	3s
6000K	...	...	...	...	6%	89.1M	3s
6050K	...	...	...	...	6%	76.3M	3s
6100K	...	...	...	...	6%	66.6M	3s
6150K	...	...	...	...	6%	32.9M	3s
6200K	...	...	...	...	6%	76.9M	3s
6250K	...	...	...	...	6%	87.1M	3s
6300K	...	...	...	...	6%	69.5M	3s
6350K	...	...	...	...	6%	66.0M	3s
6400K	...	...	...	...	6%	79.8M	3s
6450K	...	...	...	...	6%	84.5M	3s
6500K	...	...	...	...	6%	92.1M	3s
6550K	...	...	...	...	6%	76.9M	3s
6600K	...	...	...	...	6%	61.0M	3s
6650K	...	...	...	...	6%	39.0M	3s
6700K	...	...	...	...	6%	48.5M	3s
6750K	...	...	...	...	7%	48.0M	3s
6800K	...	...	...	...	7%	87.0M	3s
6850K	...	...	...	...	7%	90.3M	3s
6900K	...	...	...	...	7%	84.7M	3s
6950K	...	...	...	...	7%	78.7M	3s
7000K	...	...	...	...	7%	85.7M	3s
7050K	...	...	...	...	7%	91.8M	3s
7100K	...	...	...	...	7%	86.6M	3s
7150K	...	...	...	...	7%	35.7M	3s
7200K	...	...	...	...	7%	67.8M	3s
7250K	...	...	...	...	7%	62.7M	3s
7300K	...	...	...	...	7%	80.4M	3s
7350K	...	...	...	...	7%	75.9M	3s
7400K	...	...	...	...	7%	73.2M	3s
7450K	...	...	...	...	7%	91.0M	3s
7500K	...	...	...	...	7%	66.0M	3s
7550K	...	...	...	...	7%	76.9M	3s
7600K	...	...	...	...	7%	93.1M	3s
7650K	...	...	...	...	7%	41.3M	3s
7700K	...	...	...	...	8%	76.0M	3s
7750K	...	...	...	...	8%	69.3M	3s
7800K	...	...	...	...	8%	95.6M	3s
7850K	...	...	...	...	8%	84.4M	3s
7900K	...	...	...	...	8%	70.3M	3s
7950K	...	...	...	...	8%	54.9M	3s
8000K	...	...	...	...	8%	57.9M	3s
8050K	...	...	...	...	8%	65.1M	3s
8100K	...	...	...	...	8%	27.1M	3s
8150K	...	...	...	...	8%	58.8M	3s
8200K	...	...	...	...	8%	32.0M	3s
8250K	...	...	...	...	8%	83.9M	3s
8300K	...	...	...	...	8%	93.8M	3s

8350K	...	...	...	...	...	8%	71.7M	3s
8400K	...	...	...	...	...	8%	89.3M	3s
8450K	...	...	...	...	...	8%	87.4M	3s
8500K	...	...	...	...	...	8%	92.9M	3s
8550K	...	...	...	...	...	8%	77.6M	3s
8600K	...	...	...	...	...	8%	87.5M	3s
8650K	...	...	...	...	...	9%	94.2M	3s
8700K	...	...	...	...	...	9%	41.8M	3s
8750K	...	...	...	...	...	9%	76.7M	3s
8800K	...	...	...	...	...	9%	88.2M	3s
8850K	...	...	...	...	...	9%	89.7M	3s
8900K	...	...	...	...	...	9%	89.1M	3s
8950K	...	...	...	...	...	9%	80.9M	3s
9000K	...	...	...	...	...	9%	86.5M	2s
9050K	...	...	...	...	...	9%	89.2M	2s
9100K	...	...	...	...	...	9%	86.0M	2s
9150K	...	...	...	...	...	9%	77.1M	2s
9200K	...	...	...	...	...	9%	43.6M	2s
9250K	...	...	...	...	...	9%	79.1M	2s
9300K	...	...	...	...	...	9%	81.9M	2s
9350K	...	...	...	...	...	9%	75.5M	2s
9400K	...	...	...	...	...	9%	77.9M	2s
9450K	...	...	...	...	...	9%	79.9M	2s
9500K	...	...	...	...	...	9%	89.8M	2s
9550K	...	...	...	...	...	9%	76.4M	2s
9600K	...	...	...	...	...	10%	89.6M	2s
9650K	...	...	...	...	...	10%	88.4M	2s
9700K	...	...	...	...	...	10%	44.9M	2s
9750K	...	...	...	...	...	10%	69.2M	2s
9800K	...	...	...	...	...	10%	92.7M	2s
9850K	...	...	...	...	...	10%	81.6M	2s
9900K	...	...	...	...	...	10%	78.7M	2s
9950K	...	...	...	...	...	10%	70.6M	2s
10000K	...	...	...	...	...	10%	90.6M	2s
10050K	...	...	...	...	...	10%	89.4M	2s
10100K	...	...	...	...	...	10%	88.3M	2s
10150K	...	...	...	...	...	10%	77.3M	2s
10200K	...	...	...	...	...	10%	92.1M	2s
10250K	...	...	...	...	...	10%	42.8M	2s
10300K	...	...	...	...	...	10%	81.1M	2s
10350K	...	...	...	...	...	10%	72.7M	2s
10400K	...	...	...	...	...	10%	62.2M	2s
10450K	...	...	...	...	...	10%	86.0M	2s
10500K	...	...	...	...	...	10%	93.0M	2s
10550K	...	...	...	...	...	10%	76.3M	2s
10600K	...	...	...	...	...	11%	86.9M	2s
10650K	...	...	...	...	...	11%	93.9M	2s
10700K	...	...	...	...	...	11%	90.3M	2s

10750K	...	...	...	...	11%	41.5M	2s
10800K	...	...	...	...	11%	76.7M	2s
10850K	...	...	...	...	11%	88.6M	2s
10900K	...	...	...	...	11%	89.5M	2s
10950K	...	...	...	...	11%	77.8M	2s
11000K	...	...	...	...	11%	87.6M	2s
11050K	...	...	...	...	11%	92.2M	2s
11100K	...	...	...	...	11%	86.5M	2s
11150K	...	...	...	...	11%	76.4M	2s
11200K	...	...	...	...	11%	87.3M	2s
11250K	...	...	...	...	11%	40.7M	2s
11300K	...	...	...	...	11%	73.6M	2s
11350K	...	...	...	...	11%	71.9M	2s
11400K	...	...	...	...	11%	79.7M	2s
11450K	...	...	...	...	11%	84.0M	2s
11500K	...	...	...	...	11%	92.4M	2s
11550K	...	...	...	...	12%	76.5M	2s
11600K	...	...	...	...	12%	90.9M	2s
11650K	...	...	...	...	12%	87.8M	2s
11700K	...	...	...	...	12%	88.6M	2s
11750K	...	...	...	...	12%	40.4M	2s
11800K	...	...	...	...	12%	94.0M	2s
11850K	...	...	...	...	12%	87.2M	2s
11900K	...	...	...	...	12%	90.2M	2s
11950K	...	...	...	...	12%	79.6M	2s
12000K	...	...	...	...	12%	89.6M	2s
12050K	...	...	...	...	12%	92.7M	2s
12100K	...	...	...	...	12%	89.2M	2s
12150K	...	...	...	...	12%	74.4M	2s
12200K	...	...	...	...	12%	83.5M	2s
12250K	...	...	...	...	12%	2.24M	2s
12300K	...	...	...	...	12%	30.1M	2s
12350K	...	...	...	...	12%	57.5M	2s
12400K	...	...	...	...	12%	86.6M	2s
12450K	...	...	...	...	12%	90.9M	2s
12500K	...	...	...	...	13%	86.6M	2s
12550K	...	...	...	...	13%	77.4M	2s
12600K	...	...	...	...	13%	92.8M	2s
12650K	...	...	...	...	13%	92.8M	2s
12700K	...	...	...	...	13%	92.7M	2s
12750K	...	...	...	...	13%	71.1M	2s
12800K	...	...	...	...	13%	36.8M	2s
12850K	...	...	...	...	13%	85.6M	2s
12900K	...	...	...	...	13%	94.5M	2s
12950K	...	...	...	...	13%	76.7M	2s
13000K	...	...	...	...	13%	91.9M	2s
13050K	...	...	...	...	13%	86.7M	2s
13100K	...	...	...	...	13%	89.8M	2s

13150K	...	...	...	...	...	13%	79.1M	2s
13200K	...	...	...	...	...	13%	89.8M	2s
13250K	...	...	...	...	...	13%	83.1M	2s
13300K	...	...	...	...	...	13%	37.7M	2s
13350K	...	...	...	...	...	13%	39.5M	2s
13400K	...	...	...	...	...	13%	57.2M	2s
13450K	...	...	...	...	...	13%	55.0M	2s
13500K	...	...	...	...	...	14%	68.1M	2s
13550K	...	...	...	...	...	14%	73.2M	2s
13600K	...	...	...	...	...	14%	93.4M	2s
13650K	...	...	...	...	...	14%	90.2M	2s
13700K	...	...	...	...	...	14%	91.6M	2s
13750K	...	...	...	...	...	14%	76.2M	2s
13800K	...	...	...	...	...	14%	43.8M	2s
13850K	...	...	...	...	...	14%	86.7M	2s
13900K	...	...	...	...	...	14%	93.9M	2s
13950K	...	...	...	...	...	14%	75.5M	2s
14000K	...	...	...	...	...	14%	92.9M	2s
14050K	...	...	...	...	...	14%	75.9M	2s
14100K	...	...	...	...	...	14%	77.2M	2s
14150K	...	...	...	...	...	14%	79.1M	2s
14200K	...	...	...	...	...	14%	88.3M	2s
14250K	...	...	...	...	...	14%	94.9M	2s
14300K	...	...	...	...	...	14%	51.0M	2s
14350K	...	...	...	...	...	14%	29.7M	2s
14400K	...	...	...	...	...	14%	22.2M	2s
14450K	...	...	...	...	...	15%	41.6M	2s
14500K	...	...	...	...	...	15%	72.7M	2s
14550K	...	...	...	...	...	15%	65.5M	2s
14600K	...	...	...	...	...	15%	71.3M	2s
14650K	...	...	...	...	...	15%	63.7M	2s
14700K	...	...	...	...	...	15%	69.3M	2s
14750K	...	...	...	...	...	15%	60.4M	2s
14800K	...	...	...	...	...	15%	81.4M	2s
14850K	...	...	...	...	...	15%	36.9M	2s
14900K	...	...	...	...	...	15%	69.7M	2s
14950K	...	...	...	...	...	15%	62.0M	2s
15000K	...	...	...	...	...	15%	71.8M	2s
15050K	...	...	...	...	...	15%	73.1M	2s
15100K	...	...	...	...	...	15%	70.2M	2s
15150K	...	...	...	...	...	15%	62.2M	2s
15200K	...	...	...	...	...	15%	58.3M	2s
15250K	...	...	...	...	...	15%	49.6M	2s
15300K	...	...	...	...	...	15%	45.1M	2s
15350K	...	...	...	...	...	15%	24.7M	2s
15400K	...	...	...	...	...	16%	68.4M	2s
15450K	...	...	...	...	...	16%	86.0M	2s
15500K	...	...	...	...	...	16%	84.2M	2s

15550K	...	...	...	...	16%	77.4M	2s
15600K	...	...	...	...	16%	81.6M	2s
15650K	...	...	...	...	16%	88.5M	2s
15700K	...	...	...	...	16%	80.1M	2s
15750K	...	...	...	...	16%	62.9M	2s
15800K	...	...	...	...	16%	75.1M	2s
15850K	...	...	...	...	16%	26.7M	2s
15900K	...	...	...	...	16%	71.9M	2s
15950K	...	...	...	...	16%	36.9M	2s
16000K	...	...	...	...	16%	66.7M	2s
16050K	...	...	...	...	16%	61.1M	2s
16100K	...	...	...	...	16%	53.6M	2s
16150K	...	...	...	...	16%	41.9M	2s
16200K	...	...	...	...	16%	38.9M	2s
16250K	...	...	...	...	16%	43.9M	2s
16300K	...	...	...	...	16%	14.9M	2s
16350K	...	...	...	...	16%	52.2M	2s
16400K	...	...	...	...	17%	29.6M	2s
16450K	...	...	...	...	17%	76.5M	2s
16500K	...	...	...	...	17%	92.1M	2s
16550K	...	...	...	...	17%	72.1M	2s
16600K	...	...	...	...	17%	82.8M	2s
16650K	...	...	...	...	17%	90.0M	2s
16700K	...	...	...	...	17%	85.2M	2s
16750K	...	...	...	...	17%	78.7M	2s
16800K	...	...	...	...	17%	57.0M	2s
16850K	...	...	...	...	17%	60.4M	2s
16900K	...	...	...	...	17%	12.3M	2s
16950K	...	...	...	...	17%	66.4M	2s
17000K	...	...	...	...	17%	80.1M	2s
17050K	...	...	...	...	17%	76.8M	2s
17100K	...	...	...	...	17%	72.7M	2s
17150K	...	...	...	...	17%	65.0M	2s
17200K	...	...	...	...	17%	84.4M	2s
17250K	...	...	...	...	17%	95.4M	2s
17300K	...	...	...	...	17%	86.5M	2s
17350K	...	...	...	...	18%	77.8M	2s
17400K	...	...	...	...	18%	22.2M	2s
17450K	...	...	...	...	18%	42.8M	2s
17500K	...	...	...	...	18%	82.3M	2s
17550K	...	...	...	...	18%	62.3M	2s
17600K	...	...	...	...	18%	89.1M	2s
17650K	...	...	...	...	18%	106M	2s
17700K	...	...	...	...	18%	93.0M	2s
17750K	...	...	...	...	18%	79.3M	2s
17800K	...	...	...	...	18%	70.5M	2s
17850K	...	...	...	...	18%	55.5M	2s
17900K	...	...	...	...	18%	46.0M	2s

17950K	...	...	...	...	18%	84.7M	2s
18000K	...	...	...	...	18%	80.1M	2s
18050K	...	...	...	...	18%	39.1M	2s
18100K	...	...	...	...	18%	74.7M	2s
18150K	...	...	...	...	18%	59.5M	2s
18200K	...	...	...	...	18%	73.5M	2s
18250K	...	...	...	...	18%	74.7M	2s
18300K	...	...	...	...	19%	74.5M	2s
18350K	...	...	...	...	19%	61.7M	2s
18400K	...	...	...	...	19%	71.7M	2s
18450K	...	...	...	...	19%	36.8M	2s
18500K	...	...	...	...	19%	71.4M	2s
18550K	...	...	...	...	19%	35.3M	2s
18600K	...	...	...	...	19%	63.0M	2s
18650K	...	...	...	...	19%	74.4M	2s
18700K	...	...	...	...	19%	72.6M	2s
18750K	...	...	...	...	19%	45.1M	2s
18800K	...	...	...	...	19%	68.7M	2s
18850K	...	...	...	...	19%	54.2M	2s
18900K	...	...	...	...	19%	71.3M	2s
18950K	...	...	...	...	19%	33.7M	2s
19000K	...	...	...	...	19%	85.9M	2s
19050K	...	...	...	...	19%	84.0M	2s
19100K	...	...	...	...	19%	83.7M	2s
19150K	...	...	...	...	19%	76.9M	2s
19200K	...	...	...	...	19%	81.9M	2s
19250K	...	...	...	...	20%	86.1M	2s
19300K	...	...	...	...	20%	82.2M	2s
19350K	...	...	...	...	20%	72.3M	2s
19400K	...	...	...	...	20%	88.3M	2s
19450K	...	...	...	...	20%	39.6M	2s
19500K	...	...	...	...	20%	82.8M	2s
19550K	...	...	...	...	20%	42.1M	2s
19600K	...	...	...	...	20%	52.6M	2s
19650K	...	...	...	...	20%	72.8M	2s
19700K	...	...	...	...	20%	84.2M	2s
19750K	...	...	...	...	20%	63.9M	2s
19800K	...	...	...	...	20%	82.5M	2s
19850K	...	...	...	...	20%	88.1M	2s
19900K	...	...	...	...	20%	57.6M	2s
19950K	...	...	...	...	20%	33.5M	2s
20000K	...	...	...	...	20%	88.4M	2s
20050K	...	...	...	...	20%	84.7M	2s
20100K	...	...	...	...	20%	92.6M	2s
20150K	...	...	...	...	20%	71.1M	2s
20200K	...	...	...	...	20%	89.8M	2s
20250K	...	...	...	...	21%	83.9M	2s
20300K	...	...	...	...	21%	82.7M	2s

20350K	...	...	...	...	...	21%	75.8M	2s
20400K	...	...	...	...	...	21%	84.7M	2s
20450K	...	...	...	...	...	21%	88.7M	2s
20500K	...	...	...	...	...	21%	41.0M	2s
20550K	...	...	...	...	...	21%	73.9M	2s
20600K	...	...	...	...	...	21%	57.3M	2s
20650K	...	...	...	...	...	21%	57.0M	2s
20700K	...	...	...	...	...	21%	81.7M	2s
20750K	...	...	...	...	...	21%	76.5M	2s
20800K	...	...	...	...	...	21%	80.7M	2s
20850K	...	...	...	...	...	21%	73.8M	2s
20900K	...	...	...	...	...	21%	83.5M	2s
20950K	...	...	...	...	...	21%	69.9M	2s
21000K	...	...	...	...	...	21%	36.0M	2s
21050K	...	...	...	...	...	21%	82.0M	2s
21100K	...	...	...	...	...	21%	90.6M	2s
21150K	...	...	...	...	...	21%	72.9M	2s
21200K	...	...	...	...	...	22%	87.8M	2s
21250K	...	...	...	...	...	22%	83.8M	2s
21300K	...	...	...	...	...	22%	83.5M	2s
21350K	...	...	...	...	...	22%	76.5M	2s
21400K	...	...	...	...	...	22%	81.0M	2s
21450K	...	...	...	...	...	22%	89.2M	2s
21500K	...	...	...	...	...	22%	42.3M	2s
21550K	...	...	...	...	...	22%	70.9M	2s
21600K	...	...	...	...	...	22%	83.3M	2s
21650K	...	...	...	...	...	22%	87.4M	2s
21700K	...	...	...	...	...	22%	57.8M	2s
21750K	...	...	...	...	...	22%	45.5M	2s
21800K	...	...	...	...	...	22%	82.6M	2s
21850K	...	...	...	...	...	22%	89.8M	2s
21900K	...	...	...	...	...	22%	79.7M	2s
21950K	...	...	...	...	...	22%	57.0M	2s
22000K	...	...	...	...	...	22%	37.9M	2s
22050K	...	...	...	...	...	22%	77.5M	2s
22100K	...	...	...	...	...	22%	70.6M	2s
22150K	...	...	...	...	...	23%	77.3M	2s
22200K	...	...	...	...	...	23%	86.6M	2s
22250K	...	...	...	...	...	23%	84.7M	2s
22300K	...	...	...	...	...	23%	79.0M	2s
22350K	...	...	...	...	...	23%	68.8M	2s
22400K	...	...	...	...	...	23%	88.8M	2s
22450K	...	...	...	...	...	23%	87.0M	2s
22500K	...	...	...	...	...	23%	88.6M	2s
22550K	...	...	...	...	...	23%	37.3M	2s
22600K	...	...	...	...	...	23%	82.1M	2s
22650K	...	...	...	...	...	23%	83.2M	2s
22700K	...	...	...	...	...	23%	85.3M	2s

22750K	...	...	...	...	23%	57.8M	2s
22800K	...	...	...	...	23%	60.9M	2s
22850K	...	...	...	...	23%	82.5M	2s
22900K	...	...	...	...	23%	89.0M	2s
22950K	...	...	...	...	23%	71.7M	2s
23000K	...	...	...	...	23%	81.9M	2s
23050K	...	...	...	...	23%	40.1M	2s
23100K	...	...	...	...	23%	90.4M	2s
23150K	...	...	...	...	24%	71.1M	2s
23200K	...	...	...	...	24%	85.6M	2s
23250K	...	...	...	...	24%	84.0M	2s
23300K	...	...	...	...	24%	85.7M	2s
23350K	...	...	...	...	24%	76.9M	2s
23400K	...	...	...	...	24%	83.4M	2s
23450K	...	...	...	...	24%	88.1M	2s
23500K	...	...	...	...	24%	83.6M	2s
23550K	...	...	...	...	24%	40.0M	2s
23600K	...	...	...	...	24%	83.0M	2s
23650K	...	...	...	...	24%	92.1M	2s
23700K	...	...	...	...	24%	83.9M	2s
23750K	...	...	...	...	24%	74.8M	2s
23800K	...	...	...	...	24%	87.8M	2s
23850K	...	...	...	...	24%	84.4M	2s
23900K	...	...	...	...	24%	77.9M	2s
23950K	...	...	...	...	24%	57.1M	2s
24000K	...	...	...	...	24%	80.6M	2s
24050K	...	...	...	...	24%	42.6M	2s
24100K	...	...	...	...	25%	74.8M	2s
24150K	...	...	...	...	25%	70.8M	2s
24200K	...	...	...	...	25%	84.5M	2s
24250K	...	...	...	...	25%	84.3M	2s
24300K	...	...	...	...	25%	78.0M	2s
24350K	...	...	...	...	25%	66.6M	2s
24400K	...	...	...	...	25%	82.7M	2s
24450K	...	...	...	...	25%	84.0M	1s
24500K	...	...	...	...	25%	80.3M	1s
24550K	...	...	...	...	25%	40.5M	1s
24600K	...	...	...	...	25%	82.2M	1s
24650K	...	...	...	...	25%	81.0M	1s
24700K	...	...	...	...	25%	87.7M	1s
24750K	...	...	...	...	25%	76.5M	1s
24800K	...	...	...	...	25%	86.1M	1s
24850K	...	...	...	...	25%	82.1M	1s
24900K	...	...	...	...	25%	74.0M	1s
24950K	...	...	...	...	25%	63.0M	1s
25000K	...	...	...	...	25%	75.8M	1s
25050K	...	...	...	...	26%	87.0M	1s
25100K	...	...	...	...	26%	35.0M	1s

25150K	...	...	...	...	26%	58.7M	1s
25200K	...	...	...	...	26%	74.9M	1s
25250K	...	...	...	...	26%	86.6M	1s
25300K	...	...	...	...	26%	84.7M	1s
25350K	...	...	...	...	26%	72.9M	1s
25400K	...	...	...	...	26%	87.5M	1s
25450K	...	...	...	...	26%	76.5M	1s
25500K	...	...	...	...	26%	82.9M	1s
25550K	...	...	...	...	26%	68.3M	1s
25600K	...	...	...	...	26%	42.5M	1s
25650K	...	...	...	...	26%	70.7M	1s
25700K	...	...	...	...	26%	77.3M	1s
25750K	...	...	...	...	26%	77.1M	1s
25800K	...	...	...	...	26%	84.7M	1s
25850K	...	...	...	...	26%	88.6M	1s
25900K	...	...	...	...	26%	85.2M	1s
25950K	...	...	...	...	26%	71.8M	1s
26000K	...	...	...	...	26%	91.0M	1s
26050K	...	...	...	...	27%	76.9M	1s
26100K	...	...	...	...	27%	35.6M	1s
26150K	...	...	...	...	27%	60.4M	1s
26200K	...	...	...	...	27%	74.7M	1s
26250K	...	...	...	...	27%	85.2M	1s
26300K	...	...	...	...	27%	70.8M	1s
26350K	...	...	...	...	27%	59.4M	1s
26400K	...	...	...	...	27%	72.6M	1s
26450K	...	...	...	...	27%	78.8M	1s
26500K	...	...	...	...	27%	74.3M	1s
26550K	...	...	...	...	27%	70.5M	1s
26600K	...	...	...	...	27%	35.0M	1s
26650K	...	...	...	...	27%	66.5M	1s
26700K	...	...	...	...	27%	432K	2s
26750K	...	...	...	...	27%	73.8M	2s
26800K	...	...	...	...	27%	55.6M	2s
26850K	...	...	...	...	27%	45.6M	2s
26900K	...	...	...	...	27%	40.2M	2s
26950K	...	...	...	...	27%	50.7M	2s
27000K	...	...	...	...	28%	49.7M	2s
27050K	...	...	...	...	28%	69.0M	2s
27100K	...	...	...	...	28%	63.9M	2s
27150K	...	...	...	...	28%	35.0M	2s
27200K	...	...	...	...	28%	80.0M	2s
27250K	...	...	...	...	28%	75.3M	2s
27300K	...	...	...	...	28%	64.0M	2s
27350K	...	...	...	...	28%	77.5M	2s
27400K	...	...	...	...	28%	75.7M	2s
27450K	...	...	...	...	28%	72.7M	2s
27500K	...	...	...	...	28%	73.7M	2s

27550K	...	...	...	...	28%	70.9M	2s
27600K	...	...	...	...	28%	52.9M	2s
27650K	...	...	...	...	28%	35.5M	2s
27700K	...	...	...	...	28%	50.9M	2s
27750K	...	...	...	...	28%	77.1M	2s
27800K	...	...	...	...	28%	71.4M	2s
27850K	...	...	...	...	28%	75.5M	2s
27900K	...	...	...	...	28%	63.7M	2s
27950K	...	...	...	...	29%	78.5M	2s
28000K	...	...	...	...	29%	78.5M	2s
28050K	...	...	...	...	29%	75.4M	2s
28100K	...	...	...	...	29%	63.4M	2s
28150K	...	...	...	...	29%	40.1M	2s
28200K	...	...	...	...	29%	70.4M	2s
28250K	...	...	...	...	29%	79.5M	2s
28300K	...	...	...	...	29%	64.4M	2s
28350K	...	...	...	...	29%	76.0M	2s
28400K	...	...	...	...	29%	77.2M	2s
28450K	...	...	...	...	29%	77.8M	2s
28500K	...	...	...	...	29%	70.3M	2s
28550K	...	...	...	...	29%	86.7M	2s
28600K	...	...	...	...	29%	82.6M	2s
28650K	...	...	...	...	29%	40.8M	2s
28700K	...	...	...	...	29%	63.8M	2s
28750K	...	...	...	...	29%	136K	2s
28800K	...	...	...	...	29%	90.0M	2s
28850K	...	...	...	...	29%	129M	2s
28900K	...	...	...	...	30%	123M	2s
28950K	...	...	...	...	30%	118M	2s
29000K	...	...	...	...	30%	78.0M	2s
29050K	...	...	...	...	30%	93.4M	2s
29100K	...	...	...	...	30%	129M	2s
29150K	...	...	...	...	30%	105M	2s
29200K	...	...	...	...	30%	38.3M	2s
29250K	...	...	...	...	30%	119M	2s
29300K	...	...	...	...	30%	116M	2s
29350K	...	...	...	...	30%	128M	2s
29400K	...	...	...	...	30%	71.9M	2s
29450K	...	...	...	...	30%	123M	2s
29500K	...	...	...	...	30%	113M	2s
29550K	...	...	...	...	30%	125M	2s
29600K	...	...	...	...	30%	90.6M	2s
29650K	...	...	...	...	30%	62.5M	2s
29700K	...	...	...	...	30%	43.2M	2s
29750K	...	...	...	...	30%	102M	2s
29800K	...	...	...	...	30%	92.4M	2s
29850K	...	...	...	...	30%	115M	2s
29900K	...	...	...	...	31%	128M	2s

29950K	...	...	...	...	31%	121M	2s
30000K	...	...	...	...	31%	108M	2s
30050K	...	...	...	...	31%	117M	2s
30100K	...	...	...	...	31%	126M	2s
30150K	...	...	...	...	31%	73.7M	2s
30200K	...	...	...	...	31%	42.3M	2s
30250K	...	...	...	...	31%	121M	2s
30300K	...	...	...	...	31%	116M	2s
30350K	...	...	...	...	31%	121M	2s
30400K	...	...	...	...	31%	103M	2s
30450K	...	...	...	...	31%	124M	2s
30500K	...	...	...	...	31%	120M	2s
30550K	...	...	...	...	31%	120M	2s
30600K	...	...	...	...	31%	109M	2s
30650K	...	...	...	...	31%	120M	2s
30700K	...	...	...	...	31%	53.4M	2s
30750K	...	...	...	...	31%	222K	3s
30800K	...	...	...	...	31%	324K	3s
30850K	...	...	...	...	32%	58.6M	3s
30900K	...	...	...	...	32%	66.5M	3s
30950K	...	...	...	...	32%	121M	3s
31000K	...	...	...	...	32%	95.1M	3s
31050K	...	...	...	...	32%	129M	3s
31100K	...	...	...	...	32%	114M	3s
31150K	...	...	...	...	32%	111M	3s
31200K	...	...	...	...	32%	45.0M	3s
31250K	...	...	...	...	32%	92.3M	3s
31300K	...	...	...	...	32%	111M	3s
31350K	...	...	...	...	32%	117M	3s
31400K	...	...	...	...	32%	106M	3s
31450K	...	...	...	...	32%	126M	3s
31500K	...	...	...	...	32%	115M	3s
31550K	...	...	...	...	32%	122M	3s
31600K	...	...	...	...	32%	100M	3s
31650K	...	...	...	...	32%	126M	3s
31700K	...	...	...	...	32%	117M	3s
31750K	...	...	...	...	32%	52.0M	3s
31800K	...	...	...	...	33%	100M	3s
31850K	...	...	...	...	33%	116M	3s
31900K	...	...	...	...	33%	105M	3s
31950K	...	...	...	...	33%	114M	3s
32000K	...	...	...	...	33%	101M	3s
32050K	...	...	...	...	33%	126M	3s
32100K	...	...	...	...	33%	14.5M	3s
32150K	...	...	...	...	33%	56.7M	3s
32200K	...	...	...	...	33%	62.1M	3s
32250K	...	...	...	...	33%	44.1M	3s
32300K	...	...	...	...	33%	74.4M	3s

32350K	...	...	...	...	33%	82.9M	3s
32400K	...	...	...	...	33%	79.3M	3s
32450K	...	...	...	...	33%	90.4M	3s
32500K	...	...	...	...	33%	90.2M	3s
32550K	...	...	...	...	33%	87.1M	3s
32600K	...	...	...	...	33%	79.2M	3s
32650K	...	...	...	...	33%	82.8M	3s
32700K	...	...	...	...	33%	83.0M	3s
32750K	...	...	...	...	33%	41.1M	3s
32800K	...	...	...	...	34%	134K	4s
32850K	...	...	...	...	34%	20.4M	4s
32900K	...	...	...	...	34%	94.6M	4s
32950K	...	...	...	...	34%	77.4M	4s
33000K	...	...	...	...	34%	65.3M	4s
33050K	...	...	...	...	34%	86.1M	4s
33100K	...	...	...	...	34%	84.3M	4s
33150K	...	...	...	...	34%	85.5M	4s
33200K	...	...	...	...	34%	75.4M	4s
33250K	...	...	...	...	34%	80.7M	4s
33300K	...	...	...	...	34%	42.7M	4s
33350K	...	...	...	...	34%	83.8M	4s
33400K	...	...	...	...	34%	75.1M	4s
33450K	...	...	...	...	34%	80.0M	4s
33500K	...	...	...	...	34%	81.6M	4s
33550K	...	...	...	...	34%	79.8M	3s
33600K	...	...	...	...	34%	67.8M	3s
33650K	...	...	...	...	34%	75.0M	3s
33700K	...	...	...	...	34%	83.0M	3s
33750K	...	...	...	...	35%	87.0M	3s
33800K	...	...	...	...	35%	19.5M	3s
33850K	...	...	...	...	35%	66.5M	3s
33900K	...	...	...	...	35%	83.6M	3s
33950K	...	...	...	...	35%	91.2M	3s
34000K	...	...	...	...	35%	81.5M	3s
34050K	...	...	...	...	35%	90.1M	3s
34100K	...	...	...	...	35%	94.8M	3s
34150K	...	...	...	...	35%	97.7M	3s
34200K	...	...	...	...	35%	85.1M	3s
34250K	...	...	...	...	35%	99.0M	3s
34300K	...	...	...	...	35%	48.9M	3s
34350K	...	...	...	...	35%	90.3M	3s
34400K	...	...	...	...	35%	87.3M	3s
34450K	...	...	...	...	35%	80.2M	3s
34500K	...	...	...	...	35%	88.0M	3s
34550K	...	...	...	...	35%	88.6M	3s
34600K	...	...	...	...	35%	73.5M	3s
34650K	...	...	...	...	35%	79.0M	3s
34700K	...	...	...	...	36%	84.8M	3s

34750K	...	...	...	...	36%	94.5M	3s
34800K	...	...	...	...	36%	42.8M	3s
34850K	...	...	...	...	36%	134K	4s
34900K	...	...	...	...	36%	36.9M	4s
34950K	...	...	...	...	36%	66.2M	4s
35000K	...	...	...	...	36%	76.2M	4s
35050K	...	...	...	...	36%	87.8M	4s
35100K	...	...	...	...	36%	97.3M	4s
35150K	...	...	...	...	36%	93.4M	4s
35200K	...	...	...	...	36%	84.5M	4s
35250K	...	...	...	...	36%	80.7M	4s
35300K	...	...	...	...	36%	92.2M	4s
35350K	...	...	...	...	36%	41.1M	4s
35400K	...	...	...	...	36%	84.5M	4s
35450K	...	...	...	...	36%	93.0M	4s
35500K	...	...	...	...	36%	94.1M	4s
35550K	...	...	...	...	36%	82.0M	4s
35600K	...	...	...	...	36%	77.3M	4s
35650K	...	...	...	...	36%	94.6M	4s
35700K	...	...	...	...	37%	94.7M	4s
35750K	...	...	...	...	37%	86.2M	4s
35800K	...	...	...	...	37%	72.8M	4s
35850K	...	...	...	...	37%	43.5M	4s
35900K	...	...	...	...	37%	10.1M	4s
35950K	...	...	...	...	37%	92.5M	4s
36000K	...	...	...	...	37%	82.9M	4s
36050K	...	...	...	...	37%	99.3M	4s
36100K	...	...	...	...	37%	93.9M	4s
36150K	...	...	...	...	37%	93.6M	4s
36200K	...	...	...	...	37%	79.6M	4s
36250K	...	...	...	...	37%	79.0M	4s
36300K	...	...	...	...	37%	88.0M	4s
36350K	...	...	...	...	37%	46.8M	4s
36400K	...	...	...	...	37%	79.6M	4s
36450K	...	...	...	...	37%	98.3M	4s
36500K	...	...	...	...	37%	82.1M	4s
36550K	...	...	...	...	37%	90.5M	4s
36600K	...	...	...	...	37%	71.3M	4s
36650K	...	...	...	...	38%	86.3M	4s
36700K	...	...	...	...	38%	95.1M	4s
36750K	...	...	...	...	38%	94.1M	4s
36800K	...	...	...	...	38%	80.3M	4s
36850K	...	...	...	...	38%	48.3M	4s
36900K	...	...	...	...	38%	135K	4s
36950K	...	...	...	...	38%	33.7M	4s
37000K	...	...	...	...	38%	58.8M	4s
37050K	...	...	...	...	38%	80.2M	4s
37100K	...	...	...	...	38%	88.6M	4s

37150K	...	...	...	...	38%	89.8M	4s
37200K	...	...	...	...	38%	81.5M	4s
37250K	...	...	...	...	38%	92.9M	4s
37300K	...	...	...	...	38%	97.7M	4s
37350K	...	...	...	...	38%	92.7M	4s
37400K	...	...	...	...	38%	43.2M	4s
37450K	...	...	...	...	38%	112M	4s
37500K	...	...	...	...	38%	98.7M	4s
37550K	...	...	...	...	38%	108M	4s
37600K	...	...	...	...	39%	85.7M	4s
37650K	...	...	...	...	39%	93.8M	4s
37700K	...	...	...	...	39%	101M	4s
37750K	...	...	...	...	39%	113M	4s
37800K	...	...	...	...	39%	82.7M	4s
37850K	...	...	...	...	39%	97.7M	4s
37900K	...	...	...	...	39%	48.5M	4s
37950K	...	...	...	...	39%	103M	4s
38000K	...	...	...	...	39%	90.6M	4s
38050K	...	...	...	...	39%	105M	4s
38100K	...	...	...	...	39%	95.8M	4s
38150K	...	...	...	...	39%	103M	4s
38200K	...	...	...	...	39%	43.2M	4s
38250K	...	...	...	...	39%	64.5M	4s
38300K	...	...	...	...	39%	103M	4s
38350K	...	...	...	...	39%	90.2M	4s
38400K	...	...	...	...	39%	33.9M	4s
38450K	...	...	...	...	39%	92.3M	4s
38500K	...	...	...	...	39%	101M	4s
38550K	...	...	...	...	40%	102M	4s
38600K	...	...	...	...	40%	89.7M	4s
38650K	...	...	...	...	40%	90.8M	4s
38700K	...	...	...	...	40%	105M	4s
38750K	...	...	...	...	40%	86.9M	4s
38800K	...	...	...	...	40%	71.7M	4s
38850K	...	...	...	...	40%	102M	4s
38900K	...	...	...	...	40%	46.7M	4s
38950K	...	...	...	...	40%	133K	5s
39000K	...	...	...	...	40%	36.2M	5s
39050K	...	...	...	...	40%	68.2M	4s
39100K	...	...	...	...	40%	83.4M	4s
39150K	...	...	...	...	40%	90.2M	4s
39200K	...	...	...	...	40%	74.2M	4s
39250K	...	...	...	...	40%	92.2M	4s
39300K	...	...	...	...	40%	84.1M	4s
39350K	...	...	...	...	40%	89.0M	4s
39400K	...	...	...	...	40%	42.4M	4s
39450K	...	...	...	...	40%	81.0M	4s
39500K	...	...	...	...	40%	82.9M	4s

39550K	...	...	...	...	41%	89.6M	4s
39600K	...	...	...	...	41%	81.2M	4s
39650K	...	...	...	...	41%	91.3M	4s
39700K	...	...	...	...	41%	91.5M	4s
39750K	...	...	...	...	41%	86.1M	4s
39800K	...	...	...	...	41%	80.3M	4s
39850K	...	...	...	...	41%	89.8M	4s
39900K	...	...	...	...	41%	94.1M	4s
39950K	...	...	...	...	41%	47.5M	4s
40000K	...	...	...	...	41%	78.3M	4s
40050K	...	...	...	...	41%	80.1M	4s
40100K	...	...	...	...	41%	92.5M	4s
40150K	...	...	...	...	41%	53.3M	4s
40200K	...	...	...	...	41%	60.6M	4s
40250K	...	...	...	...	41%	89.1M	4s
40300K	...	...	...	...	41%	90.1M	4s
40350K	...	...	...	...	41%	89.9M	4s
40400K	...	...	...	...	41%	77.6M	4s
40450K	...	...	...	...	41%	46.3M	4s
40500K	...	...	...	...	42%	90.1M	4s
40550K	...	...	...	...	42%	97.2M	4s
40600K	...	...	...	...	42%	64.9M	4s
40650K	...	...	...	...	42%	88.1M	4s
40700K	...	...	...	...	42%	91.0M	4s
40750K	...	...	...	...	42%	85.8M	4s
40800K	...	...	...	...	42%	69.6M	4s
40850K	...	...	...	...	42%	79.5M	4s
40900K	...	...	...	...	42%	88.8M	4s
40950K	...	...	...	...	42%	45.4M	4s
41000K	...	...	...	...	42%	134K	5s
41050K	...	...	...	...	42%	51.6M	5s
41100K	...	...	...	...	42%	67.6M	5s
41150K	...	...	...	...	42%	82.5M	5s
41200K	...	...	...	...	42%	77.8M	5s
41250K	...	...	...	...	42%	79.5M	5s
41300K	...	...	...	...	42%	88.9M	5s
41350K	...	...	...	...	42%	90.7M	5s
41400K	...	...	...	...	42%	77.7M	5s
41450K	...	...	...	...	43%	44.6M	5s
41500K	...	...	...	...	43%	80.8M	5s
41550K	...	...	...	...	43%	92.2M	5s
41600K	...	...	...	...	43%	76.3M	5s
41650K	...	...	...	...	43%	91.5M	5s
41700K	...	...	...	...	43%	83.5M	5s
41750K	...	...	...	...	43%	93.7M	5s
41800K	...	...	...	...	43%	74.7M	5s
41850K	...	...	...	...	43%	88.4M	5s
41900K	...	...	...	...	43%	87.1M	5s

41950K	...	...	...	...	43%	85.9M	5s
42000K	...	...	...	...	43%	38.8M	4s
42050K	...	...	...	...	43%	87.2M	4s
42100K	...	...	...	...	43%	85.5M	4s
42150K	...	...	...	...	43%	59.5M	4s
42200K	...	...	...	...	43%	52.6M	4s
42250K	...	...	...	...	43%	82.3M	4s
42300K	...	...	...	...	43%	83.8M	4s
42350K	...	...	...	...	43%	82.2M	4s
42400K	...	...	...	...	43%	69.5M	4s
42450K	...	...	...	...	44%	86.1M	4s
42500K	...	...	...	...	44%	45.3M	4s
42550K	...	...	...	...	44%	89.8M	4s
42600K	...	...	...	...	44%	50.5M	4s
42650K	...	...	...	...	44%	74.3M	4s
42700K	...	...	...	...	44%	78.7M	4s
42750K	...	...	...	...	44%	74.6M	4s
42800K	...	...	...	...	44%	63.4M	4s
42850K	...	...	...	...	44%	74.3M	4s
42900K	...	...	...	...	44%	92.4M	4s
42950K	...	...	...	...	44%	85.0M	4s
43000K	...	...	...	...	44%	38.9M	4s
43050K	...	...	...	...	44%	134K	5s
43100K	...	...	...	...	44%	52.8M	5s
43150K	...	...	...	...	44%	67.1M	5s
43200K	...	...	...	...	44%	65.5M	5s
43250K	...	...	...	...	44%	78.8M	5s
43300K	...	...	...	...	44%	75.6M	5s
43350K	...	...	...	...	44%	80.6M	5s
43400K	...	...	...	...	45%	92.8M	5s
43450K	...	...	...	...	45%	98.8M	5s
43500K	...	...	...	...	45%	44.7M	5s
43550K	...	...	...	...	45%	84.2M	5s
43600K	...	...	...	...	45%	74.3M	5s
43650K	...	...	...	...	45%	89.4M	5s
43700K	...	...	...	...	45%	76.8M	5s
43750K	...	...	...	...	45%	78.8M	5s
43800K	...	...	...	...	45%	73.6M	5s
43850K	...	...	...	...	45%	84.6M	5s
43900K	...	...	...	...	45%	76.7M	5s
43950K	...	...	...	...	45%	87.9M	5s
44000K	...	...	...	...	45%	62.3M	5s
44050K	...	...	...	...	45%	44.4M	5s
44100K	...	...	...	...	45%	77.6M	5s
44150K	...	...	...	...	45%	72.2M	5s
44200K	...	...	...	...	45%	52.8M	5s
44250K	...	...	...	...	45%	70.9M	5s
44300K	...	...	...	...	45%	81.3M	5s

44350K	...	...	...	...	46%	87.5M	5s
44400K	...	...	...	...	46%	73.9M	5s
44450K	...	...	...	...	46%	90.5M	5s
44500K	...	...	...	...	46%	80.1M	5s
44550K	...	...	...	...	46%	49.1M	5s
44600K	...	...	...	...	46%	93.4M	5s
44650K	...	...	...	...	46%	91.1M	4s
44700K	...	...	...	...	46%	101M	4s
44750K	...	...	...	...	46%	92.4M	4s
44800K	...	...	...	...	46%	81.1M	4s
44850K	...	...	...	...	46%	97.1M	4s
44900K	...	...	...	...	46%	92.9M	4s
44950K	...	...	...	...	46%	89.6M	4s
45000K	...	...	...	...	46%	80.4M	4s
45050K	...	...	...	...	46%	41.9M	4s
45100K	...	...	...	...	46%	134K	5s
45150K	...	...	...	...	46%	47.9M	5s
45200K	...	...	...	...	46%	56.6M	5s
45250K	...	...	...	...	46%	90.6M	5s
45300K	...	...	...	...	46%	85.6M	5s
45350K	...	...	...	...	47%	55.7M	5s
45400K	...	...	...	...	47%	66.9M	5s
45450K	...	...	...	...	47%	88.8M	5s
45500K	...	...	...	...	47%	93.2M	5s
45550K	...	...	...	...	47%	45.6M	5s
45600K	...	...	...	...	47%	64.4M	5s
45650K	...	...	...	...	47%	91.2M	5s
45700K	...	...	...	...	47%	88.1M	5s
45750K	...	...	...	...	47%	91.6M	5s
45800K	...	...	...	...	47%	75.1M	5s
45850K	...	...	...	...	47%	95.6M	5s
45900K	...	...	...	...	47%	74.7M	5s
45950K	...	...	...	...	47%	83.7M	5s
46000K	...	...	...	...	47%	75.9M	5s
46050K	...	...	...	...	47%	94.3M	5s
46100K	...	...	...	...	47%	45.3M	5s
46150K	...	...	...	...	47%	91.0M	5s
46200K	...	...	...	...	47%	54.4M	5s
46250K	...	...	...	...	47%	62.3M	5s
46300K	...	...	...	...	48%	79.7M	5s
46350K	...	...	...	...	48%	92.9M	5s
46400K	...	...	...	...	48%	74.0M	5s
46450K	...	...	...	...	48%	86.4M	5s
46500K	...	...	...	...	48%	92.9M	5s
46550K	...	...	...	...	48%	89.5M	5s
46600K	...	...	...	...	48%	34.6M	5s
46650K	...	...	...	...	48%	105M	5s
46700K	...	...	...	...	48%	89.7M	5s

46750K	...	...	...	...	48%	89.3M	5s
46800K	...	...	...	...	48%	79.1M	5s
46850K	...	...	...	...	48%	93.0M	5s
46900K	...	...	...	...	48%	88.4M	5s
46950K	...	...	...	...	48%	85.5M	5s
47000K	...	...	...	...	48%	82.0M	5s
47050K	...	...	...	...	48%	88.9M	4s
47100K	...	...	...	...	48%	48.1M	4s
47150K	...	...	...	...	48%	134K	5s
47200K	...	...	...	...	48%	46.6M	5s
47250K	...	...	...	...	49%	91.9M	5s
47300K	...	...	...	...	49%	88.1M	5s
47350K	...	...	...	...	49%	79.2M	5s
47400K	...	...	...	...	49%	56.6M	5s
47450K	...	...	...	...	49%	88.1M	5s
47500K	...	...	...	...	49%	96.6M	5s
47550K	...	...	...	...	49%	90.9M	5s
47600K	...	...	...	...	49%	32.6M	5s
47650K	...	...	...	...	49%	59.5M	5s
47700K	...	...	...	...	49%	83.2M	5s
47750K	...	...	...	...	49%	92.5M	5s
47800K	...	...	...	...	49%	76.0M	5s
47850K	...	...	...	...	49%	83.4M	5s
47900K	...	...	...	...	49%	88.1M	5s
47950K	...	...	...	...	49%	94.8M	5s
48000K	...	...	...	...	49%	80.7M	5s
48050K	...	...	...	...	49%	91.3M	5s
48100K	...	...	...	...	49%	90.7M	5s
48150K	...	...	...	...	49%	53.8M	5s
48200K	...	...	...	...	50%	146M	5s
48250K	...	...	...	...	50%	143M	5s
48300K	...	...	...	...	50%	76.5M	5s
48350K	...	...	...	...	50%	67.8M	5s
48400K	...	...	...	...	50%	114M	5s
48450K	...	...	...	...	50%	126M	5s
48500K	...	...	...	...	50%	144M	5s
48550K	...	...	...	...	50%	147M	5s
48600K	...	...	...	...	50%	104M	5s
48650K	...	...	...	...	50%	33.5M	5s
48700K	...	...	...	...	50%	103M	5s
48750K	...	...	...	...	50%	121M	5s
48800K	...	...	...	...	50%	96.4M	5s
48850K	...	...	...	...	50%	111M	5s
48900K	...	...	...	...	50%	92.5M	5s
48950K	...	...	...	...	50%	101M	5s
49000K	...	...	...	...	50%	92.3M	5s
49050K	...	...	...	...	50%	113M	5s
49100K	...	...	...	...	50%	122M	5s

49150K	...	...	...	...	50%	48.4M	5s
49200K	...	...	...	...	51%	133K	5s
49250K	...	...	...	...	51%	78.2M	5s
49300K	...	...	...	...	51%	72.7M	5s
49350K	...	...	...	...	51%	57.1M	5s
49400K	...	...	...	...	51%	48.9M	5s
49450K	...	...	...	...	51%	84.2M	5s
49500K	...	...	...	...	51%	89.2M	5s
49550K	...	...	...	...	51%	49.0M	5s
49600K	...	...	...	...	51%	67.5M	5s
49650K	...	...	...	...	51%	42.5M	5s
49700K	...	...	...	...	51%	63.0M	5s
49750K	...	...	...	...	51%	92.8M	5s
49800K	...	...	...	...	51%	75.6M	5s
49850K	...	...	...	...	51%	91.8M	5s
49900K	...	...	...	...	51%	90.9M	5s
49950K	...	...	...	...	51%	89.1M	5s
50000K	...	...	...	...	51%	81.3M	5s
50050K	...	...	...	...	51%	74.3M	5s
50100K	...	...	...	...	51%	89.0M	5s
50150K	...	...	...	...	52%	88.5M	5s
50200K	...	...	...	...	52%	43.3M	5s
50250K	...	...	...	...	52%	55.2M	5s
50300K	...	...	...	...	52%	88.3M	5s
50350K	...	...	...	...	52%	67.7M	5s
50400K	...	...	...	...	52%	71.1M	5s
50450K	...	...	...	...	52%	63.0M	5s
50500K	...	...	...	...	52%	73.1M	5s
50550K	...	...	...	...	52%	91.1M	5s
50600K	...	...	...	...	52%	65.4M	5s
50650K	...	...	...	...	52%	104M	5s
50700K	...	...	...	...	52%	45.7M	5s
50750K	...	...	...	...	52%	89.1M	5s
50800K	...	...	...	...	52%	80.5M	5s
50850K	...	...	...	...	52%	93.5M	5s
50900K	...	...	...	...	52%	90.9M	5s
50950K	...	...	...	...	52%	85.5M	5s
51000K	...	...	...	...	52%	81.8M	5s
51050K	...	...	...	...	52%	85.7M	5s
51100K	...	...	...	...	53%	90.4M	5s
51150K	...	...	...	...	53%	82.2M	4s
51200K	...	...	...	...	53%	42.7M	4s
51250K	...	...	...	...	53%	134K	5s
51300K	...	...	...	...	53%	54.2M	5s
51350K	...	...	...	...	53%	68.6M	5s
51400K	...	...	...	...	53%	67.7M	5s
51450K	...	...	...	...	53%	90.8M	5s
51500K	...	...	...	...	53%	90.8M	5s

51550K	...	...	...	...	53%	87.0M	5s
51600K	...	...	...	...	53%	81.1M	5s
51650K	...	...	...	...	53%	90.3M	5s
51700K	...	...	...	...	53%	45.5M	5s
51750K	...	...	...	...	53%	85.9M	5s
51800K	...	...	...	...	53%	75.6M	5s
51850K	...	...	...	...	53%	86.6M	5s
51900K	...	...	...	...	53%	93.7M	5s
51950K	...	...	...	...	53%	85.4M	5s
52000K	...	...	...	...	53%	68.8M	5s
52050K	...	...	...	...	53%	87.8M	5s
52100K	...	...	...	...	54%	68.5M	5s
52150K	...	...	...	...	54%	91.2M	5s
52200K	...	...	...	...	54%	44.6M	5s
52250K	...	...	...	...	54%	74.0M	5s
52300K	...	...	...	...	54%	90.5M	5s
52350K	...	...	...	...	54%	93.2M	5s
52400K	...	...	...	...	54%	39.0M	5s
52450K	...	...	...	...	54%	70.8M	5s
52500K	...	...	...	...	54%	91.1M	5s
52550K	...	...	...	...	54%	89.1M	5s
52600K	...	...	...	...	54%	75.4M	5s
52650K	...	...	...	...	54%	95.7M	5s
52700K	...	...	...	...	54%	85.2M	5s
52750K	...	...	...	...	54%	42.3M	5s
52800K	...	...	...	...	54%	76.3M	5s
52850K	...	...	...	...	54%	85.7M	5s
52900K	...	...	...	...	54%	89.0M	5s
52950K	...	...	...	...	54%	90.1M	4s
53000K	...	...	...	...	54%	77.8M	4s
53050K	...	...	...	...	55%	95.0M	4s
53100K	...	...	...	...	55%	91.8M	4s
53150K	...	...	...	...	55%	87.4M	4s
53200K	...	...	...	...	55%	71.8M	4s
53250K	...	...	...	...	55%	45.5M	4s
53300K	...	...	...	...	55%	134K	5s
53350K	...	...	...	...	55%	43.9M	5s
53400K	...	...	...	...	55%	70.1M	5s
53450K	...	...	...	...	55%	85.5M	5s
53500K	...	...	...	...	55%	89.1M	5s
53550K	...	...	...	...	55%	73.8M	5s
53600K	...	...	...	...	55%	63.0M	5s
53650K	...	...	...	...	55%	88.3M	5s
53700K	...	...	...	...	55%	93.5M	5s
53750K	...	...	...	...	55%	44.8M	5s
53800K	...	...	...	...	55%	71.4M	5s
53850K	...	...	...	...	55%	89.6M	5s
53900K	...	...	...	...	55%	88.3M	5s

53950K	...	...	...	...	55%	96.4M	5s
54000K	...	...	...	...	56%	76.3M	5s
54050K	...	...	...	...	56%	86.7M	5s
54100K	...	...	...	...	56%	83.0M	5s
54150K	...	...	...	...	56%	84.1M	5s
54200K	...	...	...	...	56%	78.9M	5s
54250K	...	...	...	...	56%	51.3M	5s
54300K	...	...	...	...	56%	71.5M	5s
54350K	...	...	...	...	56%	92.7M	5s
54400K	...	...	...	...	56%	63.5M	5s
54450K	...	...	...	...	56%	50.1M	5s
54500K	...	...	...	...	56%	64.9M	5s
54550K	...	...	...	...	56%	90.7M	5s
54600K	...	...	...	...	56%	77.9M	4s
54650K	...	...	...	...	56%	88.4M	4s
54700K	...	...	...	...	56%	92.6M	4s
54750K	...	...	...	...	56%	88.1M	4s
54800K	...	...	...	...	56%	42.5M	4s
54850K	...	...	...	...	56%	89.5M	4s
54900K	...	...	...	...	56%	89.3M	4s
54950K	...	...	...	...	56%	65.1M	4s
55000K	...	...	...	...	57%	72.9M	4s
55050K	...	...	...	...	57%	69.2M	4s
55100K	...	...	...	...	57%	88.9M	4s
55150K	...	...	...	...	57%	89.3M	4s
55200K	...	...	...	...	57%	76.8M	4s
55250K	...	...	...	...	57%	91.2M	4s
55300K	...	...	...	...	57%	45.8M	4s
55350K	...	...	...	...	57%	134K	5s
55400K	...	...	...	...	57%	46.0M	5s
55450K	...	...	...	...	57%	76.0M	5s
55500K	...	...	...	...	57%	84.5M	5s
55550K	...	...	...	...	57%	83.5M	5s
55600K	...	...	...	...	57%	70.9M	5s
55650K	...	...	...	...	57%	85.2M	5s
55700K	...	...	...	...	57%	85.0M	5s
55750K	...	...	...	...	57%	92.1M	5s
55800K	...	...	...	...	57%	39.6M	5s
55850K	...	...	...	...	57%	84.5M	5s
55900K	...	...	...	...	57%	90.4M	5s
55950K	...	...	...	...	58%	95.0M	5s
56000K	...	...	...	...	58%	76.1M	5s
56050K	...	...	...	...	58%	94.5M	5s
56100K	...	...	...	...	58%	91.2M	5s
56150K	...	...	...	...	58%	78.7M	4s
56200K	...	...	...	...	58%	37.5M	4s
56250K	...	...	...	...	58%	85.8M	4s
56300K	...	...	...	...	58%	42.1M	4s

56350K	...	...	...	...	58%	78.5M	4s
56400K	...	...	...	...	58%	59.6M	4s
56450K	...	...	...	...	58%	60.8M	4s
56500K	...	...	...	...	58%	74.9M	4s
56550K	...	...	...	...	58%	92.3M	4s
56600K	...	...	...	...	58%	76.8M	4s
56650K	...	...	...	...	58%	93.0M	4s
56700K	...	...	...	...	58%	90.7M	4s
56750K	...	...	...	...	58%	96.1M	4s
56800K	...	...	...	...	58%	69.7M	4s
56850K	...	...	...	...	58%	43.3M	4s
56900K	...	...	...	...	59%	92.0M	4s
56950K	...	...	...	...	59%	98.0M	4s
57000K	...	...	...	...	59%	82.3M	4s
57050K	...	...	...	...	59%	99.7M	4s
57100K	...	...	...	...	59%	84.2M	4s
57150K	...	...	...	...	59%	96.2M	4s
57200K	...	...	...	...	59%	82.7M	4s
57250K	...	...	...	...	59%	88.7M	4s
57300K	...	...	...	...	59%	94.7M	4s
57350K	...	...	...	...	59%	49.0M	4s
57400K	...	...	...	...	59%	134K	5s
57450K	...	...	...	...	59%	62.6M	5s
57500K	...	...	...	...	59%	81.3M	5s
57550K	...	...	...	...	59%	86.1M	4s
57600K	...	...	...	...	59%	70.6M	4s
57650K	...	...	...	...	59%	83.2M	4s
57700K	...	...	...	...	59%	70.8M	4s
57750K	...	...	...	...	59%	59.2M	4s
57800K	...	...	...	...	59%	60.9M	4s
57850K	...	...	...	...	60%	34.0M	4s
57900K	...	...	...	...	60%	73.8M	4s
57950K	...	...	...	...	60%	86.7M	4s
58000K	...	...	...	...	60%	64.0M	4s
58050K	...	...	...	...	60%	72.3M	4s
58100K	...	...	...	...	60%	72.9M	4s
58150K	...	...	...	...	60%	95.4M	4s
58200K	...	...	...	...	60%	81.7M	4s
58250K	...	...	...	...	60%	87.3M	4s
58300K	...	...	...	...	60%	86.0M	4s
58350K	...	...	...	...	60%	47.3M	4s
58400K	...	...	...	...	60%	57.5M	4s
58450K	...	...	...	...	60%	44.0M	4s
58500K	...	...	...	...	60%	63.3M	4s
58550K	...	...	...	...	60%	92.1M	4s
58600K	...	...	...	...	60%	76.1M	4s
58650K	...	...	...	...	60%	88.9M	4s
58700K	...	...	...	...	60%	82.3M	4s

58750K	...	...	...	...	60%	109M	4s
58800K	...	...	...	...	60%	79.5M	4s
58850K	...	...	...	...	61%	88.0M	4s
58900K	...	...	...	...	61%	40.5M	4s
58950K	...	...	...	...	61%	77.7M	4s
59000K	...	...	...	...	61%	75.7M	4s
59050K	...	...	...	...	61%	89.5M	4s
59100K	...	...	...	...	61%	84.4M	4s
59150K	...	...	...	...	61%	85.8M	4s
59200K	...	...	...	...	61%	77.6M	4s
59250K	...	...	...	...	61%	89.0M	4s
59300K	...	...	...	...	61%	94.3M	4s
59350K	...	...	...	...	61%	87.4M	4s
59400K	...	...	...	...	61%	42.3M	4s
59450K	...	...	...	...	61%	134K	4s
59500K	...	...	...	...	61%	78.8M	4s
59550K	...	...	...	...	61%	95.8M	4s
59600K	...	...	...	...	61%	71.5M	4s
59650K	...	...	...	...	61%	64.7M	4s
59700K	...	...	...	...	61%	70.5M	4s
59750K	...	...	...	...	61%	79.5M	4s
59800K	...	...	...	...	62%	52.4M	4s
59850K	...	...	...	...	62%	52.9M	4s
59900K	...	...	...	...	62%	41.1M	4s
59950K	...	...	...	...	62%	81.8M	4s
60000K	...	...	...	...	62%	79.8M	4s
60050K	...	...	...	...	62%	92.8M	4s
60100K	...	...	...	...	62%	98.3M	4s
60150K	...	...	...	...	62%	62.1M	4s
60200K	...	...	...	...	62%	72.2M	4s
60250K	...	...	...	...	62%	90.5M	4s
60300K	...	...	...	...	62%	86.6M	4s
60350K	...	...	...	...	62%	93.8M	4s
60400K	...	...	...	...	62%	42.0M	4s
60450K	...	...	...	...	62%	86.9M	4s
60500K	...	...	...	...	62%	74.2M	4s
60550K	...	...	...	...	62%	49.9M	4s
60600K	...	...	...	...	62%	72.9M	4s
60650K	...	...	...	...	62%	92.8M	4s
60700K	...	...	...	...	62%	78.1M	4s
60750K	...	...	...	...	63%	54.8M	4s
60800K	...	...	...	...	63%	67.6M	4s
60850K	...	...	...	...	63%	76.5M	4s
60900K	...	...	...	...	63%	83.3M	4s
60950K	...	...	...	...	63%	39.0M	4s
61000K	...	...	...	...	63%	63.4M	4s
61050K	...	...	...	...	63%	65.0M	4s
61100K	...	...	...	...	63%	82.2M	4s

61150K	...	...	...	...	63%	83.0M	4s
61200K	...	...	...	...	63%	74.4M	4s
61250K	...	...	...	...	63%	90.1M	4s
61300K	...	...	...	...	63%	102M	4s
61350K	...	...	...	...	63%	101M	4s
61400K	...	...	...	...	63%	71.1M	4s
61450K	...	...	...	...	63%	47.5M	4s
61500K	...	...	...	...	63%	135K	4s
61550K	...	...	...	...	63%	47.0M	4s
61600K	...	...	...	...	63%	81.1M	4s
61650K	...	...	...	...	63%	97.5M	4s
61700K	...	...	...	...	63%	109M	4s
61750K	...	...	...	...	64%	102M	4s
61800K	...	...	...	...	64%	95.4M	4s
61850K	...	...	...	...	64%	103M	4s
61900K	...	...	...	...	64%	103M	4s
61950K	...	...	...	...	64%	41.0M	4s
62000K	...	...	...	...	64%	61.3M	4s
62050K	...	...	...	...	64%	87.2M	4s
62100K	...	...	...	...	64%	96.3M	4s
62150K	...	...	...	...	64%	90.6M	4s
62200K	...	...	...	...	64%	71.2M	4s
62250K	...	...	...	...	64%	84.4M	4s
62300K	...	...	...	...	64%	77.4M	4s
62350K	...	...	...	...	64%	76.6M	4s
62400K	...	...	...	...	64%	68.2M	4s
62450K	...	...	...	...	64%	46.0M	4s
62500K	...	...	...	...	64%	73.5M	4s
62550K	...	...	...	...	64%	64.8M	4s
62600K	...	...	...	...	64%	46.4M	4s
62650K	...	...	...	...	64%	59.0M	4s
62700K	...	...	...	...	65%	76.2M	4s
62750K	...	...	...	...	65%	80.0M	4s
62800K	...	...	...	...	65%	70.7M	4s
62850K	...	...	...	...	65%	89.1M	4s
62900K	...	...	...	...	65%	98.6M	4s
62950K	...	...	...	...	65%	92.0M	4s
63000K	...	...	...	...	65%	40.6M	4s
63050K	...	...	...	...	65%	73.7M	4s
63100K	...	...	...	...	65%	70.3M	4s
63150K	...	...	...	...	65%	84.8M	4s
63200K	...	...	...	...	65%	70.0M	4s
63250K	...	...	...	...	65%	82.4M	4s
63300K	...	...	...	...	65%	78.8M	4s
63350K	...	...	...	...	65%	75.9M	4s
63400K	...	...	...	...	65%	67.9M	4s
63450K	...	...	...	...	65%	78.3M	4s
63500K	...	...	...	...	65%	44.9M	4s

63550K	...	...	...	...	65%	134K	4s
63600K	...	...	...	...	65%	52.8M	4s
63650K	...	...	...	...	66%	86.0M	4s
63700K	...	...	...	...	66%	86.1M	4s
63750K	...	...	...	...	66%	88.1M	4s
63800K	...	...	...	...	66%	81.4M	4s
63850K	...	...	...	...	66%	87.6M	4s
63900K	...	...	...	...	66%	90.5M	4s
63950K	...	...	...	...	66%	86.8M	4s
64000K	...	...	...	...	66%	39.7M	4s
64050K	...	...	...	...	66%	63.1M	4s
64100K	...	...	...	...	66%	87.0M	4s
64150K	...	...	...	...	66%	87.7M	4s
64200K	...	...	...	...	66%	79.6M	4s
64250K	...	...	...	...	66%	85.3M	4s
64300K	...	...	...	...	66%	83.0M	4s
64350K	...	...	...	...	66%	42.8M	4s
64400K	...	...	...	...	66%	78.7M	4s
64450K	...	...	...	...	66%	86.2M	4s
64500K	...	...	...	...	66%	45.5M	4s
64550K	...	...	...	...	66%	87.0M	4s
64600K	...	...	...	...	66%	69.3M	4s
64650K	...	...	...	...	67%	57.3M	4s
64700K	...	...	...	...	67%	74.8M	4s
64750K	...	...	...	...	67%	88.4M	4s
64800K	...	...	...	...	67%	71.4M	4s
64850K	...	...	...	...	67%	81.1M	4s
64900K	...	...	...	...	67%	88.3M	4s
64950K	...	...	...	...	67%	91.7M	4s
65000K	...	...	...	...	67%	42.2M	4s
65050K	...	...	...	...	67%	77.6M	4s
65100K	...	...	...	...	67%	92.6M	4s
65150K	...	...	...	...	67%	88.8M	4s
65200K	...	...	...	...	67%	81.1M	4s
65250K	...	...	...	...	67%	91.3M	4s
65300K	...	...	...	...	67%	90.5M	4s
65350K	...	...	...	...	67%	90.9M	4s
65400K	...	...	...	...	67%	76.7M	4s
65450K	...	...	...	...	67%	91.9M	4s
65500K	...	...	...	...	67%	87.2M	4s
65550K	...	...	...	...	67%	45.1M	4s
65600K	...	...	...	...	68%	134K	4s
65650K	...	...	...	...	68%	44.4M	4s
65700K	...	...	...	...	68%	82.4M	4s
65750K	...	...	...	...	68%	89.4M	4s
65800K	...	...	...	...	68%	68.8M	4s
65850K	...	...	...	...	68%	82.7M	4s
65900K	...	...	...	...	68%	83.6M	4s

65950K	...	...	...	...	68%	81.8M	4s
66000K	...	...	...	...	68%	73.1M	4s
66050K	...	...	...	...	68%	43.2M	4s
66100K	...	...	...	...	68%	91.7M	4s
66150K	...	...	...	...	68%	72.9M	4s
66200K	...	...	...	...	68%	70.1M	4s
66250K	...	...	...	...	68%	72.3M	4s
66300K	...	...	...	...	68%	81.4M	4s
66350K	...	...	...	...	68%	65.5M	4s
66400K	...	...	...	...	68%	72.8M	4s
66450K	...	...	...	...	68%	74.8M	4s
66500K	...	...	...	...	68%	102M	4s
66550K	...	...	...	...	69%	47.3M	4s
66600K	...	...	...	...	69%	82.7M	4s
66650K	...	...	...	...	69%	54.4M	4s
66700K	...	...	...	...	69%	67.5M	4s
66750K	...	...	...	...	69%	94.9M	4s
66800K	...	...	...	...	69%	77.4M	4s
66850K	...	...	...	...	69%	87.9M	4s
66900K	...	...	...	...	69%	105M	4s
66950K	...	...	...	...	69%	113M	4s
67000K	...	...	...	...	69%	79.4M	4s
67050K	...	...	...	...	69%	47.8M	4s
67100K	...	...	...	...	69%	80.0M	4s
67150K	...	...	...	...	69%	99.5M	4s
67200K	...	...	...	...	69%	70.9M	4s
67250K	...	...	...	...	69%	96.7M	4s
67300K	...	...	...	...	69%	92.1M	4s
67350K	...	...	...	...	69%	101M	4s
67400K	...	...	...	...	69%	79.9M	4s
67450K	...	...	...	...	69%	91.8M	4s
67500K	...	...	...	...	70%	88.8M	4s
67550K	...	...	...	...	70%	93.5M	4s
67600K	...	...	...	...	70%	42.8M	4s
67650K	...	...	...	...	70%	134K	4s
67700K	...	...	...	...	70%	54.5M	4s
67750K	...	...	...	...	70%	79.9M	4s
67800K	...	...	...	...	70%	70.4M	4s
67850K	...	...	...	...	70%	84.1M	4s
67900K	...	...	...	...	70%	96.4M	4s
67950K	...	...	...	...	70%	93.4M	4s
68000K	...	...	...	...	70%	79.9M	4s
68050K	...	...	...	...	70%	98.5M	4s
68100K	...	...	...	...	70%	44.6M	4s
68150K	...	...	...	...	70%	91.6M	4s
68200K	...	...	...	...	70%	72.7M	4s
68250K	...	...	...	...	70%	84.5M	4s
68300K	...	...	...	...	70%	92.4M	4s

68350K	...	...	...	...	70%	95.6M	4s
68400K	...	...	...	...	70%	71.4M	4s
68450K	...	...	...	...	70%	85.8M	4s
68500K	...	...	...	...	71%	90.5M	4s
68550K	...	...	...	...	71%	101M	4s
68600K	...	...	...	...	71%	47.4M	4s
68650K	...	...	...	...	71%	90.4M	4s
68700K	...	...	...	...	71%	100M	4s
68750K	...	...	...	...	71%	92.9M	3s
68800K	...	...	...	...	71%	62.7M	3s
68850K	...	...	...	...	71%	65.0M	3s
68900K	...	...	...	...	71%	79.2M	3s
68950K	...	...	...	...	71%	90.5M	3s
69000K	...	...	...	...	71%	85.4M	3s
69050K	...	...	...	...	71%	92.0M	3s
69100K	...	...	...	...	71%	50.8M	3s
69150K	...	...	...	...	71%	83.4M	3s
69200K	...	...	...	...	71%	82.1M	3s
69250K	...	...	...	...	71%	99.0M	3s
69300K	...	...	...	...	71%	98.6M	3s
69350K	...	...	...	...	71%	99.1M	3s
69400K	...	...	...	...	71%	78.1M	3s
69450K	...	...	...	...	72%	84.0M	3s
69500K	...	...	...	...	72%	101M	3s
69550K	...	...	...	...	72%	84.1M	3s
69600K	...	...	...	...	72%	84.8M	3s
69650K	...	...	...	...	72%	49.6M	3s
69700K	...	...	...	...	72%	133K	3s
69750K	...	...	...	...	72%	58.2M	3s
69800K	...	...	...	...	72%	78.2M	3s
69850K	...	...	...	...	72%	103M	3s
69900K	...	...	...	...	72%	108M	3s
69950K	...	...	...	...	72%	67.4M	3s
70000K	...	...	...	...	72%	67.5M	3s
70050K	...	...	...	...	72%	89.1M	3s
70100K	...	...	...	...	72%	92.3M	3s
70150K	...	...	...	...	72%	42.0M	3s
70200K	...	...	...	...	72%	76.3M	3s
70250K	...	...	...	...	72%	97.8M	3s
70300K	...	...	...	...	72%	92.5M	3s
70350K	...	...	...	...	72%	97.2M	3s
70400K	...	...	...	...	73%	76.9M	3s
70450K	...	...	...	...	73%	91.1M	3s
70500K	...	...	...	...	73%	90.0M	3s
70550K	...	...	...	...	73%	92.4M	3s
70600K	...	...	...	...	73%	78.3M	3s
70650K	...	...	...	...	73%	47.3M	3s
70700K	...	...	...	...	73%	79.1M	3s

70750K	...	...	...	...	73%	90.6M	3s
70800K	...	...	...	...	73%	62.5M	3s
70850K	...	...	...	...	73%	61.9M	3s
70900K	...	...	...	...	73%	72.9M	3s
70950K	...	...	...	...	73%	88.0M	3s
71000K	...	...	...	...	73%	74.9M	3s
71050K	...	...	...	...	73%	97.1M	3s
71100K	...	...	...	...	73%	91.3M	3s
71150K	...	...	...	...	73%	50.3M	3s
71200K	...	...	...	...	73%	68.5M	3s
71250K	...	...	...	...	73%	89.1M	3s
71300K	...	...	...	...	73%	76.3M	3s
71350K	...	...	...	...	73%	101M	3s
71400K	...	...	...	...	74%	82.2M	3s
71450K	...	...	...	...	74%	99.5M	3s
71500K	...	...	...	...	74%	91.1M	3s
71550K	...	...	...	...	74%	93.1M	3s
71600K	...	...	...	...	74%	75.9M	3s
71650K	...	...	...	...	74%	92.8M	3s
71700K	...	...	...	...	74%	48.9M	3s
71750K	...	...	...	...	74%	133K	3s
71800K	...	...	...	...	74%	44.6M	3s
71850K	...	...	...	...	74%	109M	3s
71900K	...	...	...	...	74%	105M	3s
71950K	...	...	...	...	74%	116M	3s
72000K	...	...	...	...	74%	83.6M	3s
72050K	...	...	...	...	74%	112M	3s
72100K	...	...	...	...	74%	104M	3s
72150K	...	...	...	...	74%	114M	3s
72200K	...	...	...	...	74%	45.8M	3s
72250K	...	...	...	...	74%	102M	3s
72300K	...	...	...	...	74%	116M	3s
72350K	...	...	...	...	75%	106M	3s
72400K	...	...	...	...	75%	72.1M	3s
72450K	...	...	...	...	75%	89.7M	3s
72500K	...	...	...	...	75%	88.6M	3s
72550K	...	...	...	...	75%	66.6M	3s
72600K	...	...	...	...	75%	80.3M	3s
72650K	...	...	...	...	75%	92.1M	3s
72700K	...	...	...	...	75%	44.0M	3s
72750K	...	...	...	...	75%	90.9M	3s
72800K	...	...	...	...	75%	86.5M	3s
72850K	...	...	...	...	75%	99.3M	3s
72900K	...	...	...	...	75%	103M	3s
72950K	...	...	...	...	75%	74.4M	3s
73000K	...	...	...	...	75%	51.4M	3s
73050K	...	...	...	...	75%	73.9M	3s
73100K	...	...	...	...	75%	87.7M	3s

73150K	...	...	...	...	75%	93.4M	3s
73200K	...	...	...	...	75%	45.2M	3s
73250K	...	...	...	...	75%	85.1M	3s
73300K	...	...	...	...	76%	81.3M	3s
73350K	...	...	...	...	76%	85.0M	3s
73400K	...	...	...	...	76%	52.4M	3s
73450K	...	...	...	...	76%	88.1M	3s
73500K	...	...	...	...	76%	89.3M	3s
73550K	...	...	...	...	76%	90.0M	3s
73600K	...	...	...	...	76%	70.4M	3s
73650K	...	...	...	...	76%	86.6M	3s
73700K	...	...	...	...	76%	88.9M	3s
73750K	...	...	...	...	76%	44.2M	3s
73800K	...	...	...	...	76%	133K	3s
73850K	...	...	...	...	76%	66.0M	3s
73900K	...	...	...	...	76%	83.4M	3s
73950K	...	...	...	...	76%	86.7M	3s
74000K	...	...	...	...	76%	82.1M	3s
74050K	...	...	...	...	76%	83.8M	3s
74100K	...	...	...	...	76%	78.0M	3s
74150K	...	...	...	...	76%	87.4M	3s
74200K	...	...	...	...	76%	75.0M	3s
74250K	...	...	...	...	77%	47.0M	3s
74300K	...	...	...	...	77%	87.0M	3s
74350K	...	...	...	...	77%	88.7M	3s
74400K	...	...	...	...	77%	74.6M	3s
74450K	...	...	...	...	77%	93.5M	3s
74500K	...	...	...	...	77%	89.7M	3s
74550K	...	...	...	...	77%	87.4M	3s
74600K	...	...	...	...	77%	74.5M	3s
74650K	...	...	...	...	77%	87.0M	3s
74700K	...	...	...	...	77%	91.7M	3s
74750K	...	...	...	...	77%	41.9M	3s
74800K	...	...	...	...	77%	70.0M	3s
74850K	...	...	...	...	77%	84.5M	3s
74900K	...	...	...	...	77%	17.7M	3s
74950K	...	...	...	...	77%	57.8M	3s
75000K	...	...	...	...	77%	63.9M	3s
75050K	...	...	...	...	77%	3.67M	3s
75100K	...	...	...	...	77%	58.3M	3s
75150K	...	...	...	...	77%	50.8M	3s
75200K	...	...	...	...	77%	64.5M	3s
75250K	...	...	...	...	78%	37.0M	3s
75300K	...	...	...	...	78%	53.4M	3s
75350K	...	...	...	...	78%	46.8M	3s
75400K	...	...	...	...	78%	65.0M	3s
75450K	...	...	...	...	78%	60.1M	3s
75500K	...	...	...	...	78%	52.6M	3s

75550K	...	...	...	...	78%	55.5M	3s
75600K	...	...	...	...	78%	70.7M	3s
75650K	...	...	...	...	78%	69.9M	3s
75700K	...	...	...	...	78%	64.5M	3s
75750K	...	...	...	...	78%	52.7M	3s
75800K	...	...	...	...	78%	28.9M	3s
75850K	...	...	...	...	78%	141K	3s
75900K	...	...	...	...	78%	77.4M	3s
75950K	...	...	...	...	78%	91.1M	3s
76000K	...	...	...	...	78%	81.4M	3s
76050K	...	...	...	...	78%	97.7M	3s
76100K	...	...	...	...	78%	95.5M	3s
76150K	...	...	...	...	78%	91.3M	3s
76200K	...	...	...	...	79%	3.98M	3s
76250K	...	...	...	...	79%	63.9M	3s
76300K	...	...	...	...	79%	35.0M	3s
76350K	...	...	...	...	79%	57.7M	3s
76400K	...	...	...	...	79%	70.2M	3s
76450K	...	...	...	...	79%	66.9M	3s
76500K	...	...	...	...	79%	66.5M	3s
76550K	...	...	...	...	79%	49.2M	3s
76600K	...	...	...	...	79%	67.1M	3s
76650K	...	...	...	...	79%	78.3M	3s
76700K	...	...	...	...	79%	6.73M	3s
76750K	...	...	...	...	79%	48.6M	3s
76800K	...	...	...	...	79%	32.0M	3s
76850K	...	...	...	...	79%	64.0M	3s
76900K	...	...	...	...	79%	69.4M	3s
76950K	...	...	...	...	79%	65.6M	3s
77000K	...	...	...	...	79%	55.1M	3s
77050K	...	...	...	...	79%	62.8M	3s
77100K	...	...	...	...	79%	55.8M	3s
77150K	...	...	...	...	80%	57.1M	3s
77200K	...	...	...	...	80%	55.2M	3s
77250K	...	...	...	...	80%	47.5M	3s
77300K	...	...	...	...	80%	32.4M	3s
77350K	...	...	...	...	80%	66.6M	3s
77400K	...	...	...	...	80%	46.5M	3s
77450K	...	...	...	...	80%	48.6M	3s
77500K	...	...	...	...	80%	55.9M	3s
77550K	...	...	...	...	80%	63.7M	3s
77600K	...	...	...	...	80%	46.9M	3s
77650K	...	...	...	...	80%	56.8M	2s
77700K	...	...	...	...	80%	56.3M	2s
77750K	...	...	...	...	80%	60.6M	2s
77800K	...	...	...	...	80%	28.3M	2s
77850K	...	...	...	...	80%	54.5M	2s
77900K	...	...	...	...	80%	144K	3s

77950K	...	...	...	...	80%	94.7M	3s
78000K	...	...	...	...	80%	83.1M	3s
78050K	...	...	...	...	80%	98.9M	3s
78100K	...	...	...	...	80%	97.5M	3s
78150K	...	...	...	...	81%	100M	2s
78200K	...	...	...	...	81%	84.0M	2s
78250K	...	...	...	...	81%	97.9M	2s
78300K	...	...	...	...	81%	99.2M	2s
78350K	...	...	...	...	81%	38.7M	2s
78400K	...	...	...	...	81%	72.2M	2s
78450K	...	...	...	...	81%	103M	2s
78500K	...	...	...	...	81%	6.02M	2s
78550K	...	...	...	...	81%	44.9M	2s
78600K	...	...	...	...	81%	57.0M	2s
78650K	...	...	...	...	81%	61.3M	2s
78700K	...	...	...	...	81%	68.7M	2s
78750K	...	...	...	...	81%	57.6M	2s
78800K	...	...	...	...	81%	68.6M	2s
78850K	...	...	...	...	81%	38.6M	2s
78900K	...	...	...	...	81%	45.9M	2s
78950K	...	...	...	...	81%	58.0M	2s
79000K	...	...	...	...	81%	69.1M	2s
79050K	...	...	...	...	81%	66.2M	2s
79100K	...	...	...	...	82%	67.7M	2s
79150K	...	...	...	...	82%	9.49M	2s
79200K	...	...	...	...	82%	50.2M	2s
79250K	...	...	...	...	82%	60.5M	2s
79300K	...	...	...	...	82%	66.7M	2s
79350K	...	...	...	...	82%	36.5M	2s
79400K	...	...	...	...	82%	62.1M	2s
79450K	...	...	...	...	82%	70.0M	2s
79500K	...	...	...	...	82%	70.2M	2s
79550K	...	...	...	...	82%	61.0M	2s
79600K	...	...	...	...	82%	67.7M	2s
79650K	...	...	...	...	82%	63.4M	2s
79700K	...	...	...	...	82%	58.5M	2s
79750K	...	...	...	...	82%	55.4M	2s
79800K	...	...	...	...	82%	9.03M	2s
79850K	...	...	...	...	82%	56.5M	2s
79900K	...	...	...	...	82%	38.3M	2s
79950K	...	...	...	...	82%	142K	2s
80000K	...	...	...	...	82%	78.8M	2s
80050K	...	...	...	...	83%	80.1M	2s
80100K	...	...	...	...	83%	96.5M	2s
80150K	...	...	...	...	83%	88.0M	2s
80200K	...	...	...	...	83%	83.5M	2s
80250K	...	...	...	...	83%	93.5M	2s
80300K	...	...	...	...	83%	95.9M	2s

80350K	...	...	...	...	83%	95.1M	2s
80400K	...	...	...	...	83%	43.6M	2s
80450K	...	...	...	...	83%	94.4M	2s
80500K	...	...	...	...	83%	99.4M	2s
80550K	...	...	...	...	83%	6.03M	2s
80600K	...	...	...	...	83%	45.1M	2s
80650K	...	...	...	...	83%	56.2M	2s
80700K	...	...	...	...	83%	64.2M	2s
80750K	...	...	...	...	83%	64.1M	2s
80800K	...	...	...	...	83%	64.0M	2s
80850K	...	...	...	...	83%	60.2M	2s
80900K	...	...	...	...	83%	37.9M	2s
80950K	...	...	...	...	83%	64.0M	2s
81000K	...	...	...	...	83%	65.0M	2s
81050K	...	...	...	...	84%	57.3M	2s
81100K	...	...	...	...	84%	66.3M	2s
81150K	...	...	...	...	84%	66.5M	2s
81200K	...	...	...	...	84%	8.86M	2s
81250K	...	...	...	...	84%	46.1M	2s
81300K	...	...	...	...	84%	52.7M	2s
81350K	...	...	...	...	84%	61.8M	2s
81400K	...	...	...	...	84%	63.1M	2s
81450K	...	...	...	...	84%	36.4M	2s
81500K	...	...	...	...	84%	53.4M	2s
81550K	...	...	...	...	84%	63.5M	2s
81600K	...	...	...	...	84%	55.2M	2s
81650K	...	...	...	...	84%	51.5M	2s
81700K	...	...	...	...	84%	50.4M	2s
81750K	...	...	...	...	84%	62.6M	2s
81800K	...	...	...	...	84%	63.7M	2s
81850K	...	...	...	...	84%	10.8M	2s
81900K	...	...	...	...	84%	61.9M	2s
81950K	...	...	...	...	84%	219K	2s
82000K	...	...	...	...	85%	389K	2s
82050K	...	...	...	...	85%	65.0M	2s
82100K	...	...	...	...	85%	101M	2s
82150K	...	...	...	...	85%	98.2M	2s
82200K	...	...	...	...	85%	76.0M	2s
82250K	...	...	...	...	85%	109M	2s
82300K	...	...	...	...	85%	108M	2s
82350K	...	...	...	...	85%	108M	2s
82400K	...	...	...	...	85%	88.2M	2s
82450K	...	...	...	...	85%	35.7M	2s
82500K	...	...	...	...	85%	80.5M	2s
82550K	...	...	...	...	85%	96.8M	2s
82600K	...	...	...	...	85%	8.89M	2s
82650K	...	...	...	...	85%	53.0M	2s
82700K	...	...	...	...	85%	59.3M	2s

82750K	...	...	...	...	85%	54.7M	2s
82800K	...	...	...	...	85%	46.5M	2s
82850K	...	...	...	...	85%	52.3M	2s
82900K	...	...	...	...	85%	53.5M	2s
82950K	...	...	...	...	86%	53.1M	2s
83000K	...	...	...	...	86%	37.1M	2s
83050K	...	...	...	...	86%	65.9M	2s
83100K	...	...	...	...	86%	64.6M	2s
83150K	...	...	...	...	86%	56.4M	2s
83200K	...	...	...	...	86%	64.5M	2s
83250K	...	...	...	...	86%	10.6M	2s
83300K	...	...	...	...	86%	57.1M	2s
83350K	...	...	...	...	86%	60.2M	2s
83400K	...	...	...	...	86%	55.6M	2s
83450K	...	...	...	...	86%	45.8M	2s
83500K	...	...	...	...	86%	17.7M	2s
83550K	...	...	...	...	86%	50.7M	2s
83600K	...	...	...	...	86%	47.0M	2s
83650K	...	...	...	...	86%	63.6M	2s
83700K	...	...	...	...	86%	66.4M	2s
83750K	...	...	...	...	86%	63.7M	2s
83800K	...	...	...	...	86%	54.6M	2s
83850K	...	...	...	...	86%	61.7M	2s
83900K	...	...	...	...	87%	15.2M	2s
83950K	...	...	...	...	87%	58.9M	2s
84000K	...	...	...	...	87%	141K	2s
84050K	...	...	...	...	87%	48.8M	2s
84100K	...	...	...	...	87%	87.3M	2s
84150K	...	...	...	...	87%	82.0M	2s
84200K	...	...	...	...	87%	49.4M	2s
84250K	...	...	...	...	87%	33.3M	2s
84300K	...	...	...	...	87%	65.1M	2s
84350K	...	...	...	...	87%	55.5M	2s
84400K	...	...	...	...	87%	54.5M	2s
84450K	...	...	...	...	87%	67.1M	2s
84500K	...	...	...	...	87%	31.2M	2s
84550K	...	...	...	...	87%	26.9M	2s
84600K	...	...	...	...	87%	33.8M	2s
84650K	...	...	...	...	87%	16.1M	2s
84700K	...	...	...	...	87%	18.1M	2s
84750K	...	...	...	...	87%	52.0M	2s
84800K	...	...	...	...	87%	42.8M	2s
84850K	...	...	...	...	87%	16.0M	2s
84900K	...	...	...	...	88%	66.3M	2s
84950K	...	...	...	...	88%	71.5M	2s
85000K	...	...	...	...	88%	57.6M	2s
85050K	...	...	...	...	88%	25.9M	2s
85100K	...	...	...	...	88%	50.6M	2s

85150K	...	...	...	...	...	88%	50.7M	2s
85200K	...	...	...	...	...	88%	19.2M	2s
85250K	...	...	...	...	...	88%	68.9M	2s
85300K	...	...	...	...	...	88%	32.7M	2s
85350K	...	...	...	...	...	88%	31.4M	2s
85400K	...	...	...	...	...	88%	24.3M	2s
85450K	...	...	...	...	...	88%	56.5M	2s
85500K	...	...	...	...	...	88%	70.2M	2s
85550K	...	...	...	...	...	88%	18.5M	2s
85600K	...	...	...	...	...	88%	62.6M	2s
85650K	...	...	...	...	...	88%	66.9M	1s
85700K	...	...	...	...	...	88%	61.5M	1s
85750K	...	...	...	...	...	88%	62.4M	1s
85800K	...	...	...	...	...	88%	54.5M	1s
85850K	...	...	...	...	...	89%	49.4M	1s
85900K	...	...	...	...	...	89%	66.2M	1s
85950K	...	...	...	...	...	89%	63.0M	1s
86000K	...	...	...	...	...	89%	50.6M	1s
86050K	...	...	...	...	...	89%	143K	1s
86100K	...	...	...	...	...	89%	50.4M	1s
86150K	...	...	...	...	...	89%	31.8M	1s
86200K	...	...	...	...	...	89%	107M	1s
86250K	...	...	...	...	...	89%	101M	1s
86300K	...	...	...	...	...	89%	110M	1s
86350K	...	...	...	...	...	89%	115M	1s
86400K	...	...	...	...	...	89%	111M	1s
86450K	...	...	...	...	...	89%	117M	1s
86500K	...	...	...	...	...	89%	118M	1s
86550K	...	...	...	...	...	89%	126M	1s
86600K	...	...	...	...	...	89%	48.0M	1s
86650K	...	...	...	...	...	89%	107M	1s
86700K	...	...	...	...	...	89%	6.11M	1s
86750K	...	...	...	...	...	89%	56.3M	1s
86800K	...	...	...	...	...	90%	73.6M	1s
86850K	...	...	...	...	...	90%	50.0M	1s
86900K	...	...	...	...	...	90%	67.0M	1s
86950K	...	...	...	...	...	90%	82.6M	1s
87000K	...	...	...	...	...	90%	89.5M	1s
87050K	...	...	...	...	...	90%	73.5M	1s
87100K	...	...	...	...	...	90%	3.96M	1s
87150K	...	...	...	...	...	90%	61.2M	1s
87200K	...	...	...	...	...	90%	78.0M	1s
87250K	...	...	...	...	...	90%	72.8M	1s
87300K	...	...	...	...	...	90%	80.1M	1s
87350K	...	...	...	...	...	90%	81.2M	1s
87400K	...	...	...	...	...	90%	80.6M	1s
87450K	...	...	...	...	...	90%	63.2M	1s
87500K	...	...	...	...	...	90%	53.2M	1s

87550K	...	...	...	...	90%	84.4M	1s
87600K	...	...	...	...	90%	41.4M	1s
87650K	...	...	...	...	90%	71.3M	1s
87700K	...	...	...	...	90%	84.6M	1s
87750K	...	...	...	...	90%	87.2M	1s
87800K	...	...	...	...	91%	86.7M	1s
87850K	...	...	...	...	91%	70.0M	1s
87900K	...	...	...	...	91%	88.3M	1s
87950K	...	...	...	...	91%	82.3M	1s
88000K	...	...	...	...	91%	12.7M	1s
88050K	...	...	...	...	91%	84.3M	1s
88100K	...	...	...	...	91%	140K	1s
88150K	...	...	...	...	91%	57.4M	1s
88200K	...	...	...	...	91%	72.4M	1s
88250K	...	...	...	...	91%	88.5M	1s
88300K	...	...	...	...	91%	75.4M	1s
88350K	...	...	...	...	91%	81.8M	1s
88400K	...	...	...	...	91%	80.3M	1s
88450K	...	...	...	...	91%	93.6M	1s
88500K	...	...	...	...	91%	79.6M	1s
88550K	...	...	...	...	91%	91.4M	1s
88600K	...	...	...	...	91%	81.2M	1s
88650K	...	...	...	...	91%	40.0M	1s
88700K	...	...	...	...	91%	18.3M	1s
88750K	...	...	...	...	92%	9.26M	1s
88800K	...	...	...	...	92%	55.2M	1s
88850K	...	...	...	...	92%	58.8M	1s
88900K	...	...	...	...	92%	69.5M	1s
88950K	...	...	...	...	92%	75.7M	1s
89000K	...	...	...	...	92%	86.5M	1s
89050K	...	...	...	...	92%	89.4M	1s
89100K	...	...	...	...	92%	85.8M	1s
89150K	...	...	...	...	92%	36.3M	1s
89200K	...	...	...	...	92%	86.7M	1s
89250K	...	...	...	...	92%	91.4M	1s
89300K	...	...	...	...	92%	93.0M	1s
89350K	...	...	...	...	92%	81.9M	1s
89400K	...	...	...	...	92%	7.16M	1s
89450K	...	...	...	...	92%	54.5M	1s
89500K	...	...	...	...	92%	53.2M	1s
89550K	...	...	...	...	92%	71.1M	1s
89600K	...	...	...	...	92%	74.5M	1s
89650K	...	...	...	...	92%	40.4M	1s
89700K	...	...	...	...	93%	57.8M	1s
89750K	...	...	...	...	93%	76.4M	1s
89800K	...	...	...	...	93%	86.0M	1s
89850K	...	...	...	...	93%	82.3M	1s
89900K	...	...	...	...	93%	64.1M	1s

89950K	...	...	...	...	93%	68.0M	1s
90000K	...	...	...	...	93%	78.1M	1s
90050K	...	...	...	...	93%	6.42M	1s
90100K	...	...	...	...	93%	69.0M	1s
90150K	...	...	...	...	93%	142K	1s
90200K	...	...	...	...	93%	32.2M	1s
90250K	...	...	...	...	93%	75.8M	1s
90300K	...	...	...	...	93%	104M	1s
90350K	...	...	...	...	93%	114M	1s
90400K	...	...	...	...	93%	104M	1s
90450K	...	...	...	...	93%	107M	1s
90500K	...	...	...	...	93%	111M	1s
90550K	...	...	...	...	93%	115M	1s
90600K	...	...	...	...	93%	91.2M	1s
90650K	...	...	...	...	93%	113M	1s
90700K	...	...	...	...	94%	50.3M	1s
90750K	...	...	...	...	94%	94.9M	1s
90800K	...	...	...	...	94%	6.10M	1s
90850K	...	...	...	...	94%	71.7M	1s
90900K	...	...	...	...	94%	72.9M	1s
90950K	...	...	...	...	94%	53.6M	1s
91000K	...	...	...	...	94%	86.2M	1s
91050K	...	...	...	...	94%	92.2M	1s
91100K	...	...	...	...	94%	86.0M	1s
91150K	...	...	...	...	94%	75.4M	1s
91200K	...	...	...	...	94%	42.8M	1s
91250K	...	...	...	...	94%	87.2M	1s
91300K	...	...	...	...	94%	90.2M	1s
91350K	...	...	...	...	94%	77.1M	1s
91400K	...	...	...	...	94%	79.7M	1s
91450K	...	...	...	...	94%	6.92M	1s
91500K	...	...	...	...	94%	58.1M	1s
91550K	...	...	...	...	94%	51.0M	1s
91600K	...	...	...	...	94%	68.4M	1s
91650K	...	...	...	...	95%	80.1M	1s
91700K	...	...	...	...	95%	42.6M	1s
91750K	...	...	...	...	95%	69.9M	1s
91800K	...	...	...	...	95%	78.3M	1s
91850K	...	...	...	...	95%	86.6M	1s
91900K	...	...	...	...	95%	77.9M	1s
91950K	...	...	...	...	95%	81.3M	1s
92000K	...	...	...	...	95%	90.8M	1s
92050K	...	...	...	...	95%	84.4M	1s
92100K	...	...	...	...	95%	7.22M	1s
92150K	...	...	...	...	95%	67.9M	1s
92200K	...	...	...	...	95%	140K	1s
92250K	...	...	...	...	95%	101M	1s
92300K	...	...	...	...	95%	115M	1s

92350K	...	...	...	...	95%	110M	1s
92400K	...	...	...	...	95%	93.5M	1s
92450K	...	...	...	...	95%	110M	1s
92500K	...	...	...	...	95%	28.9M	1s
92550K	...	...	...	...	95%	122M	1s
92600K	...	...	...	...	96%	98.7M	1s
92650K	...	...	...	...	96%	129M	1s
92700K	...	...	...	...	96%	124M	1s
92750K	...	...	...	...	96%	47.4M	1s
92800K	...	...	...	...	96%	91.1M	1s
92850K	...	...	...	...	96%	8.71M	1s
92900K	...	...	...	...	96%	66.0M	1s
92950K	...	...	...	...	96%	87.6M	1s
93000K	...	...	...	...	96%	76.5M	0s
93050K	...	...	...	...	96%	77.5M	0s
93100K	...	...	...	...	96%	78.5M	0s
93150K	...	...	...	...	96%	90.1M	0s
93200K	...	...	...	...	96%	75.3M	0s
93250K	...	...	...	...	96%	43.0M	0s
93300K	...	...	...	...	96%	89.2M	0s
93350K	...	...	...	...	96%	94.2M	0s
93400K	...	...	...	...	96%	73.9M	0s
93450K	...	...	...	...	96%	86.6M	0s
93500K	...	...	...	...	96%	6.62M	0s
93550K	...	...	...	...	97%	58.8M	0s
93600K	...	...	...	...	97%	91.1M	0s
93650K	...	...	...	...	97%	77.5M	0s
93700K	...	...	...	...	97%	83.3M	0s
93750K	...	...	...	...	97%	36.4M	0s
93800K	...	...	...	...	97%	71.7M	0s
93850K	...	...	...	...	97%	75.5M	0s
93900K	...	...	...	...	97%	88.0M	0s
93950K	...	...	...	...	97%	97.4M	0s
94000K	...	...	...	...	97%	85.1M	0s
94050K	...	...	...	...	97%	83.0M	0s
94100K	...	...	...	...	97%	89.9M	0s
94150K	...	...	...	...	97%	6.89M	0s
94200K	...	...	...	...	97%	61.0M	0s
94250K	...	...	...	...	97%	140K	0s
94300K	...	...	...	...	97%	103M	0s
94350K	...	...	...	...	97%	133M	0s
94400K	...	...	...	...	97%	112M	0s
94450K	...	...	...	...	97%	125M	0s
94500K	...	...	...	...	97%	140M	0s
94550K	...	...	...	...	98%	142M	0s
94600K	...	...	...	...	98%	110M	0s
94650K	...	...	...	...	98%	123M	0s
94700K	...	...	...	...	98%	120M	0s


```

94750K ... .. 98% 127M 0s
94800K ... .. 98% 18.1M 0s
94850K ... .. 98% 148M 0s
94900K ... .. 98% 10.2M 0s
94950K ... .. 98% 62.0M 0s
95000K ... .. 98% 92.7M 0s
95050K ... .. 98% 74.6M 0s
95100K ... .. 98% 77.5M 0s
95150K ... .. 98% 93.5M 0s
95200K ... .. 98% 76.4M 0s
95250K ... .. 98% 71.1M 0s
95300K ... .. 98% 43.2M 0s
95350K ... .. 98% 94.2M 0s
95400K ... .. 98% 89.9M 0s
95450K ... .. 98% 95.4M 0s
95500K ... .. 99% 77.5M 0s
95550K ... .. 99% 6.79M 0s
95600K ... .. 99% 42.9M 0s
95650K ... .. 99% 64.8M 0s
95700K ... .. 99% 65.7M 0s
95750K ... .. 99% 86.0M 0s
95800K ... .. 99% 32.1M 0s
95850K ... .. 99% 80.9M 0s
95900K ... .. 99% 68.0M 0s
95950K ... .. 99% 85.2M 0s
96000K ... .. 99% 80.6M 0s
96050K ... .. 99% 81.3M 0s
96100K ... .. 99% 90.5M 0s
96150K ... .. 99% 89.6M 0s
96200K ... .. 99% 64.5M 0s
96250K ... .. 99% 7.63M 0s
96300K ... .. 99% 141K 0s
96350K ... .. 99% 23.6M 0s
96400K ... .. 99% 102M 0s
96450K ... .. 100% 117M=14s

```

2026-09-06 13:50:00 (6.60 MB/s) - 'pretrained_model/pretrained_simclr_model.pth'
saved [98808459/98808459]

```

[ ]: # Setup cell.
      %pip install thop
      import torch
      import os
      import sys
      import types
      import importlib

```

```

import pandas as pd
import numpy as np
import torch.optim as optim
import torch.nn as nn
import random
from thop import profile, clever_format
from torch.utils.data import DataLoader
from torchvision.datasets import CIFAR10
import matplotlib.pyplot as plt
%matplotlib inline

if "imp" not in sys.modules:
    imp = types.ModuleType("imp")
    imp.reload = importlib.reload
    sys.modules["imp"] = imp

%load_ext autoreload
%autoreload 2

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

```

Collecting thop

```

Downloading thop-0.1.1.post2209072238-py3-none-any.whl.metadata (2.7 kB)
Requirement already satisfied: torch in /usr/local/lib/python3.13/dist-packages
(from thop) (2.11.0+cu128)
Requirement already satisfied: filelock in /usr/local/lib/python3.13/dist-
packages (from torch->thop) (3.32.4)
Requirement already satisfied: typing-extensions>=4.10.0 in
/usr/local/lib/python3.13/dist-packages (from torch->thop) (4.16.0)
Requirement already satisfied: setuptools<82 in /usr/local/lib/python3.13/dist-
packages (from torch->thop) (75.2.0)
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.13/dist-
packages (from torch->thop) (1.14.0)
Requirement already satisfied: networkx>=2.5.1 in
/usr/local/lib/python3.13/dist-packages (from torch->thop) (3.6.1)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.13/dist-packages
(from torch->thop) (3.1.6)
Requirement already satisfied: fsspec>=0.8.5 in /usr/local/lib/python3.13/dist-
packages (from torch->thop) (2025.3.0)
Requirement already satisfied: cuda-toolkit==12.8.1 in
/usr/local/lib/python3.13/dist-packages (from cuda-toolkit[cublas,cudart,cufft,c
ufile,cupti,curand,cusolver,cusparsenvjitlink,nvrtc,nvtx]==12.8.1;
platform_system == "Linux"->torch->thop) (12.8.1)
Requirement already satisfied: cuda-bindings<13,>=12.9.4 in
/usr/local/lib/python3.13/dist-packages (from torch->thop) (12.9.7)
Requirement already satisfied: nvidia-cudnn-cu12==9.19.0.56 in
/usr/local/lib/python3.13/dist-packages (from torch->thop) (9.19.0.56)

```

Requirement already satisfied: nvidia-cusparse-cu12==0.7.1 in
 /usr/local/lib/python3.13/dist-packages (from torch->thop) (0.7.1)

Requirement already satisfied: nvidia-nccl-cu12==2.28.9 in
 /usr/local/lib/python3.13/dist-packages (from torch->thop) (2.28.9)

Requirement already satisfied: nvidia-nvshmem-cu12==3.4.5 in
 /usr/local/lib/python3.13/dist-packages (from torch->thop) (3.4.5)

Requirement already satisfied: triton==3.6.0 in /usr/local/lib/python3.13/dist-
 packages (from torch->thop) (3.6.0)

Requirement already satisfied: nvidia-cublas-cu12==12.8.4.1.* in
 /usr/local/lib/python3.13/dist-packages (from cuda-toolkit[cublas,cudart,cufft,c
 ufile,cupti,curand,cusolver,cusparsenvjitlink,nvrtc,nvtx]==12.8.1;
 platform_system == "Linux"->torch->thop) (12.8.4.1)

Requirement already satisfied: nvidia-cuda-runtime-cu12==12.8.90.* in
 /usr/local/lib/python3.13/dist-packages (from cuda-toolkit[cublas,cudart,cufft,c
 ufile,cupti,curand,cusolver,cusparsenvjitlink,nvrtc,nvtx]==12.8.1;
 platform_system == "Linux"->torch->thop) (12.8.90)

Requirement already satisfied: nvidia-cufft-cu12==11.3.3.83.* in
 /usr/local/lib/python3.13/dist-packages (from cuda-toolkit[cublas,cudart,cufft,c
 ufile,cupti,curand,cusolver,cusparsenvjitlink,nvrtc,nvtx]==12.8.1;
 platform_system == "Linux"->torch->thop) (11.3.3.83)

Requirement already satisfied: nvidia-cufile-cu12==1.13.1.3.* in
 /usr/local/lib/python3.13/dist-packages (from cuda-toolkit[cublas,cudart,cufft,c
 ufile,cupti,curand,cusolver,cusparsenvjitlink,nvrtc,nvtx]==12.8.1;
 platform_system == "Linux"->torch->thop) (1.13.1.3)

Requirement already satisfied: nvidia-cuda-cupti-cu12==12.8.90.* in
 /usr/local/lib/python3.13/dist-packages (from cuda-toolkit[cublas,cudart,cufft,c
 ufile,cupti,curand,cusolver,cusparsenvjitlink,nvrtc,nvtx]==12.8.1;
 platform_system == "Linux"->torch->thop) (12.8.90)

Requirement already satisfied: nvidia-curand-cu12==10.3.9.90.* in
 /usr/local/lib/python3.13/dist-packages (from cuda-toolkit[cublas,cudart,cufft,c
 ufile,cupti,curand,cusolver,cusparsenvjitlink,nvrtc,nvtx]==12.8.1;
 platform_system == "Linux"->torch->thop) (10.3.9.90)

Requirement already satisfied: nvidia-cusolver-cu12==11.7.3.90.* in
 /usr/local/lib/python3.13/dist-packages (from cuda-toolkit[cublas,cudart,cufft,c
 ufile,cupti,curand,cusolver,cusparsenvjitlink,nvrtc,nvtx]==12.8.1;
 platform_system == "Linux"->torch->thop) (11.7.3.90)

Requirement already satisfied: nvidia-cusparsenvjitlink-cu12==12.5.8.93.* in
 /usr/local/lib/python3.13/dist-packages (from cuda-toolkit[cublas,cudart,cufft,c
 ufile,cupti,curand,cusolver,cusparsenvjitlink,nvrtc,nvtx]==12.8.1;
 platform_system == "Linux"->torch->thop) (12.5.8.93)

Requirement already satisfied: nvidia-nvjitlink-cu12==12.8.93.* in
 /usr/local/lib/python3.13/dist-packages (from cuda-toolkit[cublas,cudart,cufft,c
 ufile,cupti,curand,cusolver,cusparsenvjitlink,nvrtc,nvtx]==12.8.1;
 platform_system == "Linux"->torch->thop) (12.8.93)

Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.8.93.* in
 /usr/local/lib/python3.13/dist-packages (from cuda-toolkit[cublas,cudart,cufft,c
 ufile,cupti,curand,cusolver,cusparsenvjitlink,nvrtc,nvtx]==12.8.1;
 platform_system == "Linux"->torch->thop) (12.8.93)

```
Requirement already satisfied: nvidia-nvtx-cu12==12.8.90.* in
/usr/local/lib/python3.13/dist-packages (from cuda-toolkit[cublas,cudart,cufft,cufile,cupti,curand,cusolver,cusparsenvjitlink,nvrtc,nvtx]==12.8.1; platform_system == "Linux"->torch->thop) (12.8.90)
Requirement already satisfied: cuda-pathfinder~=1.1 in
/usr/local/lib/python3.13/dist-packages (from cuda-bindings<13,>=12.9.4->torch->thop) (1.7.0)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in
/usr/local/lib/python3.13/dist-packages (from sympy>=1.13.3->torch->thop) (1.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.13/dist-packages (from jinja2->torch->thop) (3.0.3)
Downloading thop-0.1.1.post2209072238-py3-none-any.whl (15 kB)
Installing collected packages: thop
Successfully installed thop-0.1.1.post2209072238
```

2 Data Augmentation

Our first step is to perform data augmentation. Implement the `compute_train_transform()` function in `cs231n/simclr/data_utils.py` to apply the following random transformations:

1. Randomly resize and crop to 32x32.
2. Horizontally flip the image with probability 0.5
3. With a probability of 0.8, apply color jitter (see `compute_train_transform()` for definition)
4. With a probability of 0.2, convert the image to grayscale

Now complete `compute_train_transform()` and `CIFAR10Pair.__getitem__()` in `cs231n/simclr/data_utils.py` to apply the data augmentation transform and generate $\tilde{x}_i$ and $\tilde{x}_j$.

Test to make sure that your data augmentation code is correct:

```
[ ]: from cs231n.simclr.data_utils import *
    from cs231n.simclr.contrastive_loss import *

    answers = torch.load('simclr_sanity_check.key')

[ ]: from PIL import Image
    import torchvision
    from torchvision.datasets import CIFAR10

    def test_data_augmentation(correct_output=None):
        train_transform = compute_train_transform(seed=2147483647)
        trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
        ↪download=True, transform=train_transform)
        trainloader = torch.utils.data.DataLoader(trainset, batch_size=2,
        ↪shuffle=False, num_workers=2)
        dataiter = iter(trainloader)
```

```


images, labels = next(dataiter)
img = torchvision.utils.make_grid(images)
img = img / 2 + 0.5      # unnormalize
npimg = img.numpy()
plt.imshow(np.transpose(npimg, (1, 2, 0)))
plt.show()
output = images

print("Maximum error in data augmentation: %g"%rel_error( output.numpy(),
↪correct_output.numpy()))

# Should be less than 1e-07.
test_data_augmentation(answers['data_augmentation'])

```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers). Got range [-0.4334516..1.8768656].



Maximum error in data augmentation: 0

3 Base Encoder and Projection Head

The next steps are to apply the base encoder and projection head to the augmented samples $\tilde{x}_i$ and $\tilde{x}_j$.

The base encoder f extracts representation vectors for the augmented samples. The SimCLR paper found that using deeper and wider models improved performance and thus chose [ResNet](#) to use as the base encoder. The output of the base encoder are the representation vectors $h_i = f(\tilde{x}_i)$ and $h_j = f(\tilde{x}_j)$.

The projection head g is a small neural network that maps the representation vectors h_i and h_j to the space where the contrastive loss is applied. The paper found that using a nonlinear projection head improved the representation quality of the layer before it. Specifically, they used a MLP with one hidden layer as the projection head g . The contrastive loss is then computed based on the outputs $z_i = g(h_i)$ and $z_j = g(h_j)$.

We provide implementations of these two parts in `cs231n/simclr/model.py`. Please skim through the file and make sure you understand the implementation.

4 SimCLR: Contrastive Loss

A mini-batch of N training images yields a total of $2N$ data-augmented examples. For each positive pair (i, j) of augmented examples, the contrastive loss function aims to maximize the agreement of vectors z_i and z_j . Specifically, the loss is the normalized temperature-scaled cross entropy loss and aims to maximize the agreement of z_i and z_j relative to all other augmented examples in the batch:

$$l(i, j) = -\log \frac{\exp(\text{sim}(z_i, z_j) / \tau)}{\sum_{k=1}^{2N} \mathbb{1}_{k \neq i} \exp(\text{sim}(z_i, z_k) / \tau)}$$

where $\mathbb{1} \in \{0, 1\}$ is an indicator function that outputs 1 if $k \neq i$ and 0 otherwise. τ is a temperature parameter that determines how fast the exponentials increase.

$\text{sim}(z_i, z_j) = \frac{z_i \cdot z_j}{\|z_i\| \|z_j\|}$ is the (normalized) dot product between vectors z_i and z_j . The higher the similarity between z_i and z_j , the larger the dot product is, and the larger the numerator becomes. The denominator normalizes the value by summing across z_i and all other augmented examples k in the batch. The range of the normalized value is $(0, 1)$, where a high score close to 1 corresponds to a high similarity between the positive pair (i, j) and low similarity between i and other augmented examples k in the batch. The negative log then maps the range $(0, 1)$ to the loss values $(\text{inf}, 0)$.

The total loss is computed across all positive pairs (i, j) in the batch. Let $z = [z_1, z_2, \dots, z_{2N}]$ include all the augmented examples in the batch, where $z_1 \dots z_N$ are outputs of the left branch, and $z_{N+1} \dots z_{2N}$ are outputs of the right branch. Thus, the positive pairs are (z_k, z_{k+N}) for $\forall k \in [1, N]$.

Then, the total loss L is:

$$L = \frac{1}{2N} \sum_{k=1}^N [l(k, k+N) + l(k+N, k)]$$

NOTE: this equation is slightly different from the one in the paper. We've rearranged the ordering of the positive pairs in the batch, so the indices are different. The rearrangement makes it easier to implement the code in vectorized form.

We'll walk through the steps of implementing the loss function in vectorized form. Implement the functions `sim`, `simclr_loss_naive` in `cs231n/simclr/contrastive_loss.py`. Test your code by running the sanity checks below.

```
[ ]: from cs231n.simclr.contrastive_loss import *
      answers = torch.load('simclr_sanity_check.key')
```

```
[ ]: def test_sim(left_vec, right_vec, correct_output):
    output = sim(left_vec, right_vec).cpu().numpy()
    print("Maximum error in sim: %g"%rel_error(correct_output.numpy(), output))

    # Should be less than 1e-07.
    test_sim(answers['left'][0], answers['right'][0], answers['sim'][0])
    test_sim(answers['left'][1], answers['right'][1], answers['sim'][1])
```

Maximum error in sim: 3.81097e-08
Maximum error in sim: 0

```
[ ]: def test_loss_naive(left, right, tau, correct_output):
    naive_loss = simclr_loss_naive(left, right, tau).item()
    print("Maximum error in simclr_loss_naive: %g"%rel_error(correct_output,
    ↪naive_loss))

    # Should be less than 1e-07.
    test_loss_naive(answers['left'], answers['right'], 5.0, answers['loss']['5.0'])
    test_loss_naive(answers['left'], answers['right'], 1.0, answers['loss']['1.0'])
```

Maximum error in simclr_loss_naive: 0
Maximum error in simclr_loss_naive: 1.13123e-07

Now implement the vectorized version by implementing `sim_positive_pairs`, `compute_sim_matrix`, `simclr_loss_vectorized` in `cs231n/simclr/contrastive_loss.py`. Test your code by running the sanity checks below.

```
[ ]: def test_sim_positive_pairs(left, right, correct_output):
    sim_pair = sim_positive_pairs(left, right).cpu().numpy()
    print("Maximum error in sim_positive_pairs: %g"%rel_error(correct_output.
    ↪numpy(), sim_pair))

    # Should be less than 1e-07.
    test_sim_positive_pairs(answers['left'], answers['right'], answers['sim'])
```

Maximum error in sim_positive_pairs: 3.81097e-08

```
[ ]: def test_sim_matrix(left, right, correct_output):
    out = torch.cat([left, right], dim=0)
    sim_matrix = compute_sim_matrix(out).cpu()
    assert torch.isclose(sim_matrix, correct_output).all(), "correct: {}. got:
    ↪{}".format(correct_output, sim_matrix)
    print("Test passed!")

    test_sim_matrix(answers['left'], answers['right'], answers['sim_matrix'])
```

Test passed!

```
[ ]: def test_loss_vectorized(left, right, tau, correct_output):
    vec_loss = simclr_loss_vectorized(left, right, tau, device=left.device).
    ↪item()
    print("Maximum error in loss_vectorized: %g"%rel_error(correct_output,
    ↪vec_loss))

# Should be less than 1e-07.
test_loss_vectorized(answers['left'], answers['right'], 5.0, answers['loss']['5.
    ↪0'])
test_loss_vectorized(answers['left'], answers['right'], 1.0, answers['loss']['1.
    ↪0'])
```

Maximum error in loss_vectorized: 0

Maximum error in loss_vectorized: 5.65617e-08

5 Implement the train function

Complete the `train()` function in `cs231n/simclr/utils.py` to obtain the model's output and use `simclr_loss_vectorized` to compute the loss. (Please take a look at the `Model` class in `cs231n/simclr/model.py` to understand the model pipeline and the returned values)

```
[ ]: from cs231n.simclr.data_utils import *
    from cs231n.simclr.model import *
    from cs231n.simclr.utils import *
```

5.0.1 Train the SimCLR model

Run the following cells to load in the pretrained weights and continue to train a little bit more. This part will take ~10 minutes and will output to `pretrained_model/trained_simclr_model.pth`.

NOTE: Don't worry about logs such as `'[WARN] Cannot find rule for ...'`. These are related to another module used in the notebook. You can verify the integrity of your code changes through our provided prompts and comments.

```
[ ]: # Do not modify this cell.
feature_dim = 128
temperature = 0.5
k = 200
batch_size = 64
epochs = 1
temperature = 0.5
percentage = 0.5
pretrained_path = './pretrained_model/pretrained_simclr_model.pth'

# Prepare the data.
train_transform = compute_train_transform()
```



```

train_data = CIFAR10Pair(root='data', train=True, transform=train_transform,
    ↳download=True)
train_data = torch.utils.data.Subset(train_data, list(np.
    ↳arange(int(len(train_data)*percentage))))
train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True,
    ↳num_workers=16, pin_memory=True, drop_last=True)
test_transform = compute_test_transform()
memory_data = CIFAR10Pair(root='data', train=True, transform=test_transform,
    ↳download=True)
memory_loader = DataLoader(memory_data, batch_size=batch_size, shuffle=False,
    ↳num_workers=16, pin_memory=True)
test_data = CIFAR10Pair(root='data', train=False, transform=test_transform,
    ↳download=True)
test_loader = DataLoader(test_data, batch_size=batch_size, shuffle=False,
    ↳num_workers=16, pin_memory=True)

# Set up the model and optimizer config.
model = Model(feature_dim)
model.load_state_dict(torch.load(pretrained_path, map_location='cpu'),
    ↳strict=False)
model = model.to(device)
flops, params = profile(model, inputs=(torch.randn(1, 3, 32, 32).to(device),))
flops, params = clever_format([flops, params])
print('# Model Params: {} FLOPs: {}'.format(params, flops))
optimizer = optim.Adam(model.parameters(), lr=1e-3, weight_decay=1e-6)
c = len(memory_data.classes)

# Training loop.
results = {'train_loss': [], 'test_acc@1': [], 'test_acc@5': []} #<< -- output

if not os.path.exists('results'):
    os.mkdir('results')
best_acc = 0.0
for epoch in range(1, epochs + 1):
    train_loss = train(model, train_loader, optimizer, epoch, epochs,
    ↳batch_size=batch_size, temperature=temperature, device=device)
    results['train_loss'].append(train_loss)
    test_acc_1, test_acc_5 = test(model, memory_loader, test_loader, epoch,
    ↳epochs, c, k=k, temperature=temperature, device=device)
    results['test_acc@1'].append(test_acc_1)
    results['test_acc@5'].append(test_acc_5)

# Save statistics.
if test_acc_1 > best_acc:
    best_acc = test_acc_1

```

```
torch.save(model.state_dict(), './pretrained_model/trained_simclr_model.  
pth')
```

```
/usr/local/lib/python3.13/dist-packages/torch/utils/data/dataloader.py:424:  
UserWarning: This DataLoader will create 16 worker processes in total. Our  
suggested max number of worker in current system is 2, which is smaller than  
what this DataLoader is going to create. Please be aware that excessive worker  
creation might get DataLoader running slow or even freeze, lower the worker  
number to avoid potential slowness/freeze if necessary.
```

```
self.check_worker_number_rationality()
```

```
[INFO] Register count_convNd() for <class 'torch.nn.modules.conv.Conv2d'>.  
[INFO] Register count_normalization() for <class  
'torch.nn.modules.batchnorm.BatchNorm2d'>.  
[INFO] Register zero_ops() for <class 'torch.nn.modules.activation.ReLU'>.  
[INFO] Register zero_ops() for <class 'torch.nn.modules.container.Sequential'>.  
[INFO] Register count_adap_avgpool() for <class  
'torch.nn.modules.pooling.AdaptiveAvgPool2d'>.  
[INFO] Register count_linear() for <class 'torch.nn.modules.linear.Linear'>.  
[INFO] Register count_normalization() for <class  
'torch.nn.modules.batchnorm.BatchNorm1d'>.  
# Model Params: 24.62M FLOPs: 1.31G
```

```
0%|          | 0/390 [00:00<?, ?it/s]/usr/local/lib/python3.13/dist-  
packages/torch/utils/data/dataloader.py:432: UserWarning: This DataLoader will  
create 16 worker processes in total. Our suggested max number of worker in  
current system is 2, which is smaller than what this DataLoader is going to  
create. Please be aware that excessive worker creation might get DataLoader  
running slow or even freeze, lower the worker number to avoid potential  
slowness/freeze if necessary.
```

```
self.check_worker_number_rationality()
```

```
Train Epoch: [1/1] Loss: 3.2586: 100%|          | 390/390 [02:31<00:00,  
2.57it/s]
```

```
Feature extracting: 100%|          | 782/782 [00:49<00:00, 15.93it/s]
```

```
Test Epoch: [1/1] Acc@1:83.31% Acc@5:99.29%: 100%|          | 157/157  
[00:11<00:00, 13.19it/s]
```

6 Finetune a Linear Layer for Classification!

Now it's time to put the representation vectors to the test!

We remove the projection head from the SimCLR model and slap on a linear layer to finetune for a simple classification task. All layers before the linear layer are frozen, and only the weights in the final linear layer are trained. We compare the performance of the SimCLR + finetuning model against a baseline model, where no self-supervised learning is done beforehand, and all weights in the model are trained. You will get to see for yourself the power of self-supervised learning and how the learned representation vectors improve downstream task performance.

6.1 Baseline: Without Self-Supervised Learning

First, let's take a look at the baseline model. We'll remove the projection head from the SimCLR model and slap on a linear layer to finetune for a simple classification task. No self-supervised learning is done beforehand, and all weights in the model are trained. Run the following cells.

NOTE: Don't worry if you see low but reasonable performance.

```
[ ]: class Classifier(nn.Module):
    def __init__(self, num_class):
        super(Classifier, self).__init__()

        # Encoder.
        self.f = Model().f

        # Classifier.
        self.fc = nn.Linear(2048, num_class, bias=True)

    def forward(self, x):
        x = self.f(x)
        feature = torch.flatten(x, start_dim=1)
        out = self.fc(feature)
        return out

[ ]: # Do not modify this cell.
feature_dim = 128
temperature = 0.5
k = 200
batch_size = 128
epochs = 10
percentage = 0.1

train_transform = compute_train_transform()
train_data = CIFAR10(root='data', train=True, transform=train_transform,
    ↪download=True)
trainset = torch.utils.data.Subset(train_data, list(np.
    ↪arange(int(len(train_data)*percentage))))
train_loader = DataLoader(trainset, batch_size=batch_size, shuffle=True,
    ↪num_workers=16, pin_memory=True)
test_transform = compute_test_transform()
test_data = CIFAR10(root='data', train=False, transform=test_transform,
    ↪download=True)
test_loader = DataLoader(test_data, batch_size=batch_size, shuffle=False,
    ↪num_workers=16, pin_memory=True)

model = Classifier(num_class=len(train_data.classes)).to(device)
for param in model.f.parameters():
    param.requires_grad = False
```

```

flops, params = profile(model, inputs=(torch.randn(1, 3, 32, 32).to(device),))
flops, params = clever_format([flops, params])
print('# Model Params: {} FLOPs: {}'.format(params, flops))
optimizer = optim.Adam(model.fc.parameters(), lr=1e-3, weight_decay=1e-6)
no_pretrain_results = {'train_loss': [], 'train_acc@1': [], 'train_acc@5': [],
                       'test_loss': [], 'test_acc@1': [], 'test_acc@5': []}

best_acc = 0.0
for epoch in range(1, epochs + 1):
    train_loss, train_acc_1, train_acc_5 = train_val(model, train_loader,
    ↪optimizer, epoch, epochs, device='cuda')
    no_pretrain_results['train_loss'].append(train_loss)
    no_pretrain_results['train_acc@1'].append(train_acc_1)
    no_pretrain_results['train_acc@5'].append(train_acc_5)
    test_loss, test_acc_1, test_acc_5 = train_val(model, test_loader, None,
    ↪epoch, epochs)
    no_pretrain_results['test_loss'].append(test_loss)
    no_pretrain_results['test_acc@1'].append(test_acc_1)
    no_pretrain_results['test_acc@5'].append(test_acc_5)
    if test_acc_1 > best_acc:
        best_acc = test_acc_1

# Print the best test accuracy.
print('Best top-1 accuracy without self-supervised learning: ', best_acc)

```

```

[INFO] Register count_convNd() for <class 'torch.nn.modules.conv.Conv2d'>.
[INFO] Register count_normalization() for <class
'torch.nn.modules.batchnorm.BatchNorm2d'>.
[INFO] Register zero_ops() for <class 'torch.nn.modules.activation.ReLU'>.
[INFO] Register zero_ops() for <class 'torch.nn.modules.container.Sequential'>.
[INFO] Register count_adap_avgpool() for <class
'torch.nn.modules.pooling.AdaptiveAvgPool2d'>.
[INFO] Register count_linear() for <class 'torch.nn.modules.linear.Linear'>.
# Model Params: 23.52M FLOPs: 1.31G

```

```

Train Epoch: [1/10] Loss: 2.5539 ACC@1: 10.70% ACC@5: 51.30%: 100%|      |
40/40 [00:10<00:00,  3.67it/s]
Test Epoch: [1/10] Loss: 2.3212 ACC@1: 11.48% ACC@5: 51.60%: 100%|      |
79/79 [00:11<00:00,  6.99it/s]
Train Epoch: [2/10] Loss: 2.4299 ACC@1: 10.88% ACC@5: 51.42%: 100%|      |
40/40 [00:08<00:00,  4.58it/s]
Test Epoch: [2/10] Loss: 2.7025 ACC@1: 10.18% ACC@5: 55.11%: 100%|      |
79/79 [00:12<00:00,  6.37it/s]
Train Epoch: [3/10] Loss: 2.3950 ACC@1: 11.70% ACC@5: 53.12%: 100%|      |
40/40 [00:09<00:00,  4.13it/s]
Test Epoch: [3/10] Loss: 2.5049 ACC@1: 10.24% ACC@5: 53.42%: 100%|      |
79/79 [00:11<00:00,  7.03it/s]

```

```

Train Epoch: [4/10] Loss: 2.4029 ACC@1: 12.44% ACC@5: 54.02%: 100%|      |
40/40 [00:07<00:00,  5.11it/s]
Test Epoch: [4/10] Loss: 2.5870 ACC@1: 10.34% ACC@5: 52.39%: 100%|      |
79/79 [00:11<00:00,  6.96it/s]
Train Epoch: [5/10] Loss: 2.4127 ACC@1: 12.24% ACC@5: 54.48%: 100%|      |
40/40 [00:08<00:00,  4.98it/s]
Test Epoch: [5/10] Loss: 2.7166 ACC@1: 14.82% ACC@5: 54.43%: 100%|      |
79/79 [00:12<00:00,  6.31it/s]
Train Epoch: [6/10] Loss: 2.3939 ACC@1: 12.44% ACC@5: 54.02%: 100%|      |
40/40 [00:09<00:00,  4.11it/s]
Test Epoch: [6/10] Loss: 2.3872 ACC@1: 13.67% ACC@5: 54.45%: 100%|      |
79/79 [00:11<00:00,  6.88it/s]
Train Epoch: [7/10] Loss: 2.3648 ACC@1: 13.10% ACC@5: 54.66%: 100%|      |
40/40 [00:08<00:00,  4.80it/s]
Test Epoch: [7/10] Loss: 2.4616 ACC@1: 11.80% ACC@5: 55.54%: 100%|      |
79/79 [00:12<00:00,  6.44it/s]
Train Epoch: [8/10] Loss: 2.3864 ACC@1: 11.86% ACC@5: 55.12%: 100%|      |
40/40 [00:08<00:00,  4.68it/s]
Test Epoch: [8/10] Loss: 2.4651 ACC@1: 14.32% ACC@5: 59.70%: 100%|      |
79/79 [00:11<00:00,  6.74it/s]
Train Epoch: [9/10] Loss: 2.3793 ACC@1: 13.22% ACC@5: 56.60%: 100%|      |
40/40 [00:09<00:00,  4.27it/s]
Test Epoch: [9/10] Loss: 2.6685 ACC@1: 10.05% ACC@5: 57.28%: 100%|      |
79/79 [00:11<00:00,  6.90it/s]
Train Epoch: [10/10] Loss: 2.4030 ACC@1: 12.96% ACC@5: 57.64%: 100%|      |
40/40 [00:08<00:00,  4.87it/s]
Test Epoch: [10/10] Loss: 2.4337 ACC@1: 15.30% ACC@5: 58.28%: 100%|      |
79/79 [00:12<00:00,  6.36it/s]

Best top-1 accuracy without self-supervised learning: 15.299999999999999

```

6.2 With Self-Supervised Learning

Let's see how much improvement we get with self-supervised learning. Here, we pretrain the SimCLR model using the `simclr` loss you wrote, remove the projection head from the SimCLR model, and use a linear layer to finetune for a simple classification task.

```

[ ]: # Do not modify this cell.
feature_dim = 128
temperature = 0.5
k = 200
batch_size = 128
epochs = 10
percentage = 0.1
pretrained_path = './pretrained_model/trained_simclr_model.pth'

train_transform = compute_train_transform()

```

```

train_data = CIFAR10(root='data', train=True, transform=train_transform,
    ↳download=True)
trainset = torch.utils.data.Subset(train_data, list(np.
    ↳arange(int(len(train_data)*percentage))))
train_loader = DataLoader(trainset, batch_size=batch_size, shuffle=True,
    ↳num_workers=16, pin_memory=True)
test_transform = compute_test_transform()
test_data = CIFAR10(root='data', train=False, transform=test_transform,
    ↳download=True)
test_loader = DataLoader(test_data, batch_size=batch_size, shuffle=False,
    ↳num_workers=16, pin_memory=True)

model = Classifier(num_class=len(train_data.classes))
model.load_state_dict(torch.load(pretrained_path, map_location='cpu'),
    ↳strict=False)
model = model.to(device)
for param in model.parameters():
    param.requires_grad = False

flops, params = profile(model, inputs=(torch.randn(1, 3, 32, 32).to(device),))
flops, params = clever_format([flops, params])
print('# Model Params: {} FLOPs: {}'.format(params, flops))
optimizer = optim.Adam(model.parameters(), lr=1e-3, weight_decay=1e-6)
pretrain_results = {'train_loss': [], 'train_acc@1': [], 'train_acc@5': [],
    'test_loss': [], 'test_acc@1': [], 'test_acc@5': []}

best_acc = 0.0
for epoch in range(1, epochs + 1):
    train_loss, train_acc_1, train_acc_5 = train_val(model, train_loader,
    ↳optimizer, epoch, epochs)
    pretrain_results['train_loss'].append(train_loss)
    pretrain_results['train_acc@1'].append(train_acc_1)
    pretrain_results['train_acc@5'].append(train_acc_5)
    test_loss, test_acc_1, test_acc_5 = train_val(model, test_loader, None,
    ↳epoch, epochs)
    pretrain_results['test_loss'].append(test_loss)
    pretrain_results['test_acc@1'].append(test_acc_1)
    pretrain_results['test_acc@5'].append(test_acc_5)
    if test_acc_1 > best_acc:
        best_acc = test_acc_1

# Print the best test accuracy. You should see a best top-1 accuracy of >=70%.
print('Best top-1 accuracy with self-supervised learning: ', best_acc)

```

```

[INFO] Register count_convNd() for <class 'torch.nn.modules.conv.Conv2d'>.
[INFO] Register count_normalization() for <class
'torch.nn.modules.batchnorm.BatchNorm2d'>.

```

```
[INFO] Register zero_ops() for <class 'torch.nn.modules.activation.ReLU'>.
[INFO] Register zero_ops() for <class 'torch.nn.modules.container.Sequential'>.
[INFO] Register count_adap_avgpool() for <class
'torch.nn.modules.pooling.AdaptiveAvgPool2d'>.
[INFO] Register count_linear() for <class 'torch.nn.modules.linear.Linear'>.
# Model Params: 23.52M FLOPs: 1.31G
```


```
Train Epoch: [1/10] Loss: 1.8206 ACC@1: 64.84% ACC@5: 93.42%: 100%|      |
40/40 [00:07<00:00,  5.06it/s]
Test Epoch: [1/10] Loss: 1.3316 ACC@1: 78.01% ACC@5: 98.24%: 100%|      |
79/79 [00:11<00:00,  6.62it/s]
Train Epoch: [2/10] Loss: 1.1859 ACC@1: 75.94% ACC@5: 97.54%: 100%|      |
40/40 [00:10<00:00,  3.80it/s]
Test Epoch: [2/10] Loss: 0.9373 ACC@1: 79.36% ACC@5: 98.29%: 100%|      |
79/79 [00:11<00:00,  6.98it/s]
Train Epoch: [3/10] Loss: 0.9329 ACC@1: 76.40% ACC@5: 97.96%: 100%|      |
40/40 [00:07<00:00,  5.01it/s]
Test Epoch: [3/10] Loss: 0.7805 ACC@1: 79.80% ACC@5: 98.67%: 100%|      |
79/79 [00:11<00:00,  6.75it/s]
Train Epoch: [4/10] Loss: 0.8399 ACC@1: 76.82% ACC@5: 97.80%: 100%|      |
40/40 [00:08<00:00,  4.90it/s]
Test Epoch: [4/10] Loss: 0.7078 ACC@1: 79.70% ACC@5: 98.67%: 100%|      |
79/79 [00:11<00:00,  7.09it/s]
Train Epoch: [5/10] Loss: 0.7676 ACC@1: 77.72% ACC@5: 97.62%: 100%|      |
40/40 [00:09<00:00,  4.32it/s]
Test Epoch: [5/10] Loss: 0.6460 ACC@1: 81.01% ACC@5: 98.86%: 100%|      |
79/79 [00:11<00:00,  7.08it/s]
Train Epoch: [6/10] Loss: 0.7343 ACC@1: 77.64% ACC@5: 97.84%: 100%|      |
40/40 [00:08<00:00,  4.96it/s]
Test Epoch: [6/10] Loss: 0.6082 ACC@1: 81.50% ACC@5: 98.93%: 100%|      |
79/79 [00:11<00:00,  6.60it/s]
Train Epoch: [7/10] Loss: 0.6961 ACC@1: 77.98% ACC@5: 98.44%: 100%|      |
40/40 [00:08<00:00,  4.45it/s]
Test Epoch: [7/10] Loss: 0.5882 ACC@1: 81.71% ACC@5: 98.96%: 100%|      |
79/79 [00:11<00:00,  7.05it/s]
Train Epoch: [8/10] Loss: 0.6784 ACC@1: 78.30% ACC@5: 98.26%: 100%|      |
40/40 [00:08<00:00,  4.88it/s]
Test Epoch: [8/10] Loss: 0.5661 ACC@1: 82.15% ACC@5: 99.01%: 100%|      |
79/79 [00:11<00:00,  7.13it/s]
Train Epoch: [9/10] Loss: 0.6685 ACC@1: 78.76% ACC@5: 98.32%: 100%|      |
40/40 [00:07<00:00,  5.06it/s]
Test Epoch: [9/10] Loss: 0.5551 ACC@1: 82.23% ACC@5: 98.99%: 100%|      |
79/79 [00:11<00:00,  6.66it/s]
Train Epoch: [10/10] Loss: 0.6445 ACC@1: 79.06% ACC@5: 98.18%: 100%|      |
40/40 [00:09<00:00,  4.09it/s]
Test Epoch: [10/10] Loss: 0.5403 ACC@1: 82.37% ACC@5: 99.01%: 100%|      |
79/79 [00:11<00:00,  7.05it/s]
```

Best top-1 accuracy with self-supervised learning: 82.37

6.2.1 Plot your Comparison

Plot the test accuracies between the baseline model (no pretraining) and same model pretrained with self-supervised learning.

```
[ ]: plt.plot(no_pretrain_results['test_acc@1'], label="Without Pretrain")
plt.plot(pretrain_results['test_acc@1'], label="With Pretrain")
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Test Top-1 Accuracy')
plt.legend()
plt.show()
```



```
[ ]:
```


DDPM

September 10, 2026

```
[2]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment3/'
FOLDERNAME = "cs231n/assignments/assignment3/"
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# # This downloads the Emoji dataset to your Drive
# # if it doesn't already exist.
# %cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
# !bash get_datasets.sh
# %cd /content/drive/My\ Drive/$FOLDERNAME
```

Mounted at /content/drive

1 Denoising Diffusion Probabilistic Models

So far, we have explored discriminative models, which are trained to produce a labeled output. These range from straightforward image classification to sentence generation where the problem was still framed as classification, with labels in vocabulary space and a recurrence mechanism capturing multi-word labels. Now, we will expand our repertoire by building a generative model capable of creating novel images that resemble a given set of training images.

There are many types of generative models, including Generative Adversarial Networks (GANs), autoregressive models, normalizing flow models, and Variational Autoencoders (VAEs), all of which can synthesize striking images. However, in 2020, Ho et al. introduced Denoising Diffusion Probabilistic Models (DDPMs) by combining diffusion probabilistic models with denoising score matching. This resulted in a generative model that is both simple to train and powerful enough to generate complex, high-quality images. The following provides a high-level overview of DDPMs. For more details, please refer to the course slides and the original DDPM paper [1].

2 Forward Process

Let $q(x_0)$ be the distribution of clean dataset images. We define a forward noising process as a Markov chain of small noising steps:

$$q(x_t|x_{t-1}) \sim N(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$

where the stepwise variance $(\beta_1, \dots, \beta_T)$ determines the noise schedule. Due to the properties of Gaussian distributions, we can express $q(x_t|x_0)$ in closed form as:

$$q(x_t|x_0) \sim N(x_t; \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I)$$

where $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$. If the noise schedule $(\beta_1, \dots, \beta_T)$ is set properly, the final distribution $q(x_T)$ becomes indistinguishable from pure Gaussian noise $N(0, I)$.

Recall that sampling from a Gaussian distribution $x \sim N(\mu, \sigma^2)$ is equivalent to computing $\sigma * \epsilon + \mu$ where $\epsilon \sim N(0, 1)$. Hence, sampling from $q(x_t|x_{t-1})$ or $q(x_t|x_0)$ is straightforward given x_{t-1} or x_0 respectively. Because of this, the forward process is simple and does not require learning.

3 Reverse Process

The reverse process reconstructs a clean image x_0 from pure noise x_T through multiple steps. Let $p(x_{t-1}|x_t)$ denote the reverse step of $q(x_t|x_{t-1})$. The first key insight is that learning to reverse each individual denoising step is easier than reversing the entire forward process in one go. In other words, learning $p(x_{t-1}|x_t)$ for each t is easier than directly learning $p(x_0|x_T)$.

However, learning $p(x_{t-1}|x_t)$ is still challenging. Although $q(x_t|x_{t-1})$ is Gaussian, $p(x_{t-1}|x_t)$ could take any complex form and is almost certainly not Gaussian. Modeling and sampling from an arbitrary distribution is significantly harder than working with a simple parametric distribution like a Gaussian.

The second key insight is that if the stepwise noise β_t in the forward process is small enough, then the reverse step $p(x_{t-1}|x_t)$ is also close to a Gaussian distribution. Thus, we only need to estimate its mean and variance. In practice, setting the variance of $p(x_{t-1}|x_t)$ to match β_t (the same as in the forward step) works well. Consequently, learning the reverse process reduces to learning the mean $\mu(x_t, t; \theta)$ where θ represents the parameters of a neural network.

4 Denoising Objective

Generative models are optimized by minimizing the expected negative log-likelihood $\mathbb{E}[-\log p_\theta(x_0)]$ of the dataset samples. The likelihood of each sample can be expressed as: $p_\theta(x_0) = p(x_T) \prod_{t=1}^T p(x_{t-1}|x_t)$. Since this objective is intractable in many cases, various classes of generative models optimize the variational lower bound on the negative log-likelihood.

Ho et al. demonstrated that this objective is equivalent to minimizing the following denoising loss

$$\mathbb{E}_{t, x_0, \epsilon} [\|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2]$$

where t is uniform between 1 and T , x_0 is clean sample, ϵ is sampled from standard gaussian $N(0, I)$, and ϵ_θ is a neural network model trained to predict the noise ϵ from the input noisy sample $x_t = \sqrt{\alpha_t}x_0 + \sqrt{1 - \alpha_t}\epsilon$. In other words, ϵ_θ learns to denoise the input noisy image. Note that this is equivalent to predicting the clean sample, since the noise can be recovered from the noisy image and the clean sample following the equation $x_t = \sqrt{\alpha_t}x_0 + \sqrt{1 - \alpha_t}\epsilon$.

[1] Denoising Diffusion Probabilistic Models. Jonathan Ho, Ajay Jain, Pieter Abbeel. [Link](#)

5 In This Notebook...

We will implement and train a DDPM model to generate small 32 x 32 emoji images conditioned on text prompts. First, we will implement the forward noising process based on Eq. (4) of the paper [1]. Then we will build a UNet model that takes x_t and t as inputs (optionally with other conditioning like text-prompt) and outputs a tensor of the same shape as x_t . Finally, we will implement the denoising objective and train our DDPM model.

We use the text encoder from a pretrained CLIP[2] model to encode input text into a 512-dimensional vector. To speed up training, we've already pre-encoded the text data from the training set.

[2] Learning transferable visual models from natural language supervision. Radford et. al. [Link](#)

[3]: `!pip install git+https://github.com/openai/CLIP.git`

```
Collecting git+https://github.com/openai/CLIP.git
  Cloning https://github.com/openai/CLIP.git to /tmp/pip-req-build-hnqnqcl
  Running command git clone --filter=blob:none --quiet
https://github.com/openai/CLIP.git /tmp/pip-req-build-hnqnqcl
  Resolved https://github.com/openai/CLIP.git to commit
d05afc436d78f1c48dc0dbf8e5980a9d471f35f6
  Preparing metadata (setup.py) ... done
Collecting ftfy (from clip==1.0)
  Downloading ftfy-6.3.1-py3-none-any.whl.metadata (7.3 kB)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-
packages (from clip==1.0) (24.2)
Requirement already satisfied: regex in /usr/local/lib/python3.11/dist-packages
(from clip==1.0) (2024.11.6)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages
(from clip==1.0) (4.67.1)
Requirement already satisfied: torch in /usr/local/lib/python3.11/dist-packages
(from clip==1.0) (2.6.0+cu124)
Requirement already satisfied: torchvision in /usr/local/lib/python3.11/dist-
packages (from clip==1.0) (0.21.0+cu124)
Requirement already satisfied: wcwidth in /usr/local/lib/python3.11/dist-
packages (from ftfy->clip==1.0) (0.2.13)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-
packages (from torch->clip==1.0) (3.18.0)
Requirement already satisfied: typing-extensions>=4.10.0 in
/usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (4.14.1)
```

Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (3.5)

Requirement already satisfied: jinja2 in /usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (3.1.6)

Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (2025.3.2)

Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch->clip==1.0)
 Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)

Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch->clip==1.0)
 Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)

Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch->clip==1.0)
 Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)

Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch->clip==1.0)
 Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)

Collecting nvidia-cublas-cu12==12.4.5.8 (from torch->clip==1.0)
 Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)

Collecting nvidia-cufft-cu12==11.2.1.3 (from torch->clip==1.0)
 Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)

Collecting nvidia-curand-cu12==10.3.5.147 (from torch->clip==1.0)
 Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)

Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch->clip==1.0)
 Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)

Collecting nvidia-cusparselt-cu12==0.6.2 in /usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (0.6.2)

Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (2.21.5)

Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (12.4.127)

Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch->clip==1.0)
 Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)

Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (3.2.0)

Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (1.13.1)

Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch->clip==1.0)

```

(1.3.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages
(from torchvision->clip==1.0) (2.0.2)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in
/usr/local/lib/python3.11/dist-packages (from torchvision->clip==1.0) (11.2.1)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.11/dist-packages (from jinja2->torch->clip==1.0) (3.0.2)
Downloading ftfy-6.3.1-py3-none-any.whl (44 kB)
44.8/44.8 kB
4.8 MB/s eta 0:00:00
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl
(363.4 MB)
363.4/363.4 MB
5.2 MB/s eta 0:00:00
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (13.8 MB)
13.8/13.8 MB
105.1 MB/s eta 0:00:00
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (24.6 MB)
24.6/24.6 MB
70.8 MB/s eta 0:00:00
Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (883 kB)
883.7/883.7 kB
65.8 MB/s eta 0:00:00
Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl
(664.8 MB)
664.8/664.8 MB
3.2 MB/s eta 0:00:00
Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl
(211.5 MB)
211.5/211.5 MB
6.9 MB/s eta 0:00:00
Downloading nvidia_curand_cu12-10.3.5.147-py3-none-
manylinux2014_x86_64.whl (56.3 MB)
56.3/56.3 MB
19.4 MB/s eta 0:00:00
Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-
manylinux2014_x86_64.whl (127.9 MB)
127.9/127.9 MB
10.7 MB/s eta 0:00:00
Downloading nvidia_cusparsparse_cu12-12.3.1.170-py3-none-
manylinux2014_x86_64.whl (207.5 MB)
207.5/207.5 MB
7.5 MB/s eta 0:00:00
Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (21.1 MB)

```

21.1/21.1 MB

80.6 MB/s eta 0:00:00

Building wheels for collected packages: clip

Building wheel for clip (setup.py) ... done

Created wheel for clip: filename=clip-1.0-py3-none-any.whl size=1369490
sha256=45c949500712c0405322a44abac512e7ab54f04ac606e89300a65111684cc4d8

Stored in directory: /tmp/pip-ephem-wheel-cache-
65ye7l9m/wheels/3f/7c/a4/9b490845988bf7a4db33674d52f709f088f64392063872eb9a
Successfully built clip

Installing collected packages: nvidia-nvjitlink-cu12, nvidia-curand-cu12,
nvidia-cufft-cu12, nvidia-cuda-runtime-cu12, nvidia-cuda-nvrtc-cu12, nvidia-
cuda-cupti-cu12, nvidia-cublas-cu12, ftfy, nvidia-cuspars-cu12, nvidia-cudnn-
cu12, nvidia-cusolver-cu12, clip

Attempting uninstall: nvidia-nvjitlink-cu12

Found existing installation: nvidia-nvjitlink-cu12 12.5.82

Uninstalling nvidia-nvjitlink-cu12-12.5.82:

Successfully uninstalled nvidia-nvjitlink-cu12-12.5.82

Attempting uninstall: nvidia-curand-cu12

Found existing installation: nvidia-curand-cu12 10.3.6.82

Uninstalling nvidia-curand-cu12-10.3.6.82:

Successfully uninstalled nvidia-curand-cu12-10.3.6.82

Attempting uninstall: nvidia-cufft-cu12

Found existing installation: nvidia-cufft-cu12 11.2.3.61

Uninstalling nvidia-cufft-cu12-11.2.3.61:

Successfully uninstalled nvidia-cufft-cu12-11.2.3.61

Attempting uninstall: nvidia-cuda-runtime-cu12

Found existing installation: nvidia-cuda-runtime-cu12 12.5.82

Uninstalling nvidia-cuda-runtime-cu12-12.5.82:

Successfully uninstalled nvidia-cuda-runtime-cu12-12.5.82

Attempting uninstall: nvidia-cuda-nvrtc-cu12

Found existing installation: nvidia-cuda-nvrtc-cu12 12.5.82

Uninstalling nvidia-cuda-nvrtc-cu12-12.5.82:

Successfully uninstalled nvidia-cuda-nvrtc-cu12-12.5.82

Attempting uninstall: nvidia-cuda-cupti-cu12

Found existing installation: nvidia-cuda-cupti-cu12 12.5.82

Uninstalling nvidia-cuda-cupti-cu12-12.5.82:

Successfully uninstalled nvidia-cuda-cupti-cu12-12.5.82

Attempting uninstall: nvidia-cublas-cu12

Found existing installation: nvidia-cublas-cu12 12.5.3.2

Uninstalling nvidia-cublas-cu12-12.5.3.2:

Successfully uninstalled nvidia-cublas-cu12-12.5.3.2

Attempting uninstall: nvidia-cuspars-cu12

Found existing installation: nvidia-cuspars-cu12 12.5.1.3

Uninstalling nvidia-cuspars-cu12-12.5.1.3:

Successfully uninstalled nvidia-cuspars-cu12-12.5.1.3

Attempting uninstall: nvidia-cudnn-cu12

Found existing installation: nvidia-cudnn-cu12 9.3.0.75

Uninstalling nvidia-cudnn-cu12-9.3.0.75:

```

    Successfully uninstalled nvidia-cudnn-cu12-9.3.0.75
Attempting uninstall: nvidia-cusolver-cu12
    Found existing installation: nvidia-cusolver-cu12 11.6.3.83
    Uninstalling nvidia-cusolver-cu12-11.6.3.83:
        Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83
Successfully installed clip-1.0 ftfy-6.3.1 nvidia-cublas-cu12-12.4.5.8 nvidia-
cuda-cupti-cu12-12.4.127 nvidia-cuda-nvrtc-cu12-12.4.127 nvidia-cuda-runtime-
cu12-12.4.127 nvidia-cudnn-cu12-9.1.0.70 nvidia-cufft-cu12-11.2.1.3 nvidia-
curand-cu12-10.3.5.147 nvidia-cusolver-cu12-11.6.1.9 nvidia-cusparse-
cu12-12.3.1.170 nvidia-nvjitlink-cu12-12.4.127

```

```

[4]: import sys
import types
import importlib

if "imp" not in sys.modules:
    imp = types.ModuleType("imp")
    imp.reload = importlib.reload
    sys.modules["imp"] = imp

%load_ext autoreload
%autoreload 2

import numpy as np
import torch
import random
import matplotlib.pyplot as plt
import torchvision.utils as tv_utils
from cs231n.emoji_dataset import EmojiDataset
from cs231n.gaussian_diffusion import GaussianDiffusion

def rel_error(x, y):
    """Returns relative error."""
    return np.max(np.abs(x - y) / (np.maximum(1e-10, np.abs(x) + np.abs(y))))

```

```

[5]: # First, let's load and visualize the dataset.
# Each sample of the dataset is a tuple: (image, {"text_emb": <tensor>, "text":
↳<string>})
# We will use a pretrained text-encoder to encode text into an embedding vector.
# We have pre-encoded the dataset texts into embeddings for faster training.
image_size = 32
dataset = EmojiDataset(image_size)

```

emoji_data.npz already downloaded.
text_embeddings.pt already downloaded.

```
[6]: def visualize_samples(dataset, num_samples=25, grid_size=(5, 5)):
    # Randomly sample indices
    indices = random.sample(range(len(dataset)), num_samples)
    samples = [dataset[i] for i in indices]

    # Inspect one sample
    img_shape = list(samples[0][0].shape)
    emb_shape = list(samples[0][1]["text_emb"].shape)
    print(f"One sample: (image: {img_shape}, {{ \"text_emb\": {emb_shape},
↪ \"text\": string }})")

    # Extract images and texts
    images = torch.stack([sample[0] for sample in samples]) # Stack images
    texts = [sample[1]["text"] for sample in samples] # Extract text
↪ descriptions

    # Create a grid of images
    grid_img = tv_utils.make_grid(images, nrow=grid_size[1], padding=2)

    # Convert to numpy for plotting
    grid_img = grid_img.permute(1, 2, 0).numpy()

    # Plot the images
    fig, ax = plt.subplots(figsize=(10, 10))
    ax.imshow(grid_img)
    ax.axis("off")

    # Add text annotations
    grid_w, grid_h = grid_size
    img_w, img_h = grid_img.shape[1] // grid_w, grid_img.shape[0] // grid_h

    for i, text in enumerate(texts):
        row, col = divmod(i, grid_w)
        x, y = col * img_w, row * img_h
        ax.text(x+5, y+5, text[:30], fontsize=8, color='white',
↪ bbox=dict(facecolor='black', alpha=0.5))

    plt.show()

visualize_samples(dataset)
```

One sample: (image: [3, 32, 32], { "text_emb": [512], "text": string })



5.1 q_sample

Now we will define the forward noising process. Read through the `GaussianDiffusion` class in `cs231n/gaussian_diffusion.py`. Consult the original DDPM paper[1] for the equations. Implement `q_sample` method and test it below. You should see zero relative error.

```
[7]: # Test GaussianDiffusion.q_sample method
sz = 2
b = 3

diffusion = GaussianDiffusion(
    model=None,
    image_size=sz,
```

```

        timesteps=1000,
        beta_schedule="sigmoid",
    )

t = torch.tensor([0, 300, 999]).long()
x_start = torch.linspace(-0.9, 0.6, b*3*sz*sz).view(b, 3, sz, sz)
noise = torch.linspace(-0.7, 0.8, b*3*sz*sz).view(b, 3, sz, sz)
x_t = diffusion.q_sample(x_start, t, noise)

expected_x_t = np.array([
    [
        [[-0.9119949, -0.86840147], [-0.8248081, -0.7812148]],
        [[-0.7376214, -0.694028], [-0.65043473, -0.6068413]],
        [[-0.563248, -0.51965463], [-0.47606122, -0.43246788]],
    ],
    [
        [[-0.42800453, -0.37039882], [-0.31279305, -0.2551873]],
        [[-0.19758154, -0.1399758], [-0.08237009, -0.024764337]],
        [[0.032841414, 0.090447165], [0.14805292, 0.20565866]],
    ],
    [
        [[0.32864183, 0.37152246], [0.41440308, 0.45728368]],
        [[0.50016433, 0.5430449], [0.5859255, 0.6288062]],
        [[0.67168677, 0.7145674], [0.757448, 0.8003287]],
    ],
]).astype(np.float32)

# Should see zero relative error
print("x_t error: ", rel_error(x_t.numpy(), expected_x_t))

```

x_t error: 0.0

```

[8]: # Let's visualize the noisy images at various timesteps.
diffusion = GaussianDiffusion(
    model=None,
    image_size=image_size,
    timesteps=1000,
)

B = 10
img = dataset[770][0] # 3 x H x W
x_start = img[None].repeat(B, 1, 1, 1) # B x 3 x H x W
noise = torch.randn_like(x_start) # B x 3 x H x W
t = torch.linspace(0, 1000-1, B).long()


x_start = diffusion.normalize(x_start)
x_t = diffusion.q_sample(x_start, t, noise)

```

```

x_t = diffusion.unnormalize(x_t).clamp(0, 1)
grid_img = tv_utils.make_grid(x_t, nrow=5, padding=2)
grid_img = grid_img.permute(1, 2, 0).cpu().numpy()
fig, ax = plt.subplots(figsize=(10, 10))
ax.imshow(grid_img)
ax.axis("off")
plt.show()

```



A diffusion model can be trained to predict either the clean image or the noise, as one can be derived from the other (explained in ‘Denoising Objective’ section above). Implement `predict_start_from_noise` and `predict_noise_from_start` methods and test them below. You should see relative error less than $1e-5$.

```

[9]: # Test `predict_noise_from_start` and `predict_start_from_noise`
sz = 2
b = 3

diffusion = GaussianDiffusion(
    model=None,
    image_size=sz,
    timesteps=1000,
    beta_schedule="sigmoid",
)

t = torch.tensor([1, 300, 998]).long()
x_start = torch.linspace(-0.91, 0.67, b*3*sz*sz).view(b, 3, sz, sz)
noise = torch.linspace(-0.73, 0.81, b*3*sz*sz).view(b, 3, sz, sz)
x_t = diffusion.q_sample(x_start, t, noise)

pred_noise = diffusion.predict_noise_from_start(x_t, t, x_start)
pred_x_start = diffusion.predict_start_from_noise(x_t, t, noise)

```

```
# Should relative errors around 1e-5 or less
print("noise error: ", rel_error(pred_noise.numpy(), noise.numpy()))
print("x_start error: ", rel_error(pred_x_start.numpy(), x_start.numpy()))
```

```
noise error: 1.0600407e-06
x_start error: 1.8902663e-06
```

5.2 UNet Model


Now that we have defined the forward process, let's define the UNet model for denoising the input image. UNet is a neural network architecture designed for image-to-image tasks like segmentation, style transfer, etc. It consists of an encoder (or downsampling module) that transforms the input image into hierarchical features with decreasing spatial resolution and increasing feature dimensions. The decoder (or upsampling module) then upsamples the features by progressively restoring the spatial resolution, mirroring the encoder's structure. At each decoder layer, features from the corresponding encoder layer are concatenated, providing a direct pathway for finer details. This approach reduces the burden on the bottleneck layers, allowing them to focus on capturing high-level representations rather than memorizing fine details.

We will use UNet in this case because both our input and output are aligned images with same dimensions: C x H x W. Each ResNet block in the UNet will also take an additional input vector called context for conditioning. We will generate the context vector by encoding the diffusion timestep and text-prompt.

Run the cell below to get a rough outline of the UNet architecture. Each red box represents a ResNet block containing 2 or 3 convolutional layers that maintain the spatial resolution of the feature maps. The context vector input to every ResNet block is omitted for clarity. The shape written below each box indicates the output tensor shape after that block. Additional arrows illustrate the skip connections, which enable the U-Net to preserve fine-grained details in the output. For example, the output of `layer1_block1` with shape (d, h, w) will be concatenated with the output of `layer4_block1`, also with shape (d, h, w), before being passed to `layer4_block2`. Therefore, `layer4_block2` will receive an input of shape (2*d, h, w).

```
[10]: from IPython.display import Image
      Image(f'/content/drive/My Drive/{FOLDERNAME}/unet.png')
```

[10]:



Implement the `Unet.__init__` method in `cs231n/unet.py` to define the upsampling and downsampling blocks of the UNet model, and then test it below. If your implementation is correct, you should not see any error. Calling `Unet(dim=d, condition_dim=condition_dim, dim_mults=(2,4))` should successfully create a UNet model corresponding to the architecture shown in the figure above.

```
[11]: from cs231n.unet import Unet, ResnetBlock, Downsample, Upsample

dim = 2
condition_dim = 4
dim_mults = (1, 2, 4)
unet = Unet(dim=dim, condition_dim=condition_dim, dim_mults=dim_mults)

# Check number of layers
assert len(unet.downs) == len(dim_mults), "Number of Unet downsampling blocks is wrong."
assert len(unet.ups) == len(dim_mults), "Number of Unet upsampling blocks is wrong."

# Check layers
try:
    expected_downs_dims = [2, 2, 8, 2, 2, 8, 2, 2, 8, 2, 2, 8, 4, 4, 8, 4, 4, 8]
    downs_dims = [
        d for m in unet.downs for d in [
```

```

        m[0].dim, m[0].dim_out, m[0].context_dim, m[1].dim, m[1].dim_out,
        ↪m[1].context_dim,
    ]
]
assert downs_dims == expected_downs_dims, "Dimensions don't match"
except Exception as e:
    raise RuntimeError("Downsampling blocks wrongly configured") from e

try:
    expected_ups_dims = [8, 4, 8, 8, 4, 8, 4, 2, 8, 4, 2, 8, 4, 2, 8, 4, 2, 8]
    ups_dims = [
        d for m in unet.ups for d in [
            m[1].dim, m[1].dim_out, m[1].context_dim, m[2].dim, m[2].dim_out,
            ↪m[2].context_dim,
        ]
    ]
    assert ups_dims == expected_ups_dims, "Dimensions don't match"
except Exception as e:
    raise RuntimeError("Upsampling blocks wrongly configured") from e

# Check number of parameters
num_params = sum(p.numel() for p in unet.parameters())
expected_num_params = 6499
assert num_params == expected_num_params, "Unet model creation is wrong!"

```

Fill in `Unet.forward` method in `cs231n/unet.py` and test it below. For now, don't worry about `Unet.cfg_forward` method. You should see relative error less than $1e-6$.

```

[12]: np.random.seed(231)
      torch.manual_seed(231)

      dim = 4
      condition_dim = 4
      dim_mults = (2, 4)
      unet = Unet(dim=dim, condition_dim=condition_dim, dim_mults=dim_mults)
      unet.eval()

      b = 2
      h = w = 4
      inp_x = torch.randn(b, 3, h, w)
      inp_text_emb = torch.randn(b, condition_dim)
      inp_t = torch.tensor([8, 25]).long()
      out = unet.forward(x=inp_x, time=inp_t,
                        ↪model_kwargs={"text_emb": inp_text_emb}).detach().numpy()

      expected_out = np.array(

```

```

[[[ 0.3530089, 0.49302575, 0.55152094, 0.20567769],
 [ 1.1295223, 0.05811183, 0.2143015, 0.23898259],
 [ 0.44839963, 1.0291728, 0.18519637, 0.21956968],
 [ 0.7035912, 0.716347, 0.9166326, -0.11552593]],
 [[ 0.04585427, -0.26847702, 0.08526254, 0.2396591 ],
 [ 0.2858743, -0.29726404, 0.02462834, 0.2350314 ],
 [ 0.3121434, 0.3962962, 0.65541625, -0.300183 ],
 [ 0.3409412, 0.35369393, 0.37445623, 0.10297439]],
 [[ 0.49440542, 0.72878885, 0.6669705, 0.56822944],
 [ 1.0084232, 0.19204213, 0.61240125, 0.575651 ],
 [ 0.6452671, 0.91168964, 0.6458106, 0.25952405],
 [ 0.7147298, 0.7707498, 0.88941365, 0.35422024]]],
 [[ 0.3186316, -0.17632714, 0.440247, 0.05124736],
 [ 0.6866561, 0.3019496, 0.18366373, 0.42941707],
 [-0.09776017, 0.55849403, 0.30689716, 0.12954405],
 [ 0.3667726, 0.740231, 0.43682802, 0.08275545]],
 [[-0.29714173, -0.28515863, -0.02289355, 0.05637485],
 [-0.10230345, 0.182563, 0.8107456, 0.18209176],
 [ 0.57062995, 0.07061654, 0.2864037, 0.17969263],
 [ 0.31491554, 0.5652035, 0.3810383, 0.05889529]],
 [[ 0.14559639, 0.36408228, 0.5657894, 0.24066216],
 [ 0.4431491, 0.38140664, 1.1457629, 0.40314198],
 [ 0.6673274, 0.5922724, 0.44549578, 0.5544277 ],
 [ 0.5405652, 0.88729244, 0.66340685, 0.23414826]]]])

print("forward error: ", rel_error(out, expected_out))

```

forward error: 5.831791042516234e-06

6 p_losses

Now that we have model implementation done, let's write the DDPM's denoising training step. As mentioned before, optimizing the denoising loss is equivalent to minimizing the expected negative log likelihood of the dataset. Complete the `GaussianDiffusion.p_losses` method in `cs231n/gaussian_diffusion.py` and test it below. You should see relative error less than $1e-6$.

```

[13]: np.random.seed(231)
      torch.manual_seed(231)

      dim = 4
      condition_dim = 4
      dim_mults = (2, 4)
      unet = Unet(dim=dim, condition_dim=condition_dim, dim_mults=dim_mults)
      unet.eval()

      h = w = 4
      b = 3

```

```

diffusion = GaussianDiffusion(
    model=unet,
    image_size=h,
    timesteps=1000,
    beta_schedule="sigmoid",
    objective="pred_x_start",
)

inp_x = torch.randn(b, 3, h, w)
inp_model_kwargs = {"text_emb": torch.randn(b, condition_dim)}
out = diffusion.p_losses(inp_x, inp_model_kwargs)
expected_out = 30.0160

print("forward error: ", rel_error(out.item(), expected_out))

```

forward error: 7.432147198281673e-07

6.1 p_sample

There is one final ingredient remaining now. DDPM generates samples by iteratively performing the reverse process. Each iteration of this reverse process involves sampling from $p(x_{t-1}|x_t)$. Open `cs231n/gaussian_diffusion.py` and implement `p_sample` method by following Equation (6) from the paper. This equation describes sampling from the posterior of the forward process, conditioned on x_t and x_0 , where x_0 can be derived from the denoising model's output. We have already implemented `sample` method that iteratively calls `p_sample` to generate images from input texts.

Test your implementation of `p_sample` below, you should see relative errors less than $1e-6$.

```

[14]: np.random.seed(231)
      torch.manual_seed(231)

      dim = 4
      condition_dim = 4
      dim_mults = (2,)
      unet = Unet(dim=dim, condition_dim=condition_dim, dim_mults=dim_mults)
      unet.eval()

      h = w = 4
      b = 1
      inp_x_t = torch.randn(b, 3, h, w)
      inp_model_kwargs = {"text_emb": torch.randn(b, condition_dim)}
      t = 231

      # test 1
      diffusion = GaussianDiffusion(
          model=unet,
          image_size=h,
          timesteps=1000,

```



```

        beta_schedule="sigmoid",
        objective="pred_x_start",
    )
    out = diffusion.p_sample(inp_x_t, t, inp_model_kwargs).detach().numpy()

    expected_out = np.array(
        [[[[ 1.1249855,  0.098786, -0.674154,  1.3266845 ],
            [-0.21832949,  0.4874724,  1.0712924, -0.7355726 ],
            [-0.19198017,  1.0347139,  0.03363481,  0.47500697],
            [-1.0178545, -0.52259684, -0.11972852, -0.02105119]],
          [[ 0.06438226, -1.5796007,  0.43629852, -0.67392266],
            [-1.1564192, -1.7009112,  0.878653, -1.2563533 ],
            [-0.45194352,  1.0274173,  0.62859684, -0.49643248],
            [-0.18070872,  0.51325583, -0.74465716, -0.5262963 ]],
          [[ 0.76330817,  1.5878401,  1.9983996,  0.6618041 ],
            [ 0.16782297,  0.5638586, -0.8150751,  1.9507835 ],
            [ 0.7401889,  1.2696184,  1.8468435, -1.4796975 ],
            [ 0.6636643, -0.83933365,  0.78857577,  0.13521622]]]])

    print("forward error: ", rel_error(out, expected_out))

# test 2
diffusion = GaussianDiffusion(
    model=unet,
    image_size=h,
    timesteps=1000,
    beta_schedule="cosine",
    objective="pred_noise",
)
out = diffusion.p_sample(inp_x_t, t, inp_model_kwargs).detach().numpy()

expected_out = np.array(
    [[[[ 1.1139543,  0.11015256, -0.7337392,  1.3292073 ],
        [-0.286727,  0.48420003,  1.0737748, -0.764173 ],
        [-0.22848499,  1.0443672,  0.07089197,  0.50183666],
        [-1.0453997, -0.5474187, -0.06630269, -0.11730157]],
      [[ 0.10919069, -1.5161378,  0.38542387, -0.6693473 ],
        [-1.0730051, -1.6837747,  0.87086374, -1.2265388 ],
        [-0.48162907,  1.0328836,  0.6336425, -0.47346604],
        [-0.11769794,  0.5335539, -0.71324545, -0.47889465]],
      [[ 0.8311781,  1.6457081,  1.9775101,  0.563749 ],
        [ 0.25127694,  0.57862943, -0.8148284,  1.8744429 ],
        [ 0.75546896,  1.1215659,  1.9582558, -1.535511 ],
        [ 0.62724066, -0.8814953,  0.77542645,  0.12667023]]]])

    print("forward error: ", rel_error(out, expected_out))

```

forward error: 6.96344541990084e-08

forward error: 6.096590549306921e-08

6.2 Training

We have all ingredients needed for DDPM training and we can train the model on our Emoji dataset. You don't have to code anything here but we encourage you to look at the training code at `cs231n/ddpm_trainer.py`.

For the rest of the notebook, we will use a pretrained model from `cs231n/exp/pretrained` folder which is already trained for many iterations on this dataset. However, you are free to train your own model on colab GPU (make sure to change the `results_folder`). Note that it may take more than 12 hours on T4 GPU before you start seeing a reasonable generation.

```
[15]: from cs231n.ddpm_trainer import Trainer

dim = 48
image_size = 32
results_folder = f"/content/drive/My Drive/{FOLDERNAME}/cs231n/exp/pretrained"
condition_dim = 512

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

model = Unet(
    dim=dim,
    dim_mults=(1, 2, 4, 8),
    condition_dim=condition_dim,
)
print("Number of parameters:", sum(p.numel() for p in model.parameters()))

diffusion = GaussianDiffusion(
    model,
    image_size=image_size,
    timesteps=100, # number of diffusion steps
    objective="pred_noise", # "pred_x_start" or "pred_noise"
)

dataset = EmojiDataset(image_size)

trainer = Trainer(
    diffusion,
    dataset,
    device,
    train_batch_size=256,
    weight_decay=0.0,
    train_lr=1e-3,
    train_num_steps=50000,
    results_folder=results_folder,
)

# You are not required to train it yourself.
```

```
# trainer.train()
```

Number of parameters: 12444963
emoji_data.npz already downloaded.
text_embeddings.pt already downloaded.

```
[16]: # Instead, we will load a pretrained model.  
trainer.download_pretrained()  
trainer.load(70000)
```

Pretrained model already downloaded.
loading model from /content/drive/My
Drive/cs231n/assignments/assignment3//cs231n/exp/pretrained/model-70000.pt.

```
[17]: # Helper function to get CLIP text embedding during inference time.  
from cs231n.emoji_dataset import ClipEmbed  
clip_embedder = ClipEmbed(device)  
def get_text_emb(text):  
    return trainer.ds.embed_new_text(text, clip_embedder)  
  
# Helper function to visualize generations.  
def show_images(img):  
    # img: B x T x 3 x H x W  
    plt.figure(figsize=(10, 10))  
    img2 = img.clamp(0, 1).permute(0, 3, 1, 4, 2).flatten(0, 1).flatten(1, 2).  
    ↪cpu().numpy()  
    plt.imshow(img2)  
    plt.axis('off')  
  
    plt.show()
```

100% | 338M/338M [00:02<00:00, 131MiB/s]

6.3 Sampling

Run the cell below to visualize emoji generations conditioned on a text prompt. Feel free to modify the prompt to explore different generations. Since our emoji dataset is quite small, it is insufficient to train a fully generalizable text-to-image model. Because of that, generations for unseen prompts may be poor or may not be faithful to the input text (this may also happen for seen examples but less common).

For faster sampling, you may use a GPU runtime. If you switch the runtime type, be sure to rerun the entire notebook.

```
[18]: # text = "crying face" # seen example, good generations  
text = "face with cowboy hat" # seen example, good generations  
# text = "crying face with cowboy hat" # unseen example, bad generations  
text_emb = get_text_emb(text)  
text_emb = text_emb[None].expand(5, -1).to(device)
```

```

img = trainer.diffusion_model.sample(
    batch_size=5,
    model_kwargs={"text_emb": text_emb},
    return_all_timesteps=True
)
show_images(img[:, ::20])

```

sampling loop time step: 0%| | 0/100 [00:00<?, ?it/s]



6.4 Classifier Free Guidance

Generative models are typically evaluated on fidelity (i.e., the quality or realism of the generated samples) and diversity (the variability or coverage of the sample space). For conditional generative models, fidelity additionally refers to how faithfully the generated samples adhere to the input conditions. These two metrics are often in tension with each other, leading to a trade-off between them. Ho et al. introduced a simple technique called classifier-free guidance [3], which allows explicit

control over this trade-off.

In classifier-free guidance, during the training of a conditional diffusion model $\epsilon_\theta(x_t, t, c)$, the condition c is randomly dropped (i.e. replaced with $c = \phi$) with some probability (typically 0.1 to 0.2). During each denoising step of sampling, the prediction is updated as:

$$\epsilon_\theta(x_t, t, c) \leftarrow (w + 1)\epsilon_\theta(x_t, t, c) - w\epsilon_\theta(x_t, t, \phi)$$

where w is a positive scalar (the guidance scale). In other words, we perform two predictions during each denoising step, one conditional and one unconditional, and combine them linearly to favor the conditional generation. w is a hyperparameter that is tuned to optimize a model-specific evaluation metric. A higher w makes the generations more faithful to the condition but tends to reduce their diversity.

[3] Classifier-Free Diffusion Guidance. Jonathan Ho, Tim Salimans. [Link](#)

Implement the classifier-free guidance in `Unet.cfg_forward` method in `cs231n/unet.py` and test it below. You should see relative error less than $1e-6$.

```
[19]: np.random.seed(231)
      torch.manual_seed(231)

      dim = 4
      condition_dim = 4
      dim_mults = (2, 4)
      unet = Unet(dim=dim, condition_dim=condition_dim, dim_mults=dim_mults)
      unet.eval()

      b = 2
      h = w = 4
      inp_x = torch.randn(b, 3, h, w)
      inp_text_emb = torch.randn(b, condition_dim)
      inp_t = torch.tensor([8, 25]).long()
      out = unet.forward(x=inp_x, time=inp_t,
                        model_kwargs={"text_emb": inp_text_emb, "cfg_scale": 2.31}
                        ).detach().numpy()
      expected_out = np.array(
          [[[[ 0.3063889,  0.49175858,  0.53406465,  0.17793494],
              [ 1.0605793,  0.07291106,  0.26333052,  0.2807781 ],
              [ 0.47007525,  1.0033083,  0.15596265,  0.21308547],
              [ 0.6635294,  0.69716704,  0.7332697, -0.14285703]],
            [[-0.14788617, -0.17531544,  0.06930217,  0.24703366],
              [ 0.2998472, -0.25145096,  0.06968095,  0.28924388],
              [ 0.311122,  0.35599577,  0.6346719, -0.30859607],
              [ 0.3706389,  0.33177835,  0.29829532,  0.11159261]],
            [[ 0.34566414,  0.79019654,  0.6489645,  0.5683719 ],
              [ 0.98117805,  0.176887,  0.6220584,  0.6494926 ],
              [ 0.66813314,  0.90341353,  0.65361,  0.21999073],
              [ 0.73401093,  0.75344837,  0.71971107,  0.32417077]]],
            [[ 0.46322435, -1.3915588,  1.0951779,  0.37392217],
```

```
[ 0.6877824, 0.14802843, -0.7159017, 0.22476494],
[-0.28032526, 0.50201464, 0.0972808, 0.23250426],
[ 0.38375473, 0.86657, 0.4424622, 0.2161349 ]],
[[-0.34920466, -0.48507357, 0.06529156, 0.29454708],
[-0.62153226, 0.0083245, 1.7356712, -0.0690068 ],
[ 0.7745638, -0.11548355, 0.237822, 0.22680345],
[ 0.06337905, 0.7001358, 0.36441565, 0.17163567]],
[[ 0.16040441, -0.2551211, 0.8887521, 0.11380547],
[-0.12803352, 0.2783994, 1.4620805, 0.19044566],
[ 0.9112923, 0.75666, 0.3765695, 0.66629255],
[ 0.3461262, 1.0881201, 0.6392708, 0.31749305]]]]])
```

```
print("forward error: ", rel_error(out, expected_out))
```

Classifier-free guidance scale: 2.31
forward error: 4.9889353274740925e-05

Run the cell below to visualize emoji generations using classifier-free guidance. Feel free to modify the “cfg_scale” parameter value as well. As mentioned earlier, since our model does not generalize well, you may or may not observe faithful generations even with a high guidance scale.

```
[20]: # text = "crying face" # seen example, good generations
text = "face with cowboy hat" # seen example, good generations
# text = "crying face with cowboy hat" # unseen example, bad generations
text_emb = get_text_emb(text)
text_emb = text_emb[None].expand(5, -1).to(device)

img = trainer.diffusion_model.sample(
    batch_size=5,
    model_kwargs={"text_emb": text_emb, "cfg_scale": 1},
    return_all_timesteps=True
)
show_images(img[:, ::20])
```

sampling loop time step: 0% | 0/100 [00:00<?, ?it/s]

Classifier-free guidance scale: 1



[20] :

CLIP_DINO

September 10, 2026

```
[1]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive', force_remount=True)

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment3/'
FOLDERNAME = "cs231n/assignments/assignment3/"
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))
```

Mounted at /content/drive

```
[2]: # This downloads the COCO dataset to your Drive if it doesn't already exist
# (you should already have this dataset from a previous notebook!)
# Uncomment the following if you don't have it.
%cd /content/drive/My\ Drive/$FOLDERNAME/cs231n/datasets/
[!] bash get_coco_captioning.sh
%cd /content/drive/My\ Drive/$FOLDERNAME
```

```
/content/drive/My Drive/cs231n/assignments/assignment3/cs231n/datasets
bash: get_coco_captioning.sh: No such file or directory
/content/drive/My Drive/cs231n/assignments/assignment3
```

```
[3]: # Some useful python libraries
[!] pip install ftty regex tqdm
[!] pip install git+https://github.com/openai/CLIP.git
[!] pip install decord
```

Collecting ftty

Downloading ftty-6.3.1-py3-none-any.whl.metadata (7.3 kB)

Requirement already satisfied: regex in /usr/local/lib/python3.11/dist-packages (2024.11.6)

Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages


```

(4.67.1)
Requirement already satisfied: wcwidth in /usr/local/lib/python3.11/dist-
packages (from ftfy) (0.2.13)
Downloading ftfy-6.3.1-py3-none-any.whl (44 kB)
                                0.0/44.8 kB
? eta -:--:--
                                44.8/44.8 kB 4.5

MB/s eta 0:00:00
Installing collected packages: ftfy
Successfully installed ftfy-6.3.1
Collecting git+https://github.com/openai/CLIP.git
  Cloning https://github.com/openai/CLIP.git to /tmp/pip-req-build-mdo33lhx
  Running command git clone --filter=blob:none --quiet
https://github.com/openai/CLIP.git /tmp/pip-req-build-mdo33lhx
  Resolved https://github.com/openai/CLIP.git to commit
d05afc436d78f1c48dc0dbf8e5980a9d471f35f6
  Preparing metadata (setup.py) ... done
Requirement already satisfied: ftfy in /usr/local/lib/python3.11/dist-packages
(from clip==1.0) (6.3.1)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-
packages (from clip==1.0) (24.2)
Requirement already satisfied: regex in /usr/local/lib/python3.11/dist-packages
(from clip==1.0) (2024.11.6)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages
(from clip==1.0) (4.67.1)
Requirement already satisfied: torch in /usr/local/lib/python3.11/dist-packages
(from clip==1.0) (2.6.0+cu124)
Requirement already satisfied: torchvision in /usr/local/lib/python3.11/dist-
packages (from clip==1.0) (0.21.0+cu124)
Requirement already satisfied: wcwidth in /usr/local/lib/python3.11/dist-
packages (from ftfy->clip==1.0) (0.2.13)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-
packages (from torch->clip==1.0) (3.18.0)
Requirement already satisfied: typing-extensions>=4.10.0 in
/usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (4.14.1)
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-
packages (from torch->clip==1.0) (3.5)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.11/dist-packages
(from torch->clip==1.0) (3.1.6)
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages
(from torch->clip==1.0) (2025.3.2)
Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch->clip==1.0)
  Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch->clip==1.0)
  Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)

```

```

Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch->clip==1.0)
  Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch->clip==1.0)
  Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-
manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cublas-cu12==12.4.5.8 (from torch->clip==1.0)
  Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cufft-cu12==11.2.1.3 (from torch->clip==1.0)
  Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-curand-cu12==10.3.5.147 (from torch->clip==1.0)
  Downloading nvidia_curand_cu12-10.3.5.147-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch->clip==1.0)
  Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-
manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cusparselt-cu12==12.3.1.170 (from torch->clip==1.0)
  Downloading nvidia_cusparselt_cu12-12.3.1.170-py3-none-
manylinux2014_x86_64.whl.metadata (1.6 kB)
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in
/usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (0.6.2)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in
/usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in
/usr/local/lib/python3.11/dist-packages (from torch->clip==1.0) (12.4.127)
Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch->clip==1.0)
  Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl.metadata (1.5 kB)
Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist-
packages (from torch->clip==1.0) (3.2.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-
packages (from torch->clip==1.0) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in
/usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch->clip==1.0)
(1.3.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages
(from torchvision->clip==1.0) (2.0.2)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in
/usr/local/lib/python3.11/dist-packages (from torchvision->clip==1.0) (11.2.1)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.11/dist-packages (from jinja2->torch->clip==1.0) (3.0.2)
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl (363.4
MB)

363.4/363.4 MB
5.2 MB/s eta 0:00:00
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-

```

```

manylinux2014_x86_64.whl (13.8 MB)
13.8/13.8 MB
94.5 MB/s eta 0:00:00
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (24.6 MB)
24.6/24.6 MB
79.2 MB/s eta 0:00:00
Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (883 kB)
883.7/883.7 kB
78.8 MB/s eta 0:00:00
Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl
(664.8 MB)
664.8/664.8 MB
3.1 MB/s eta 0:00:00
Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl
(211.5 MB)
211.5/211.5 MB
7.7 MB/s eta 0:00:00
Downloading nvidia_curand_cu12-10.3.5.147-py3-none-
manylinux2014_x86_64.whl (56.3 MB)
56.3/56.3 MB
16.2 MB/s eta 0:00:00
Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-
manylinux2014_x86_64.whl (127.9 MB)
127.9/127.9 MB
10.8 MB/s eta 0:00:00
Downloading nvidia_cusparses_cu12-12.3.1.170-py3-none-
manylinux2014_x86_64.whl (207.5 MB)
207.5/207.5 MB
8.0 MB/s eta 0:00:00
Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-
manylinux2014_x86_64.whl (21.1 MB)
21.1/21.1 MB
45.5 MB/s eta 0:00:00
Building wheels for collected packages: clip
  Building wheel for clip (setup.py) ... done
  Created wheel for clip: filename=clip-1.0-py3-none-any.whl size=1369490
sha256=016dc066a4522381806e96c5ff3e72dbd09e18d9de7d20db427a37a4a838e43e
  Stored in directory: /tmp/pip-ephem-wheel-cache-
hrw2_geu/wheels/3f/7c/a4/9b490845988bf7a4db33674d52f709f088f64392063872eb9a
Successfully built clip
Installing collected packages: nvidia-nvjitlink-cu12, nvidia-curand-cu12,
nvidia-cufft-cu12, nvidia-cuda-runtime-cu12, nvidia-cuda-nvrtc-cu12, nvidia-
cuda-cupti-cu12, nvidia-cublas-cu12, nvidia-cusparses-cu12, nvidia-cudnn-cu12,
nvidia-cusolver-cu12, clip
  Attempting uninstall: nvidia-nvjitlink-cu12
    Found existing installation: nvidia-nvjitlink-cu12 12.5.82

```

```

Uninstalling nvidia-nvjitlink-cu12-12.5.82:
  Successfully uninstalled nvidia-nvjitlink-cu12-12.5.82
Attempting uninstall: nvidia-curand-cu12
  Found existing installation: nvidia-curand-cu12 10.3.6.82
Uninstalling nvidia-curand-cu12-10.3.6.82:
  Successfully uninstalled nvidia-curand-cu12-10.3.6.82
Attempting uninstall: nvidia-cufft-cu12
  Found existing installation: nvidia-cufft-cu12 11.2.3.61
Uninstalling nvidia-cufft-cu12-11.2.3.61:
  Successfully uninstalled nvidia-cufft-cu12-11.2.3.61
Attempting uninstall: nvidia-cuda-runtime-cu12
  Found existing installation: nvidia-cuda-runtime-cu12 12.5.82
Uninstalling nvidia-cuda-runtime-cu12-12.5.82:
  Successfully uninstalled nvidia-cuda-runtime-cu12-12.5.82
Attempting uninstall: nvidia-cuda-nvrtc-cu12
  Found existing installation: nvidia-cuda-nvrtc-cu12 12.5.82
Uninstalling nvidia-cuda-nvrtc-cu12-12.5.82:
  Successfully uninstalled nvidia-cuda-nvrtc-cu12-12.5.82
Attempting uninstall: nvidia-cuda-cupti-cu12
  Found existing installation: nvidia-cuda-cupti-cu12 12.5.82
Uninstalling nvidia-cuda-cupti-cu12-12.5.82:
  Successfully uninstalled nvidia-cuda-cupti-cu12-12.5.82
Attempting uninstall: nvidia-cublas-cu12
  Found existing installation: nvidia-cublas-cu12 12.5.3.2
Uninstalling nvidia-cublas-cu12-12.5.3.2:
  Successfully uninstalled nvidia-cublas-cu12-12.5.3.2
Attempting uninstall: nvidia-cusparse-cu12
  Found existing installation: nvidia-cusparse-cu12 12.5.1.3
Uninstalling nvidia-cusparse-cu12-12.5.1.3:
  Successfully uninstalled nvidia-cusparse-cu12-12.5.1.3
Attempting uninstall: nvidia-cudnn-cu12
  Found existing installation: nvidia-cudnn-cu12 9.3.0.75
Uninstalling nvidia-cudnn-cu12-9.3.0.75:
  Successfully uninstalled nvidia-cudnn-cu12-9.3.0.75
Attempting uninstall: nvidia-cusolver-cu12
  Found existing installation: nvidia-cusolver-cu12 11.6.3.83
Uninstalling nvidia-cusolver-cu12-11.6.3.83:
  Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83
Successfully installed clip-1.0 nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-
cu12-12.4.127 nvidia-cuda-nvrtc-cu12-12.4.127 nvidia-cuda-runtime-cu12-12.4.127
nvidia-cudnn-cu12-9.1.0.70 nvidia-cufft-cu12-11.2.1.3 nvidia-curand-
cu12-10.3.5.147 nvidia-cusolver-cu12-11.6.1.9 nvidia-cusparse-cu12-12.3.1.170
nvidia-nvjitlink-cu12-12.4.127
Collecting decord
  Downloading decord-0.6.0-py3-none-manylinux2010_x86_64.whl.metadata (422
bytes)
Requirement already satisfied: numpy>=1.14.0 in /usr/local/lib/python3.11/dist-
packages (from decord) (2.0.2)

```

Downloading decord-0.6.0-py3-none-manylinux2010_x86_64.whl (13.6 MB)

13.6/13.6 MB

42.3 MB/s eta 0:00:00

Installing collected packages: decord

Successfully installed decord-0.6.0

1 State-of-the-Art Pretrained Image Models

In the previous exercise, you learned about [SimCLR](#) and how contrastive self-supervised learning can be used to learn meaningful image representations. In this notebook, we will explore two more recent models that also aim to learn high-quality visual representations and have demonstrated strong and robust performance on a variety of downstream tasks.

First, we will examine the [CLIP](#) model. Like SimCLR, CLIP uses a contrastive learning objective, but instead of contrasting two augmented views of the same image, it contrasts two different modalities: text and image. To train CLIP, OpenAI collected a large dataset of ~400M image-text pairs from the internet, including sources like Wikipedia and image alt text. The resulting model learns rich, high-level image features and has achieved impressive zero-shot performance on many vision benchmarks.

Next, we will explore [DINO](#), a self-supervised learning method for vision tasks that applies contrastive learning in a self-distillation framework with multi-crop augmentation strategy. The authors showed that the features learned by DINO ViTs are fine-grained and semantically rich with explicit information about the semantic segmentation of the image.

2 CLIP

As explained above, CLIP’s training objective incorporates both text and images, building upon the principles of contrastive learning. Consider this quote from the SimCLR notebook: >The goal of the contrastive loss is to maximize agreement between the final vectors $z_i = g(h_i)$ and $z_j = g(h_j)$.

Similarly, CLIP is trained to maximize agreement between two vectors. However, because these vectors come from different modalities, CLIP uses two separate encoders: a transformer-based Text Encoder and a Vision Transformer (ViT)-based Image Encoder. Note that some smaller, more efficient versions of CLIP use a ResNet as the Image Encoder instead of a ViT.


Run the cell below to visualize the training and inference pipeline of CLIP.

During the pretraining phase, each batch consists of multiple images along with their corresponding captions. Each image is independently processed by an Image Encoder—typically a visual model like a Vision Transformer (ViT) or a Convolutional Neural Network (ConvNet)—which produces an image embedding I_n . Likewise, each caption is independently processed by a Text Encoder to generate a corresponding text embedding T_n . Next, we compute the pairwise similarities between all image-text combinations, meaning each image is compared with every caption, and vice versa. The training objective is to maximize the similarity scores along the diagonal of the resulting similarity matrix – that is, the scores for the matching image-caption pairs (I_n, T_n) . Through backpropagation, the model learns to assign higher similarity scores to true matches than to mismatched pairs.

Through this setup, CLIP effectively learns to represent images and texts in a shared latent space.

In this space, semantic concepts are encoded in a modality-independent way, enabling meaningful cross-modal comparisons between visual and textual inputs.

```
[4]: from IPython.display import Image as ColabImage
ColabImage(f'/content/drive/My Drive/{FOLDERNAME}/CLIP.png')
```



Inline Question 1

Why does CLIP's learning depend on the batch size? If the batch size is fixed, what strategy can we use to learn rich image features?

Your Answer :

CLIP is trained with a symmetric contrastive (InfoNCE) loss computed within each batch: for a batch of N image-text pairs, each image has exactly one positive (its own caption) and $N - 1$ negatives (the other captions in the batch). A larger batch therefore provides more negative pairs per update, which gives a richer and harder discrimination signal, better separates the embedding space, and exposes the model to more diverse concepts per step — this is why zero-shot performance was observed to scale log-linearly with batch size. With a fixed (small) batch size, we can artificially enlarge the pool of negatives using a memory bank or a MoCo-style momentum queue that stores feature embeddings from previous batches as extra negatives; we can also cache text embeddings (they are cheap to store), or apply multiple augmented crops per image so that each sample contributes several views. In short: decouple the number of effective negatives from the physical batch size.

3 Loading COCO dataset

We'll use the same captioning dataset you used to train your RNN captioning model, but instead of generating the captions let's see if we can match each image to the correct caption.

```
[5]: import sys
import types
import importlib

if "imp" not in sys.modules:
```

```

imp = types.ModuleType("imp")
imp.reload = importlib.reload
sys.modules["imp"] = imp

%load_ext autoreload
%autoreload 2

import time, os, json
import numpy as np
import matplotlib.pyplot as plt
import torch
import clip
import torch
from tqdm.auto import tqdm

from PIL import Image
from cs231n.clip_dino import *

def rel_error(x, y):
    """Returns relative error."""
    return np.max(np.abs(x - y) / (np.maximum(1e-10, np.abs(x) + np.abs(y))))

```

```

[6]: from cs231n.coco_utils import load_coco_data, sample_coco_minibatch,
      ↪ decode_captions
      from cs231n.image_utils import image_from_url

```

```

[7]: # Load COCO data from disk into a dictionary.
      # this is the same dataset you used for the RNN captioning notebook :)
      data = load_coco_data(pca_features=True)

      # Print out all the keys and values from the data dictionary.
      for k, v in data.items():
          if type(v) == np.ndarray:
              print(k, type(v), v.shape, v.dtype)
          else:
              print(k, type(v), len(v))

```

```

base_dir /content/drive/MyDrive/cs231n/assignments/assignment3/cs231n/datasets/
coco_captioning

```

```

train_captions <class 'numpy.ndarray'> (400135, 17) int32
train_image_idxes <class 'numpy.ndarray'> (400135,) int32
val_captions <class 'numpy.ndarray'> (195954, 17) int32
val_image_idxes <class 'numpy.ndarray'> (195954,) int32
train_features <class 'numpy.ndarray'> (82783, 512) float32
val_features <class 'numpy.ndarray'> (40504, 512) float32
idx_to_word <class 'list'> 1004
word_to_idx <class 'dict'> 1004

```

```
train_urls <class 'numpy.ndarray'> (82783,) <U63
val_urls <class 'numpy.ndarray'> (40504,) <U63
```

```
[8]: # we're just using the loaded captions from COCO, so we need to decode them and
      ↪ get rid of the special tokens.
      decoded_captions= []
      for caption in data['val_captions']:
          caption = decode_captions(caption, data['idx_to_word'])\
              .replace('<START>', '')\
              .replace('<END>', '')\
              .replace('<UNK>', '')\
              .strip()
          decoded_captions.append(caption)
```

```
[9]: # lets get 10 examples
      mask = np.array([135428, 122586, 122814, 133173, 176639, 163828, 98169, 6931,
                       19488, 175760])
      first_captions = [decoded_captions[elem] for elem in mask]

      img_idxes = data['val_image_idxes'][mask] # the images the captions refer to
      first_images = [image_from_url(data['val_urls'][j]) for j in img_idxes]
```

```
[10]: for i, (caption, image) in enumerate(zip(first_captions, first_images)):
        plt.imshow(image)
        plt.axis('off')
        caption_str = caption
        plt.title(caption_str)
        plt.show()
```

Output hidden; open in <https://colab.research.google.com> to view.

4 Running the CLIP Model

First we'll use the pretrained CLIP model to extract features from the texts and images separately.

```
[11]: device = "cuda" if torch.cuda.is_available() else "cpu"
      clip_model, clip_preprocess = clip.load("ViT-B/32", device=device)
```

```
100%|          | 338M/338M [00:03<00:00, 102MiB/s]
```

```
[12]: # You can check the model layers by printing the model.
      # CLIP's model code is available at https://github.com/openai/CLIP/tree/main/
      ↪ clip
      print(clip_model)
```

```
CLIP(
  (visual): VisionTransformer(
    (conv1): Conv2d(3, 768, kernel_size=(32, 32), stride=(32, 32), bias=False)
```



```

(ln_pre): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
(transformer): Transformer(
  (resblocks): Sequential(
    (0): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
      )
      (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
      (mlp): Sequential(
        (c_fc): Linear(in_features=768, out_features=3072, bias=True)
        (gelu): QuickGELU()
        (c_proj): Linear(in_features=3072, out_features=768, bias=True)
      )
      (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    )
    (1): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
      )
      (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
      (mlp): Sequential(
        (c_fc): Linear(in_features=768, out_features=3072, bias=True)
        (gelu): QuickGELU()
        (c_proj): Linear(in_features=3072, out_features=768, bias=True)
      )
      (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    )
    (2): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
      )
      (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
      (mlp): Sequential(
        (c_fc): Linear(in_features=768, out_features=3072, bias=True)
        (gelu): QuickGELU()
        (c_proj): Linear(in_features=3072, out_features=768, bias=True)
      )
      (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    )
    (3): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
      )
      (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)

```

```

        (mlp): Sequential(
          (c_fc): Linear(in_features=768, out_features=3072, bias=True)
          (gelu): QuickGELU()
          (c_proj): Linear(in_features=3072, out_features=768, bias=True)
        )
        (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
      )
    (4): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
      )
      (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
      (mlp): Sequential(
        (c_fc): Linear(in_features=768, out_features=3072, bias=True)
        (gelu): QuickGELU()
        (c_proj): Linear(in_features=3072, out_features=768, bias=True)
      )
      (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    )
    (5): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
      )
      (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
      (mlp): Sequential(
        (c_fc): Linear(in_features=768, out_features=3072, bias=True)
        (gelu): QuickGELU()
        (c_proj): Linear(in_features=3072, out_features=768, bias=True)
      )
      (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    )
    (6): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
      )
      (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
      (mlp): Sequential(
        (c_fc): Linear(in_features=768, out_features=3072, bias=True)
        (gelu): QuickGELU()
        (c_proj): Linear(in_features=3072, out_features=768, bias=True)
      )
      (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    )
    (7): ResidualAttentionBlock(
      (attn): MultiheadAttention(

```

```

        (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
    )
    (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    (mlp): Sequential(
      (c_fc): Linear(in_features=768, out_features=3072, bias=True)
      (gelu): QuickGELU()
      (c_proj): Linear(in_features=3072, out_features=768, bias=True)
    )
    (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
  )
  (8): ResidualAttentionBlock(
    (attn): MultiheadAttention(
      (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
    )
    (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    (mlp): Sequential(
      (c_fc): Linear(in_features=768, out_features=3072, bias=True)
      (gelu): QuickGELU()
      (c_proj): Linear(in_features=3072, out_features=768, bias=True)
    )
    (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
  )
  (9): ResidualAttentionBlock(
    (attn): MultiheadAttention(
      (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
    )
    (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    (mlp): Sequential(
      (c_fc): Linear(in_features=768, out_features=3072, bias=True)
      (gelu): QuickGELU()
      (c_proj): Linear(in_features=3072, out_features=768, bias=True)
    )
    (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
  )
  (10): ResidualAttentionBlock(
    (attn): MultiheadAttention(
      (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
    )
    (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    (mlp): Sequential(
      (c_fc): Linear(in_features=768, out_features=3072, bias=True)
      (gelu): QuickGELU()
      (c_proj): Linear(in_features=3072, out_features=768, bias=True)
    )
  )

```

```

        (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    )
    (11): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=768,
out_features=768, bias=True)
      )
      (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
      (mlp): Sequential(
        (c_fc): Linear(in_features=768, out_features=3072, bias=True)
        (gelu): QuickGELU()
        (c_proj): Linear(in_features=3072, out_features=768, bias=True)
      )
      (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    )
  )
  (ln_post): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
)
(transformer): Transformer(
  (resblocks): Sequential(
    (0): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=512,
out_features=512, bias=True)
      )
      (ln_1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
      (mlp): Sequential(
        (c_fc): Linear(in_features=512, out_features=2048, bias=True)
        (gelu): QuickGELU()
        (c_proj): Linear(in_features=2048, out_features=512, bias=True)
      )
      (ln_2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
    )
    (1): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=512,
out_features=512, bias=True)
      )
      (ln_1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
      (mlp): Sequential(
        (c_fc): Linear(in_features=512, out_features=2048, bias=True)
        (gelu): QuickGELU()
        (c_proj): Linear(in_features=2048, out_features=512, bias=True)
      )
      (ln_2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
    )
    (2): ResidualAttentionBlock(

```

```

        (attn): MultiheadAttention(
          (out_proj): NonDynamicallyQuantizableLinear(in_features=512,
out_features=512, bias=True)
        )
        (ln_1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
        (mlp): Sequential(
          (c_fc): Linear(in_features=512, out_features=2048, bias=True)
          (gelu): QuickGELU()
          (c_proj): Linear(in_features=2048, out_features=512, bias=True)
        )
        (ln_2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
      )
    (3): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=512,
out_features=512, bias=True)
      )
      (ln_1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
      (mlp): Sequential(
        (c_fc): Linear(in_features=512, out_features=2048, bias=True)
        (gelu): QuickGELU()
        (c_proj): Linear(in_features=2048, out_features=512, bias=True)
      )
      (ln_2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
    )
    (4): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=512,
out_features=512, bias=True)
      )
      (ln_1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
      (mlp): Sequential(
        (c_fc): Linear(in_features=512, out_features=2048, bias=True)
        (gelu): QuickGELU()
        (c_proj): Linear(in_features=2048, out_features=512, bias=True)
      )
      (ln_2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
    )
    (5): ResidualAttentionBlock(
      (attn): MultiheadAttention(
        (out_proj): NonDynamicallyQuantizableLinear(in_features=512,
out_features=512, bias=True)
      )
      (ln_1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
      (mlp): Sequential(
        (c_fc): Linear(in_features=512, out_features=2048, bias=True)
        (gelu): QuickGELU()
        (c_proj): Linear(in_features=2048, out_features=512, bias=True)

```

```

    )
    (ln_2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
)
(6): ResidualAttentionBlock(
  (attn): MultiheadAttention(
    (out_proj): NonDynamicallyQuantizableLinear(in_features=512,
out_features=512, bias=True)
  )
  (ln_1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
  (mlp): Sequential(
    (c_fc): Linear(in_features=512, out_features=2048, bias=True)
    (gelu): QuickGELU()
    (c_proj): Linear(in_features=2048, out_features=512, bias=True)
  )
  (ln_2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
)
(7): ResidualAttentionBlock(
  (attn): MultiheadAttention(
    (out_proj): NonDynamicallyQuantizableLinear(in_features=512,
out_features=512, bias=True)
  )
  (ln_1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
  (mlp): Sequential(
    (c_fc): Linear(in_features=512, out_features=2048, bias=True)
    (gelu): QuickGELU()
    (c_proj): Linear(in_features=2048, out_features=512, bias=True)
  )
  (ln_2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
)
(8): ResidualAttentionBlock(
  (attn): MultiheadAttention(
    (out_proj): NonDynamicallyQuantizableLinear(in_features=512,
out_features=512, bias=True)
  )
  (ln_1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
  (mlp): Sequential(
    (c_fc): Linear(in_features=512, out_features=2048, bias=True)
    (gelu): QuickGELU()
    (c_proj): Linear(in_features=2048, out_features=512, bias=True)
  )
  (ln_2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
)
(9): ResidualAttentionBlock(
  (attn): MultiheadAttention(
    (out_proj): NonDynamicallyQuantizableLinear(in_features=512,
out_features=512, bias=True)
  )
  (ln_1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)

```

```

        (mlp): Sequential(
          (c_fc): Linear(in_features=512, out_features=2048, bias=True)
          (gelu): QuickGELU()
          (c_proj): Linear(in_features=2048, out_features=512, bias=True)
        )
        (ln_2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
      )
      (10): ResidualAttentionBlock(
        (attn): MultiheadAttention(
          (out_proj): NonDynamicallyQuantizableLinear(in_features=512,
out_features=512, bias=True)
        )
        (ln_1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
        (mlp): Sequential(
          (c_fc): Linear(in_features=512, out_features=2048, bias=True)
          (gelu): QuickGELU()
          (c_proj): Linear(in_features=2048, out_features=512, bias=True)
        )
        (ln_2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
      )
      (11): ResidualAttentionBlock(
        (attn): MultiheadAttention(
          (out_proj): NonDynamicallyQuantizableLinear(in_features=512,
out_features=512, bias=True)
        )
        (ln_1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
        (mlp): Sequential(
          (c_fc): Linear(in_features=512, out_features=2048, bias=True)
          (gelu): QuickGELU()
          (c_proj): Linear(in_features=2048, out_features=512, bias=True)
        )
        (ln_2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
      )
    )
  )
  (token_embedding): Embedding(49408, 512)
  (ln_final): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
)

```

```

[13]: # First, we encode the captions into vectors in the shared embedding space.
      # Since we're using a Transformer as the text encoder, we need to tokenize the
      ↪text first.
      text_tokens = clip.tokenize(first_captions).to(device)
      with torch.no_grad():
        text_features = clip_model.encode_text(text_tokens)

      # Sanity check, print the shape

```

```
print(text_features.shape)
```

```
# Note: For your implementations, use the above clip_model.encode_text() ↵  
↪function. Avoid using clip_model.forward().
```

```
torch.Size([10, 512])
```

```
[14]: # Then, we encode the images into the same embedding space.
```

```
processed_images = [  
    clip_preprocess(Image.fromarray(img)).unsqueeze(0)  
    for img in first_images  
]  
images_tensor = torch.cat(processed_images, dim=0).to(device)
```

```
with torch.no_grad():  
    image_features = clip_model.encode_image(images_tensor)
```

```
# sanity check, print the shape  
print(image_features.shape)
```

```
# Note: For your implementations, use the above clip_model.encode_image() ↵  
↪function. Avoid using clip_model.forward().
```

```
torch.Size([10, 512])
```

Open `cs231n/clip_dino.py` and implement `get_similarity_no_loop` to compute similarity scores between text features and image features. Test your implementation below, you should see relative errors less than $1e-5$.

```
[15]: from cs231n.clip_dino import get_similarity_no_loop
```

```
torch.manual_seed(231)  
np.random.seed(231)  
M, N, D = 5, 6, 10
```

```
test_text_features = torch.randn(N, D)  
test_image_features = torch.randn(M, D)  
out = get_similarity_no_loop(test_text_features, test_image_features)
```

```
expected_out = np.array([  
    [ 0.1867811 , -0.23494351,  0.44155994, -0.18950461,  0.00100103],  
    [ 0.17905031, -0.25469488, -0.64330417,  0.25097957, -0.09327742],  
    [-0.4407011 , -0.4365381 ,  0.32857686, -0.3765278 ,  0.01049389],  
    [ 0.24815483,  0.42157224, -0.08459304,  0.14132318, -0.26935193],  
    [ 0.02309848, -0.01441101,  0.5469337 ,  0.6018773 ,  0.21581158],  
    [ 0.41579214, -0.014449 , -0.7242257 ,  0.39348006,  0.0822239 ],  
]).astype(np.float32)
```

```
print("relative error: ", rel_error(out.numpy(), expected_out))
```


relative error: 1.3374006e-06

```
[16]: # Let's visualize the similarities between our batch of images and their
      ↪ captions.

similarities = get_similarity_no_loop(text_features, image_features).cpu().
      ↪detach().numpy()

plt.figure(figsize=(20, 14))
plt.imshow(similarities, vmin=0.1, vmax=0.3)
plt.yticks(range(len(text_features)), first_captions, fontsize=18)
plt.xticks([])
for i, image in enumerate(first_images):
    plt.imshow(image, extent=(i - 0.5, i + 0.5, -1.6, -0.6), origin="lower")
for x in range(similarities.shape[1]):
    for y in range(similarities.shape[0]):
        plt.text(x, y, f"{similarities[y, x]:.2f}", ha="center", va="center",
        ↪size=12)

for side in ["left", "top", "right", "bottom"]:
    plt.gca().spines[side].set_visible(False)

plt.xlim([-0.5, len(image_features) - 0.5])
plt.ylim([len(text_features) + 0.5, -2])

plt.title("Cosine similarity between text and image features", size=20)
plt.show()
```



5 Zero Shot Classifier

You will be able to see a high similarity between matching image-caption pairs above. We can leverage this property to design an image classifier that doesn't require any labeled data (i.e., a zero-shot classifier). Each class can be represented using an appropriate natural language description, and any input image will be classified into the class whose description has the highest similarity with the image in CLIP's embedding space.

Implement `clip_zero_shot_classifier` in `cs231n/clip_dino.py` and test it below. You should be able to see the following predictions:

['a person', 'an animal', 'an animal', 'food', 'a person', 'a landscape', 'other', 'other', 'other', 'a person']

```
[17]: from cs231n.clip_dino import clip_zero_shot_classifier

classes = ["a person", "an animal", "food", "a landscape", "other"]

pred_classes = clip_zero_shot_classifier(
    clip_model, clip_preprocess, first_images, classes, device)

print(pred_classes)
```

['a person', 'an animal', 'an animal', 'food', 'a person', 'a landscape', 'other', 'an animal', 'other', 'a person']

Run the cell below to visualize the predictions. As you can see, CLIP offers a straightforward way to perform reasonable zero-shot classification across any class taxonomy.

CLIP was the first model to outperform standard supervised training on ImageNet classification without using any ImageNet images or labels (The original CLIP paper has many such interesting experiments and analysis).

```
[18]: # Visualize the zero shot predictions
for i, (pred_class, image) in enumerate(zip(pred_classes, first_images)):
    plt.imshow(image)
    plt.axis('off')
    plt.title(pred_class)
    plt.show()
```

Output hidden; open in <https://colab.research.google.com> to view.

6 Image Retrieval using CLIP

Just as we used CLIP to retrieve the matching class name for each image, we can also use it to retrieve matching images from text inputs (semantic image retrieval). Implement the

CLIPImageRetriever in cs231n/clip_dino.py and test it by running the two cells below. The expected top 2 outputs for each query are provided in the comments.

```
[19]: from cs231n.clip_dino import CLIPImageRetriever
clip_retriever = CLIPImageRetriever(clip_model, clip_preprocess, first_images, device)
```

```
[20]: query = "sports" # tennis, skateboard
# query = "black and white" # bathroom, zerbias
img_indices = clip_retriever.retrieve(query)

for img_index in img_indices:
    plt.imshow(first_images[img_index])
    plt.axis('off')
    plt.show()
```





Inline Question 2

CLIP learns to align image and text representations in a shared latent space using a contrastive loss. How would you extend this idea to more than two modalities?

Your Answer :

Give each modality its own encoder, project everything into one shared embedding space $\mathbb{R}^d$, and replace the two-way contrastive matrix with a multi-way objective. Concretely: (1) sum standard pairwise InfoNCE losses over every modality pair (image–text, image–audio, text–audio, ...), i.e. $\mathcal{L} = \sum_{i < j} \mathcal{L}_{ij}$; or (2) use a generalized N -way contrastive loss where each anchor’s positives are the samples jointly paired with it in the other modalities, and all unpaired in-batch samples serve as negatives, with a shared temperature τ . An efficient simplification (the ImageBind idea) is to bind every auxiliary modality directly to a reference modality (e.g., images): emergent alignment between the auxiliary modalities follows from transitivity of the embedding spaces. Either way, a semantically coherent tuple should end up close together in the shared space while mismatched combinations are pushed apart.

7 DINO

As mentioned earlier, models trained with vanilla contrastive learning methods such as SimCLR and CLIP require very large batch sizes. This makes them computationally expensive and limits their accessibility. Subsequent works, like BYOL, propose an alternative approach that avoids the need for numerous negative samples by using a student-teacher framework. This method performs surprisingly well and was later adopted by DINO .

Similar to SimCLR, DINO is trained to maximize the agreement between two vectors derived from different views of the same image. However, unlike SimCLR, DINO uses two separate encoders which are trained differently. The student network is updated via backpropagation to match the outputs of the teacher network. The teacher network is not updated via backpropagation; instead, its weights are updated using an exponential moving average (EMA) of the student's weights. This means that the teacher model evolves more slowly and provides a stable target for the student to learn from.

Run the cell below to visualize the DINO training pipeline.

```
[21]: from IPython.display import Image as ColabImage
      ColabImage(f'/content/drive/My Drive/{FOLDERNAME}/dino.gif')
```

Output hidden; open in <https://colab.research.google.com> to view.

```
[22]: # first let's get rid of the CLIP model that's currently using memory
      del clip_model
      # Uncomment the following if you are using GPU runtime
      # torch.cuda.empty_cache()
      # torch.cuda.ipc_collect()
```

```
[23]: # Load smallest dino model. ViT-S/8. Here ViT-S has ~22M parameters and
      # works on 8x8 patches.
      dino_model = torch.hub.load('facebookresearch/dino:main', 'dino_vits8')
      dino_model.eval().to(device)
```

Downloading: "https://github.com/facebookresearch/dino/zipball/main" to
/root/.cache/torch/hub/main.zip
Downloading: "https://dl.fbaipublicfiles.com/dino/dino_deitsmall8_pretrain/dino_deitsmall8_pretrain.pth" to
/root/.cache/torch/hub/checkpoints/dino_deitsmall8_pretrain.pth
100%| | 82.7M/82.7M [00:03<00:00, 26.6MB/s]

```
[23]: VisionTransformer(
  (patch_embed): PatchEmbed(
    (proj): Conv2d(3, 384, kernel_size=(8, 8), stride=(8, 8))
  )
  (pos_drop): Dropout(p=0.0, inplace=False)
  (blocks): ModuleList(
    (0-11): 12 x Block(
      (norm1): LayerNorm((384,)), eps=1e-06, elementwise_affine=True)
      (attn): Attention(
        (qkv): Linear(in_features=384, out_features=1152, bias=True)
        (attn_drop): Dropout(p=0.0, inplace=False)
        (proj): Linear(in_features=384, out_features=384, bias=True)
        (proj_drop): Dropout(p=0.0, inplace=False)
      )
      (drop_path): Identity()
```

```

(norm2): LayerNorm((384,), eps=1e-06, elementwise_affine=True)
(mlp): Mlp(
  (fc1): Linear(in_features=384, out_features=1536, bias=True)
  (act): GELU(approximate='none')
  (fc2): Linear(in_features=1536, out_features=384, bias=True)
  (drop): Dropout(p=0.0, inplace=False)
)
)
)
(norm): LayerNorm((384,), eps=1e-06, elementwise_affine=True)
(head): Identity()
)

```

```

[24]: # the image we will be playing around with
sample_image = Image.fromarray(first_images[0]).convert("RGB")
sample_image

```

[24]:



8 DINO Attention Maps

Since the loaded DINO checkpoint is based on the ViT architecture, we can visualize what each attention head is focusing on. The code below generates heatmaps showing which patches of the original image the [CLS] token attends to across the various heads in the final layer. Although this model was trained using a self-supervised objective without any explicit instruction to recognize


“structure” in images, still...

Do you notice any patterns?

```
[25]: # Preprocess
from torchvision import transforms as T
transform = T.Compose([
    T.Resize((480, 480)),
    T.ToTensor(),
    T.Normalize((0.485, 0.456, 0.406), (0.229, 0.224, 0.225)),
])
img_tensor = transform(sample_image)
w, h = img_tensor.shape[1:]
img_tensor = img_tensor[None].to(device)

# Extract attention
with torch.no_grad():
    attn = dino_model.get_last_selfattention(img_tensor)[0, :, 0, 1:]
nh, tokens = attn.shape
w_feat, h_feat = w // 8, h // 8
attn = attn.reshape(nh, w_feat, h_feat)
attn = torch.nn.functional.interpolate(attn.unsqueeze(0), scale_factor=8,
    ↪mode="nearest")[0].cpu().numpy()

# Plot attention heads
fig, axes = plt.subplots(1, nh, figsize=(3 * nh, 3))
for i in range(nh):
    ax = axes[i] if nh > 1 else axes
    ax.imshow(attn[i], cmap='inferno')
    ax.axis('off')
plt.show()
```



```
[26]: # Extract patch token features and discard [CLS] token.
with torch.no_grad():
    all_tokens = dino_model.get_intermediate_layers(img_tensor, n=1)[0] # (1, ↪
    ↪1+N, D)
    patch_tokens = all_tokens[:, 1:, :] # (N, D)

print(img_tensor.shape)
print(all_tokens.shape)
```

```
print(patch_tokens.shape)
```

```
torch.Size([1, 3, 480, 480])
```

```
torch.Size([1, 3601, 384])
```

```
torch.Size([1, 3600, 384])
```

Inline Question 3

How do we get the tensor shapes printed above? Explain your answer.

Your Answer :

A DINO ViT divides each (resized) frame into a non-overlapping grid of patches of size $P \times P$ ($P = 8$ or 16), so an $H \times W$ image produces $\frac{H}{P} \cdot \frac{W}{P}$ patch tokens (H and W are first made divisible by P by resizing/cropping). Dropping the CLS token, `process_frames` returns the remaining patch embeddings as a matrix of shape $[\frac{H}{P} \cdot \frac{W}{P}, D]$, where D is the ViT embedding dimension (e.g., $D = 384$ for ViT-S/8). `process_masks` then downsamples the ground-truth mask with exactly the same $P \times P$ grid and assigns each patch the majority class label inside it, giving one integer label per patch: a vector of shape $[\frac{H}{P} \cdot \frac{W}{P}]$. So both tensors have one row/entry per spatial patch — this 1-to-1 correspondence is what lets us train a patch-level classifier (or retrieve patches by similarity) on top of DINO features.

9 DINO Features

To understand what the model is encoding in each patch, we can visualize the contents of each patch token. Since these embeddings are high-dimensional and difficult to interpret directly, we'll use PCA to identify the directions of highest variance in the feature space.

In the next cell, we visualize the three principal directions of variance in the feature space. This reveals the dominant structure that the patch embeddings are capturing.

```
[27]: from sklearn.decomposition import PCA

np.random.seed(231)

# PCA
pca = PCA(n_components=3)
patch_pca = pca.fit_transform(patch_tokens.cpu().numpy()[0])

# Normalize PCA components to [0, 1] for RGB display
patch_rgb = (patch_pca - patch_pca.min(0)) / (patch_pca.max(0) - patch_pca.
↪min(0))

# Reshape to image grid (60x60, 3)
patch_rgb_img = patch_rgb.reshape(60, 60, 3)

# Show as image
plt.figure(figsize=(6, 6))
plt.imshow(patch_rgb_img)
```



```
plt.axis('off')
plt.title("Patch Embeddings (PCA → RGB)")
plt.show()
```



Inline Question 4

What kind of structure do you see in the visualization above? What does it imply when a region consistently appears in a specific color? What does it mean when two regions have distinctly different color? Remember that PCA reveals the directions of highest variance in the feature space across all patches. A patch's color reflects its distinct feature content.

Your Answer :

The PCA maps each patch's DINO feature onto the top-3 principal components, rendered as RGB: each patch feature $f_i \in \mathbb{R}^D$ is projected onto the first three principal directions and scaled to $[0, 255]$. We see spatially coherent, semantically meaningful regions — object boundaries (animal vs. background, body parts, etc.) emerge without any supervision, because DINO's self-distillation makes neighboring patches of the same object produce nearly identical features. When a region appears in one consistent color, all its patches have similar feature vectors, i.e., they encode the

same semantic content — effectively an unsupervised segmentation of that object/part. When two regions show distinctly different colors, their features point in very different directions along the highest-variance axes, meaning they are semantically very different (different objects or strongly different parts). In short, the color structure reveals that DINO patch features cluster by object identity and semantics, not just by low-level appearance.

10 A Simple Segmentation Model over DINO Features

In the previous section, we saw that DINO features can provide surprisingly good segmentation cues. Now, let's put that idea to the test by training a simple segmentation model on the [DAVIS dataset](#). The DAVIS dataset (Densely Annotated Video Segmentation) was created for video object segmentation tasks. It provides frame-by-frame, pixel-level annotations of objects within videos. For this experiment, we'll train our model using the annotations from just a single frame of a video and see how well it performs on the remaining frames of the same.

Our model will be intentionally minimal: we'll extract DINO features per patch and train a lightweight per-patch classifier using only the patches from that one annotated frame. Typically, you would train on the full dataset and evaluate on a separate validation set containing different videos. But here, we will test the one-shot capabilities of DINO features.

```
[28]: !wget -c -q --show-progress \
      https://data.vision.ee.ethz.ch/csergi/share/davis/DAVIS-2017-trainval-480p.
      ↪zip \
      -O /content/DAVIS-2017-trainval-480p.zip
```

```
/content/DAVIS-2017 100%[=====>] 794.19M 24.0MB/s in 35s
```

```
[29]: import os, io, zipfile
import numpy as np
import torch
import cv2
from PIL import Image
from torchvision import transforms as T
from cs231n.clip_dino import create_segmentation_overlay

class DavisDataset:

    def __init__(self):
        zip_path = "/content/DAVIS-2017-trainval-480p.zip"
        assert os.path.exists(zip_path), "zip cell"
        self.zf = zipfile.ZipFile(zip_path)
        names = self.zf.namelist()

        val_txt = [n for n in names if n.endswith("ImageSets/2017/val.txt")]
        assert val_txt, "zip cell"
        self.root = val_txt[0].split("/ImageSets/")[0]
        self.val_videos = [v.strip() for v in
```

```

        self.zf.read(val_txt[0]).decode().splitlines() if v.
↪strip()]
    print(f"{len(self.val_videos)} validation videos")

    self.img_tsfm = T.Compose([
        T.Resize((480, 480)),
        T.ToTensor(),
        T.Normalize((0.485, 0.456, 0.406), (0.229, 0.224, 0.225)),
    ])

    def get_sample(self, index):
        video = self.val_videos[index]
        f_pre = f"{self.root}/JPEGImages/480p/{video}/"
        m_pre = f"{self.root}/Annotations/480p/{video}/"

        frame_names = sorted(n for n in self.zf.namelist()
                               if n.startswith(f_pre) and n.endswith(".jpg"))
        mask_names = sorted(n for n in self.zf.namelist()
                              if n.startswith(m_pre) and n.endswith(".png"))
        frames = np.stack([np.array(Image.open(io.BytesIO(self.zf.read(n))))
↪convert("RGB"))
                               for n in frame_names])
        masks = np.stack([np.array(Image.open(io.BytesIO(self.zf.read(n)))
                                   ) for n in mask_names])
        print(f"video {video} {len(frames)} frames")
        return frames, masks

    def process_frames(self, frames, dino_model, device):
        res = []
        for f in frames:
            f = self.img_tsfm(Image.fromarray(f))[None].to(device)
            with torch.no_grad():
                tok = dino_model.get_intermediate_layers(f, n=1)[0]
            res.append(tok[0, 1:])
        res = torch.stack(res)
        return res

    def process_masks(self, masks, device):
        res = []
        for m in masks:
            m = cv2.resize(m, (60, 60), cv2.INTER_NEAREST)
            res.append(torch.from_numpy(m).long().flatten(-2, -1))
        res = torch.stack(res).to(device)
        return res

    def mask_frame_overlay(self, processed_mask, frame):
        H, W = frame.shape[:2]

```

```

        mask = processed_mask.detach().cpu().numpy().reshape((60, 60))
        mask = cv2.resize(mask.astype(np.uint8), (W, H),
                           interpolation=cv2.INTER_NEAREST)
        overlay = create_segmentation_overlay(mask, frame.copy())
        return overlay

davis_ds = DavisDataset()

frames, masks = davis_ds.get_sample(7)
num_classes = masks.max() + 1

print(frames.shape, masks.shape, num_classes)

```

30 validation videos
video cows 104 frames
(104, 480, 854, 3) (104, 480, 854) 2

```

[30]: # Get DINO patch features and corresponding class labels for a middle frame
train_fi = 40
X_train = davis_ds.process_frames(frames[train_fi:train_fi+1], dino_model,
    ↪device)[0]
Y_train = davis_ds.process_masks(masks[train_fi:train_fi+1], device)[0]
print(X_train.shape, Y_train.shape)

```

```
torch.Size([3600, 384]) torch.Size([3600])
```

Complete the implementation of the `DINOSegmentation` class in `cs231n/clip_dino.py`, and test it by running the two cells below. You should achieve a mean IoU greater than 0.45 on the first test frame and greater than 0.50 on the last test frame. To prevent overfitting on the training patch features, consider designing a very lightweight model (e.g., a linear layer or a 2-layer MLP) and applying appropriate weight decay.

You may use GPU runtime to speed up training and evaluation. Make sure to rerun the entire notebook if you change runtime type.

```

[31]: from cs231n.clip_dino import DINOSegmentation, compute_iou
torch.manual_seed(231)
np.random.seed(231)
dino_segmentation = DINOSegmentation(device, num_classes)
dino_segmentation.train(X_train, Y_train, num_iters=500)

# Test on first, middle, and last frame
ious = []
test_fis = [0, train_fi, 98]
gt_viz = []
pred_viz = []

```

```

for fi in test_fis:
    X_test = davis_ds.process_frames(frames[fi:fi+1], dino_model, device)[0]
    Y_test = davis_ds.process_masks(masks[fi:fi+1], device)[0]
    Y_pred = dino_segmentation.inference(X_test)
    iou = compute_iou(Y_pred, Y_test, num_classes)
    ious.append(iou)

    gt_viz.append(davis_ds.mask_frame_overlay(Y_test, frames[fi]))
    pred_viz.append(davis_ds.mask_frame_overlay(Y_pred, frames[fi]))

gt_viz = np.concatenate(gt_viz, 1)
pred_viz = np.concatenate(pred_viz, 1)

```

```

[32]: print(f"Mean IoU on first test frames: {ious[0]:.3f}") # should be >0.45
      print(f"Mean IoU on last test frames: {ious[2]:.3f}") # should be >0.50

```

Mean IoU on first test frames: 0.955

Mean IoU on last test frames: 0.955

Now let's visualize the results. Run the two cells below to display the ground truth and predicted segmentation masks for the first, middle, and last frames. Note that the middle frame is part of the training set, while the other frames are unseen.

```

[33]: Image.fromarray(gt_viz)

```

Output hidden; open in <https://colab.research.google.com> to view.

```

[34]: Image.fromarray(pred_viz)

```

Output hidden; open in <https://colab.research.google.com> to view.

Now run the following three cells to evaluate and visualize the entire video. You should achieve a mean IoU greater than 0.55. The saved visualization video may take some time to process in Google Drive, but you can download it to your computer and view it locally.

```

[35]: # Run on all frames
      ious = []
      gt_viz = []
      pred_viz = []
      for fi in range(len(frames)):
          if fi % 20 == 0:
              print(f"{fi} / {len(frames)}")
              X_test = davis_ds.process_frames(frames[fi:fi+1], dino_model, device)[0]
              Y_test = davis_ds.process_masks(masks[fi:fi+1], device)[0]
              Y_pred = dino_segmentation.inference(X_test)
              iou = compute_iou(Y_pred, Y_test, num_classes)
              ious.append(iou)

              gt_viz.append(davis_ds.mask_frame_overlay(Y_test, frames[fi]))

```

```

    pred_viz.append(davis_ds.mask_frame_overlay(Y_pred, frames[fi]))

gt_viz = np.stack(gt_viz) # T x H x W x 3
pred_viz = np.stack(pred_viz) # T x H x W x 3
final_viz = np.concatenate([gt_viz, pred_viz], -2) # T x H x 2W x 3

```

```

0 / 104
20 / 104
40 / 104
60 / 104
80 / 104
100 / 104

```

```
[36]: print(f"Mean IoU on all frames: {sum(ious) / len(ious):.3f}") # should be >0.55
```

Mean IoU on all frames: 0.959

```
[37]: def write_video_from_array(array, output_path, fps = 12):
    T, H, W, _ = array.shape
    fourcc = cv2.VideoWriter_fourcc(*'mp4v')
    out = cv2.VideoWriter(output_path, fourcc, fps, (W, H))
    for i in range(T):
        frame = array[i]
        out.write(frame)
    out.release()
    print(f"Video saved to {output_path}")

# It might take a while to process in google drive but you can just download it_
↪and watch on your computer
write_video_from_array(final_viz, f"/content/drive/My Drive/{FOLDERNAME}/
↪dino_res.mp4")

```

Video saved to /content/drive/My
Drive/cs231n/assignments/assignment3//dino_res.mp4

Inline Question 5 -

If you train a segmentation model on CLIP ViT's patch features, do you expect it to perform better or worse than DINO? Why should that be the case?

Your Answer :

I expect it to perform worse than DINO (at least with vanilla CLIP). CLIP is trained with a global image-text contrastive loss: the objective only shapes the global pooled representation so that images match their captions, so patch-level features receive no direct localization supervision; moreover, CLIP typically sees low-resolution (224×224) square crops and uses attention pooling, which further blurs spatial detail. DINO, in contrast, is self-supervised on images alone with multi-crop self-distillation, which forces patch-level features to be discriminative and spatially consistent — as seen in the PCA visualization, its features already delineate object boundaries. This is

exactly the motivation of CLIP-DINO-style methods: combine DINO’s dense, spatially precise patch features with CLIP’s open-vocabulary semantics. CLIP’s patch features can be improved with dense-prediction adaptations, but out of the box DINO is the better backbone for segmentation.

[37]:

